



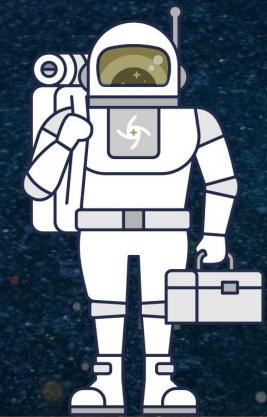
STARGATE

Apache Cassandra et Kubernetes

DataStax
Developers



K8SSANDRA



Cédrick Lunven

Director Developer Relations



- Trainer
- Public Speaker
- Developers Support
- Developer Applications
- Developer Tooling

SDK



- Creator of ff4j (ff4j.org)
- Maintainer for 8 years+



- Happy developer for 14 years
- Spring Petclinic Reactive & Starters
- Implementing APIs for 8 years

DataStax Developers Advocates Team



Cedrick
Lunven

Aleksandr
Volochnev

Jack
Fryer

Kirsten
Hunter

Stefano
Lottini

David
Gilardi

Ryan
Welford

Rags
Srinivas

Sonia
Siganporia

R

S

Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

05



K8ssandra
Tour d'horizon

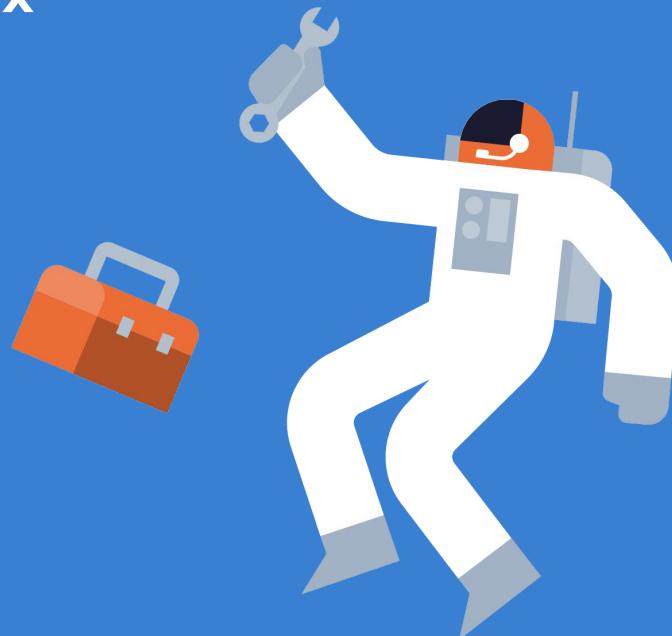


DataStax

Formation CapGemini 5 octobre

I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. Scalability
3. Distributed Data
4. Consistency
5. Use Cases
6. CQL



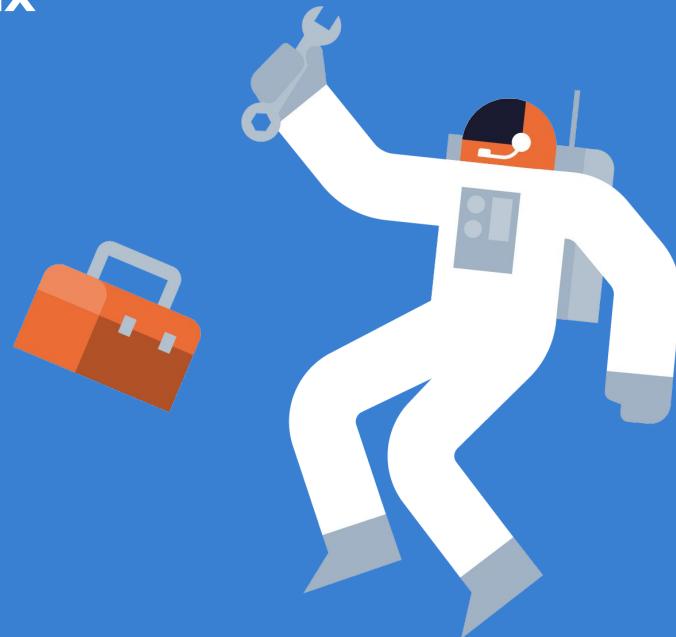


DataStax

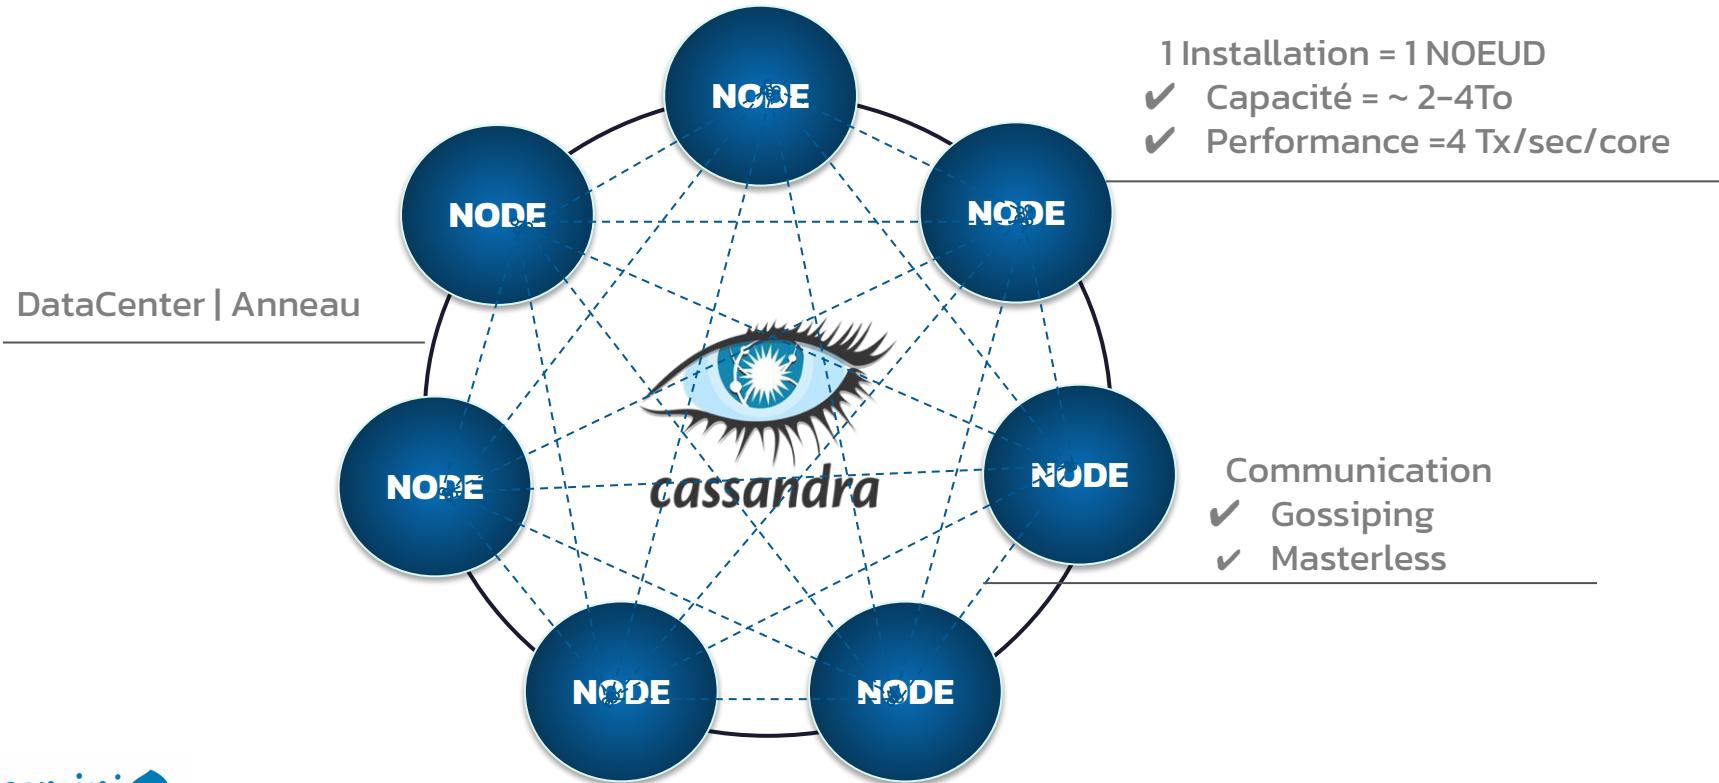
Formation CapGemini 5 octobre

I - Apache Cassandra™ Les fondamentaux

1. **Introduction**
2. Scalability
3. Distributed Data
4. Consistency
5. Use Cases
6. CQL



Apache Cassandra™ = Base de données NoSql Distribuée



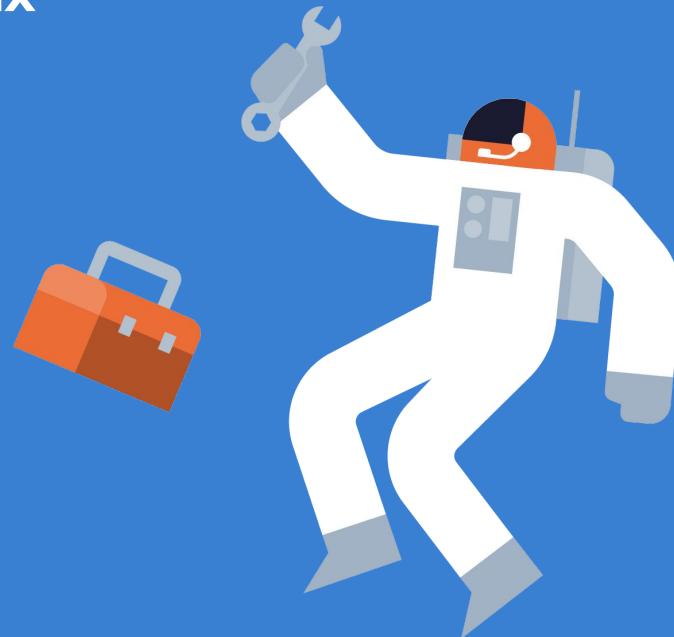


DataStax

Formation CapGemini 5 octobre

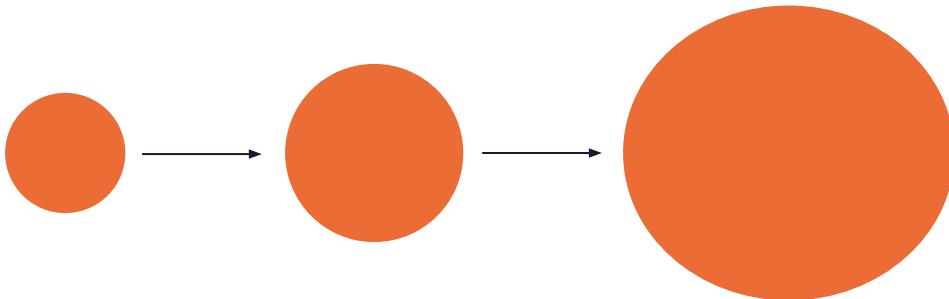
I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. **Scalability**
3. Distributed Data
4. Consistency
5. Use Cases
6. CQL



Vertical Scalability?

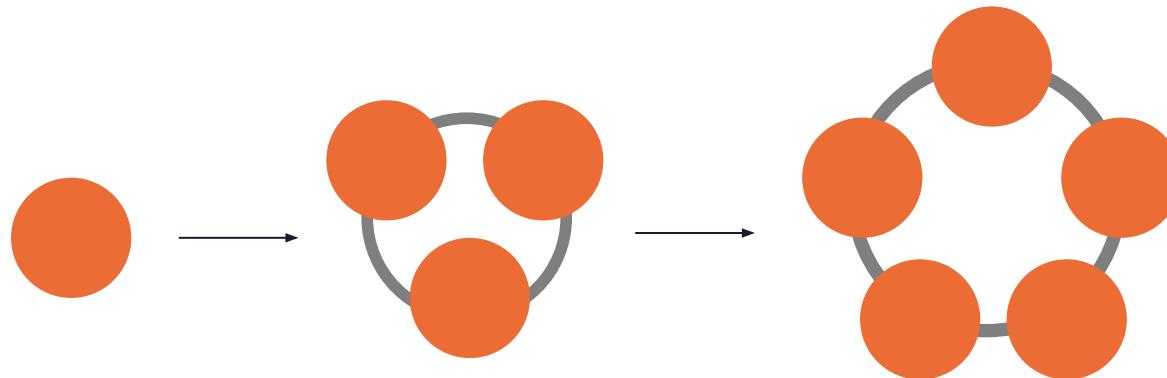
Vertical scaling requires one large expensive machine



- + Easier to get
- Expensive
- Still SPoF

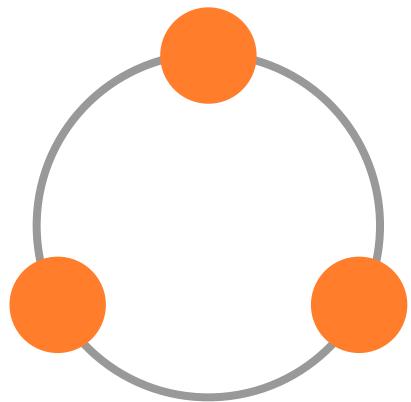
Horizontal Scalability!

Horizontal scaling requires multiple less-expensive commodity hardware

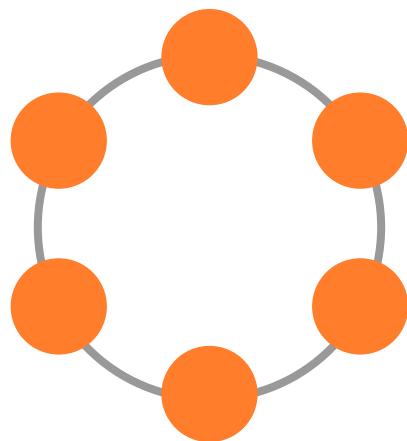


- + Cheaper
- + Removes SPoF
- Logistically harder

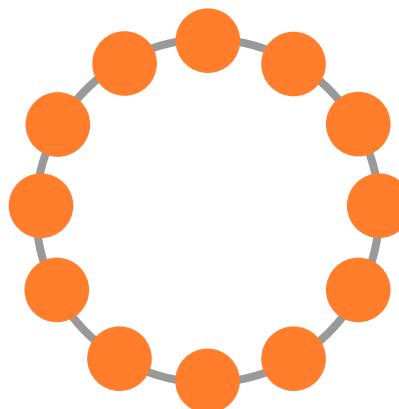
Cassandra Scales Horizontally



100,000
transactions/second



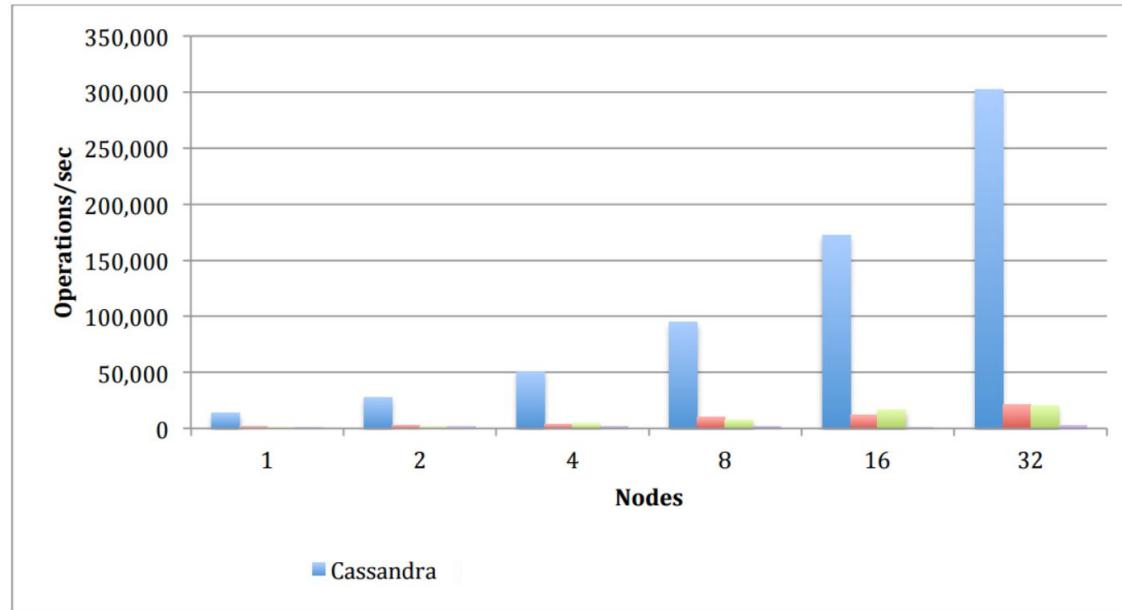
200,000
transactions/second



400,000
transactions/second

Apache Cassandra™ = Scalabilité Linéaire

Balanced Read/Write Mix

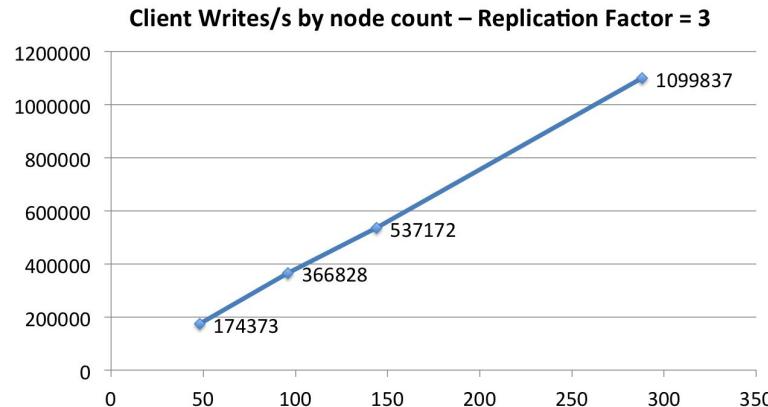




Cassandra Scales Linearly

Recommended reading: netflixtechblog.com/benchmarking-cassandra-scalability-on-aws-over-a-million-writes-per-second-39f45f066c9e

Scale-Up Linearity





Apache Cassandra™ = Scalabilité inégalée

Apache Cassandra @ Netflix

- 98% of streaming data is stored in Apache Cassandra
- Data ranges from customer details to viewing history to billing and payments
- Foundational datastore for serving millions of operations per second

- 30 million ops/sec on most active single cluster
- 500 TB most dense single cluster
- 9216 CPUs in biggest cluster

O(100) Clusters
O(10000) Instances
O(10,000,000) Replications per second
O(100,000,000) Operations per second
O(1,000,000,000,000,000) Petabytes of data

Apple Scale

- 160K+ Apache Cassandra instances
- 100+ PB stored
- Several million ops / sec
- 1000s of clusters



DataStax

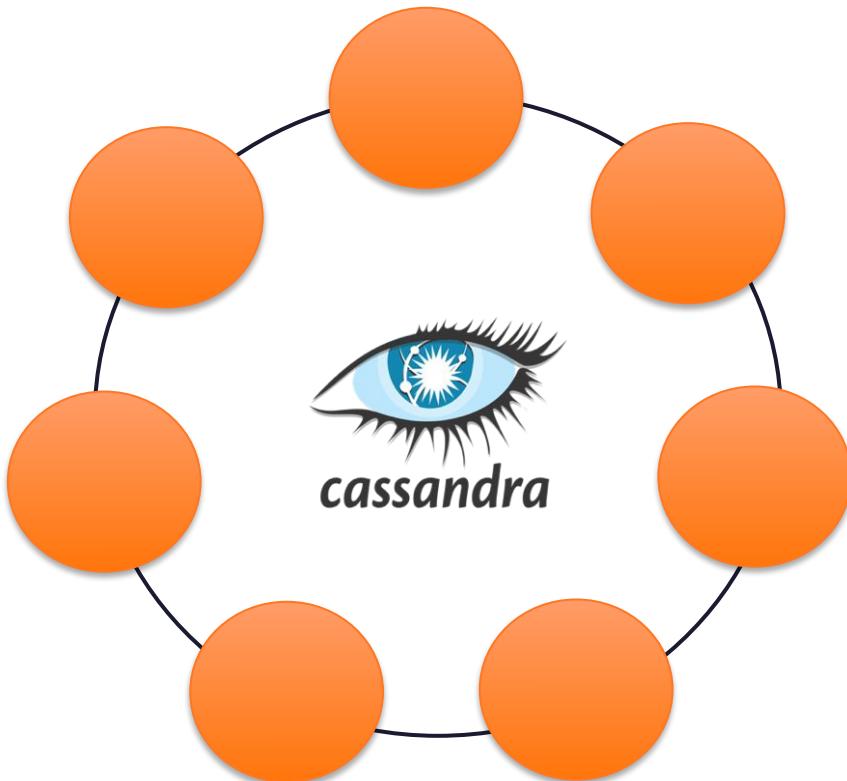
Formation CapGemini 5 octobre

I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. Scalability
3. **Distributed Data**
4. Consistency
5. Use Cases
6. CQL



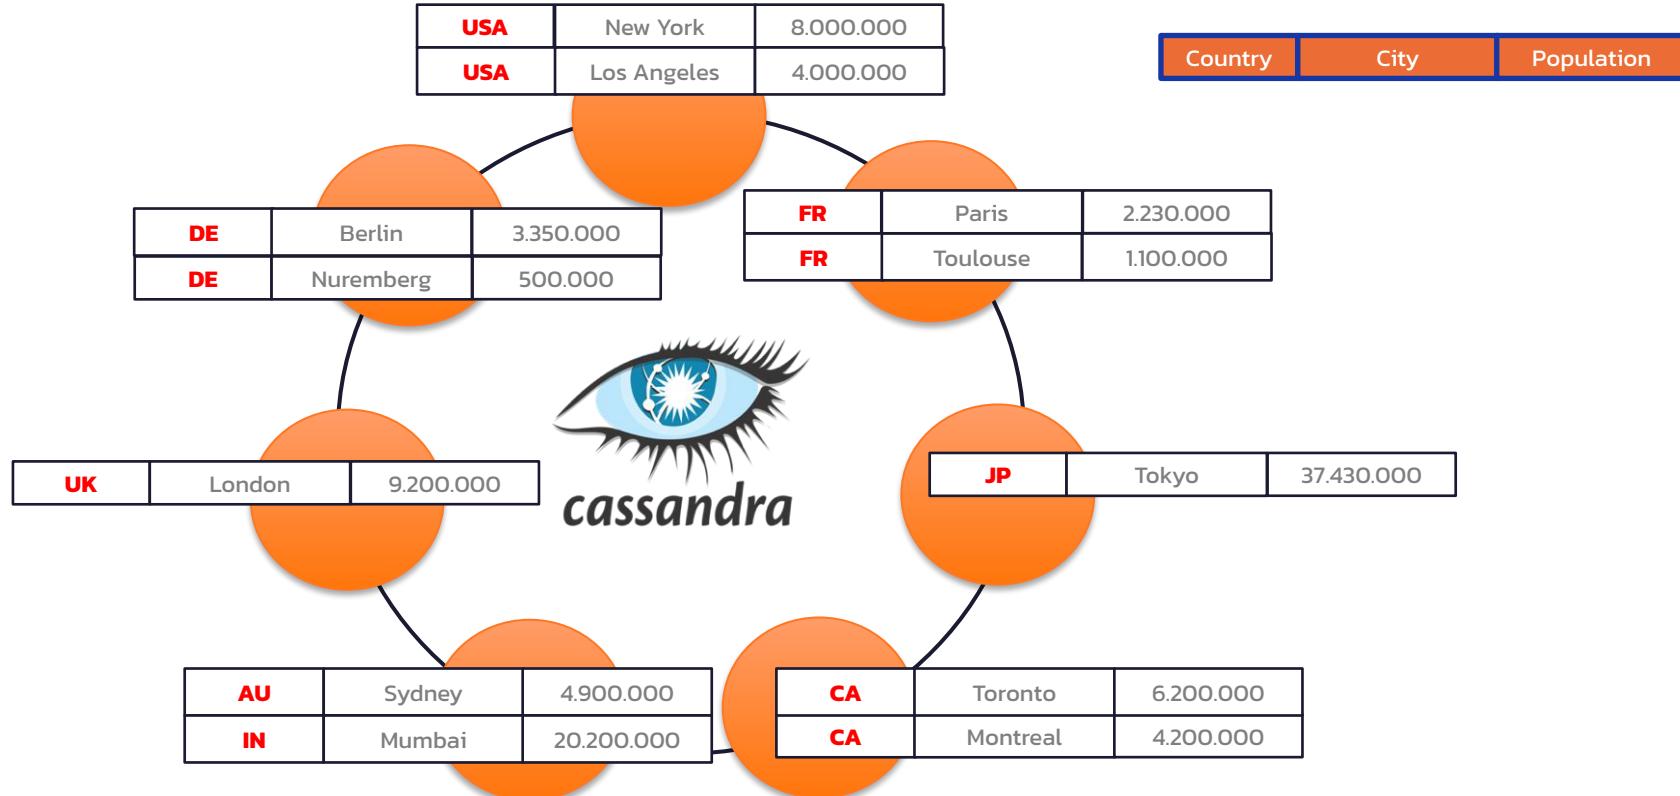
Data is Distributed



| Country | City | Population |
|---------|-------------|------------|
| USA | New York | 8.000.000 |
| USA | Los Angeles | 4.000.000 |
| FR | Paris | 2.230.000 |
| DE | Berlin | 3.350.000 |
| UK | London | 9.200.000 |
| AU | Sydney | 4.900.000 |
| DE | Nuremberg | 500.000 |
| CA | Toronto | 6.200.000 |
| CA | Montreal | 4.200.000 |
| FR | Toulouse | 1.100.000 |
| JP | Tokyo | 37.430.000 |
| IN | Mumbai | 20.200.000 |

Partition Key

Data is Distributed



Partitioning

| CO | City | Population |
|----|-----------|------------|
| AU | Sydney | 4.900.000 |
| CA | Toronto | 6.200.000 |
| CA | Montreal | 4.200.000 |
| DE | Berlin | 3.350.000 |
| DE | Nuremberg | 500.000 |

Partition Key

Partitioner
Murmur3 Hashing

| CO | City | Population |
|----|-----------|------------|
| 59 | Sydney | 4.900.000 |
| 12 | Toronto | 6.200.000 |
| 12 | Montreal | 4.200.000 |
| 45 | Berlin | 3.350.000 |
| 45 | Nuremberg | 500.000 |

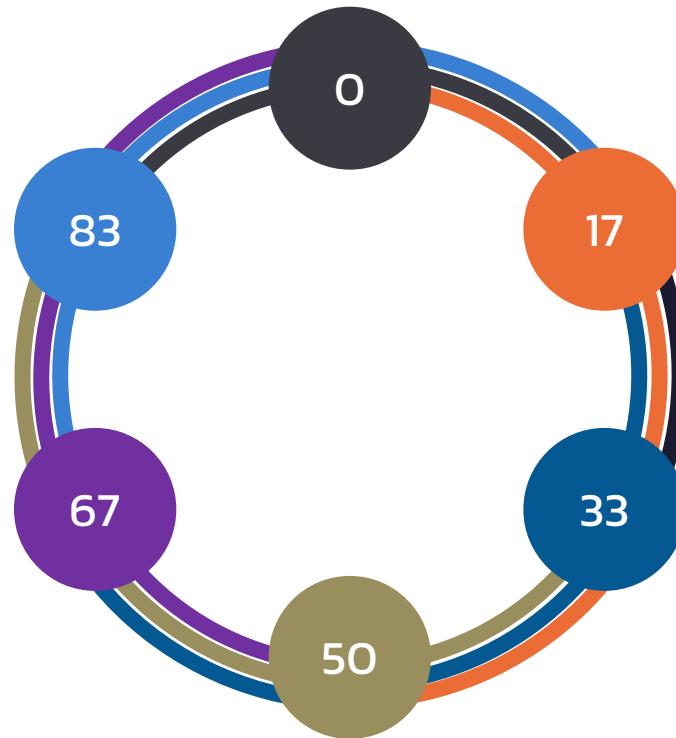
Tokens



Data is Replicated

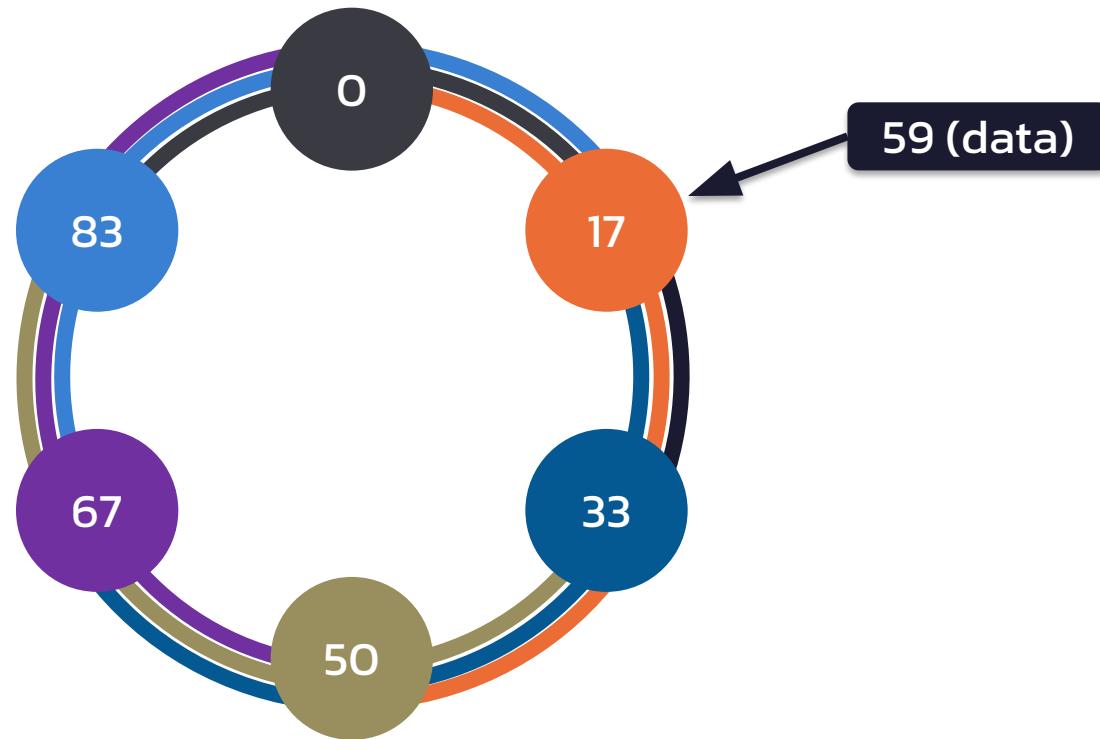
RF = 3

Replication Factor 3
means that every
row is stored on 3
different nodes

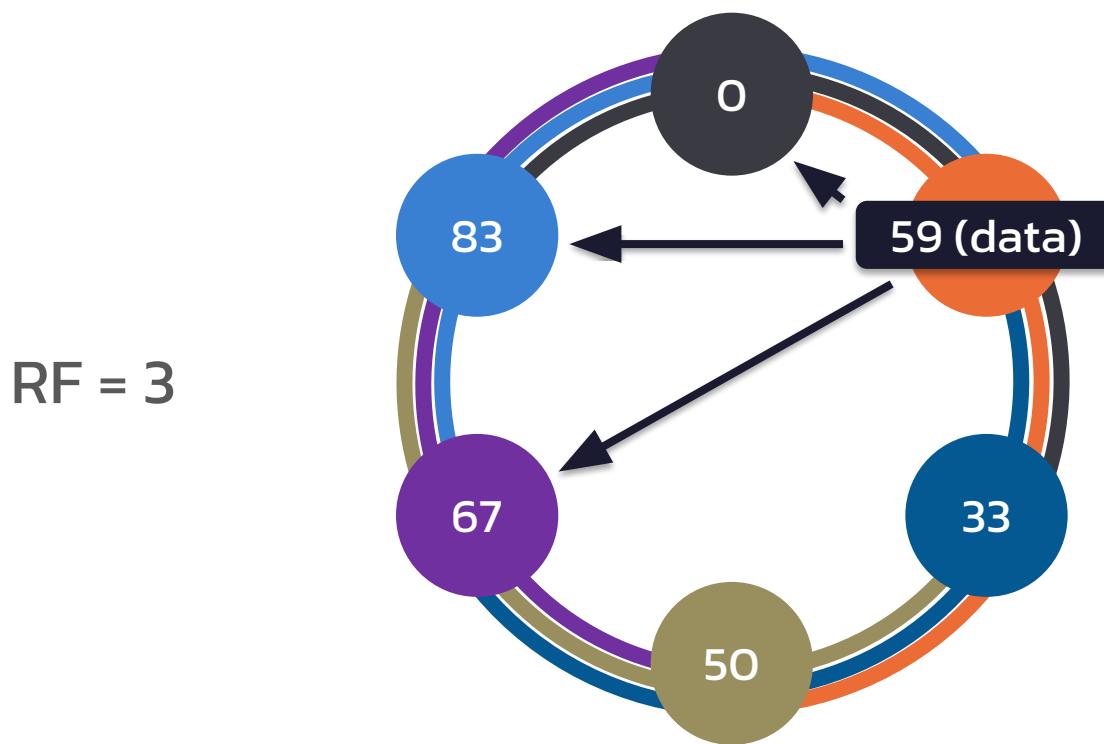


Replication within the Ring

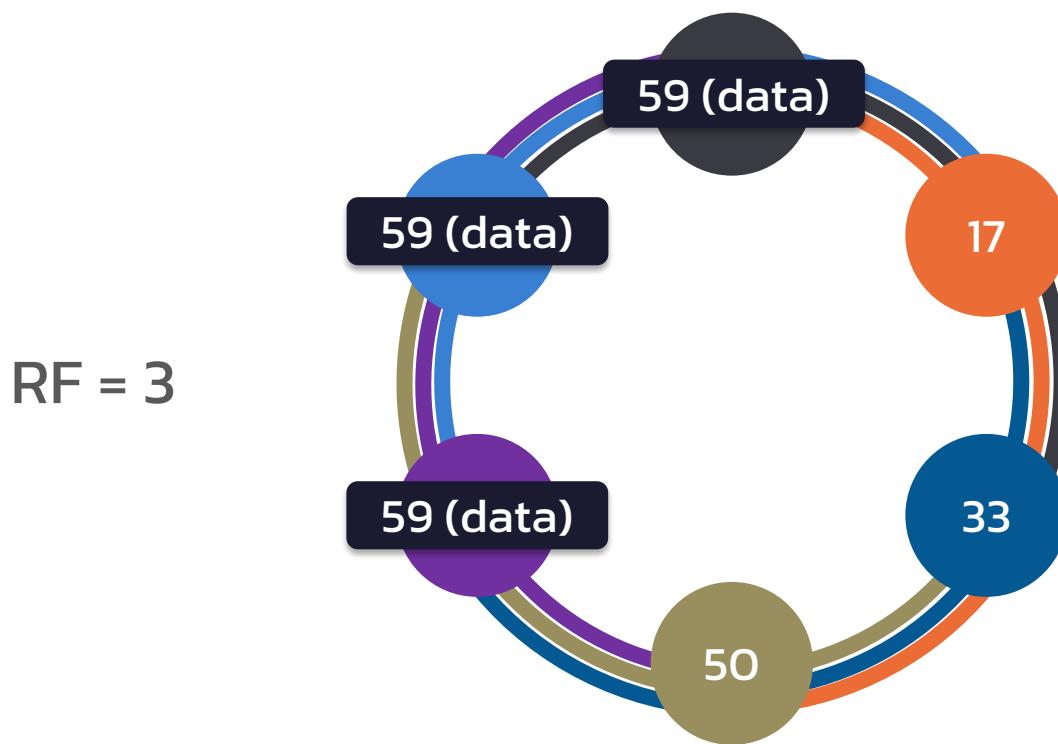
RF = 3



Replication within the Ring

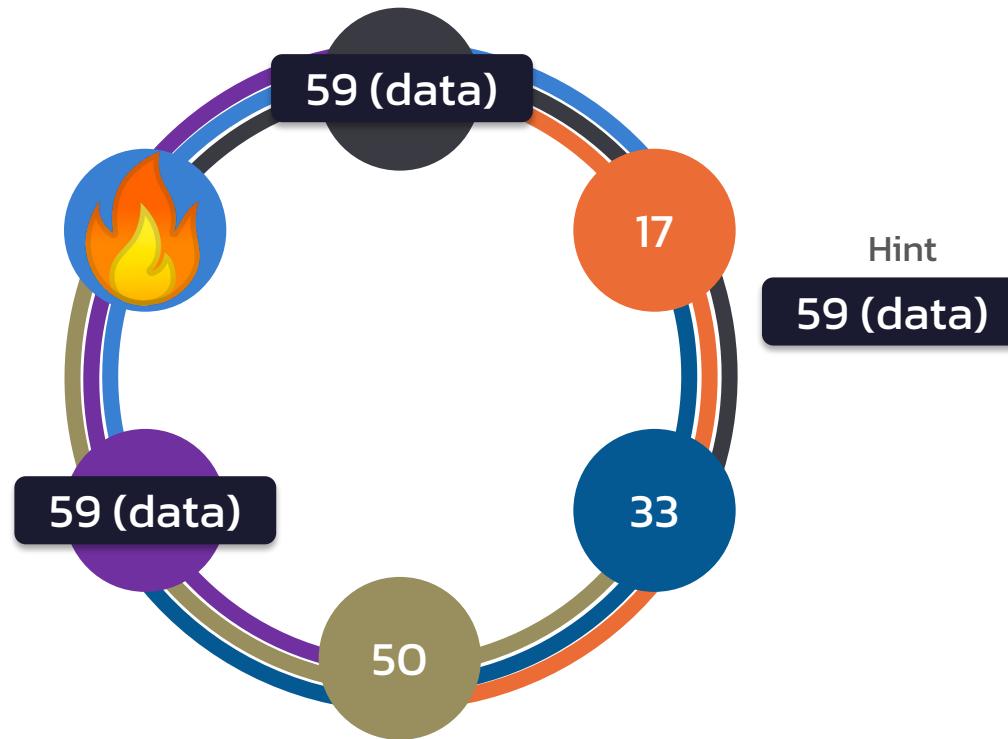


Replication within the Ring



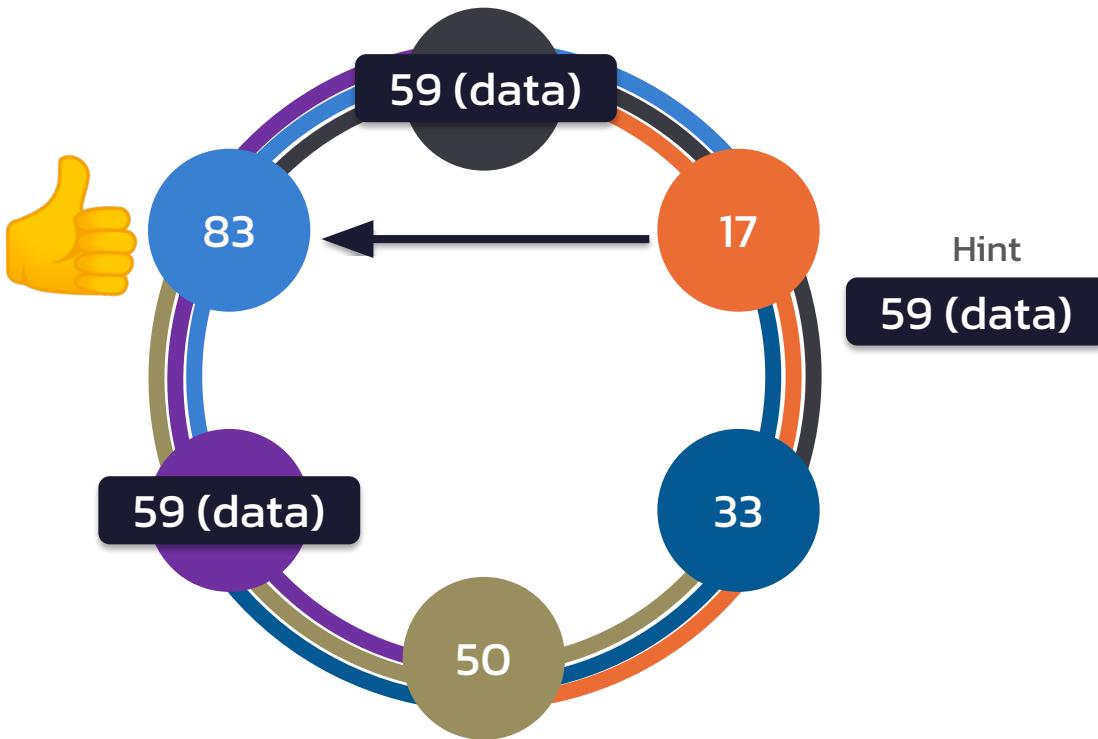
Node Failure

RF = 3



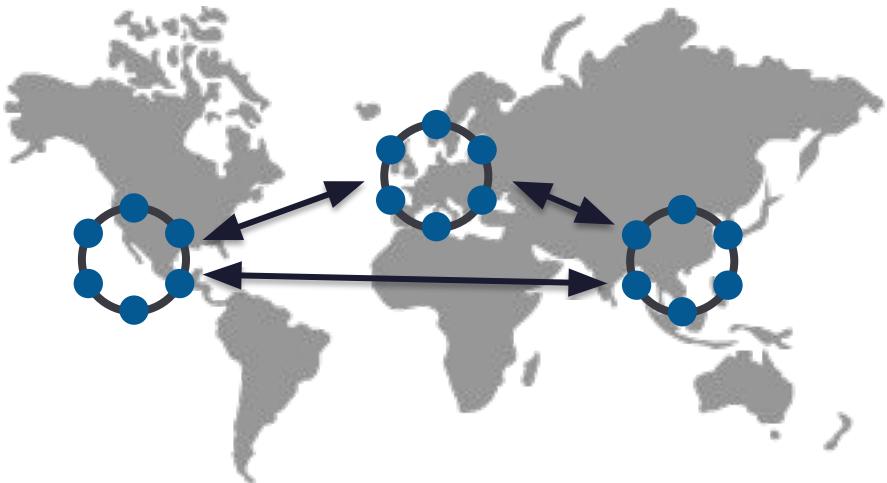
Node Failure Recovered

RF = 3

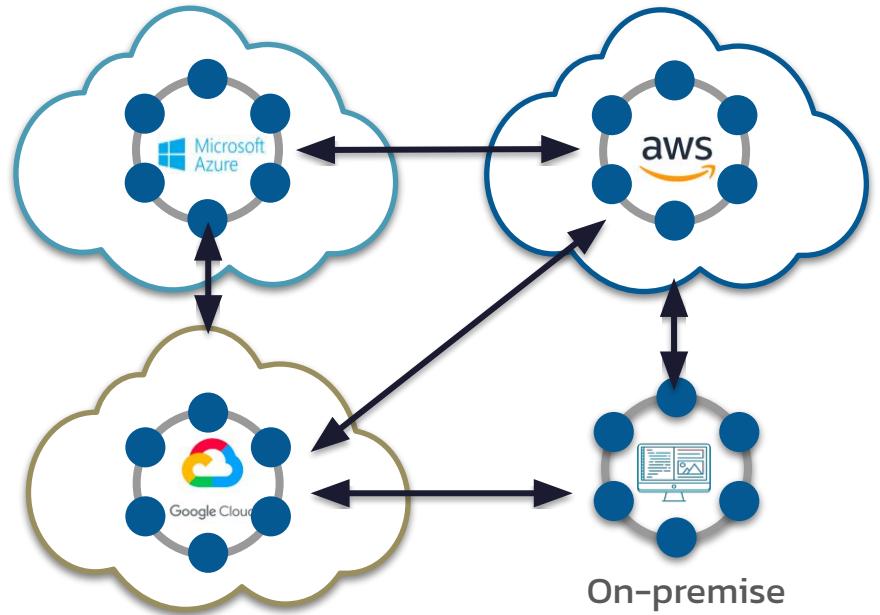


Data Distributed Everywhere

- Geographic Distribution



- Hybrid-Cloud and Multi-Cloud



Normalization

"Database normalization is the process of structuring a relational database in accordance with a series of so-called normal forms in order to reduce data redundancy and improve data integrity. It was first proposed by Edgar F. Codd as part of his relational model."

| Employees | | |
|-----------|-----------|----------|
| userId | firstName | lastName |
| 1 | Edgar | Codd |
| 2 | Raymond | Boyce |

| Departments | |
|--------------|-------------|
| departmentId | department |
| 1 | Engineering |
| 2 | Math |

PROS: Simple write, Data Integrity

CONS: Slow read, Complex Queries

Denormalization

"Denormalization is a strategy used on a database to increase performance. In computing, denormalization is the process of trying to improve the read performance of a database, at the expense of losing some write performance, by adding redundant copies of data"

PROS: Quick Read, Simple Queries

CONS: Multiple Writes, Manual Integrity

Employees

| userId | firstName | lastName | department |
|--------|-----------|----------|-------------|
| 1 | Edgar | Codd | Engineering |
| 2 | Raymond | Boyce | Math |
| 3 | Sage | Lahja | Math |
| 4 | Juniper | Jones | Botany |

Keyspaces

```
CREATE KEYSPACE users
  WITH REPLICATION = {
    'class' : 'NetworkTopologyStrategy',
    'us-west-1' : 3,
    'eu-central-1' : 5
};
```

keyspace

replication strategy

Replication factor by data center

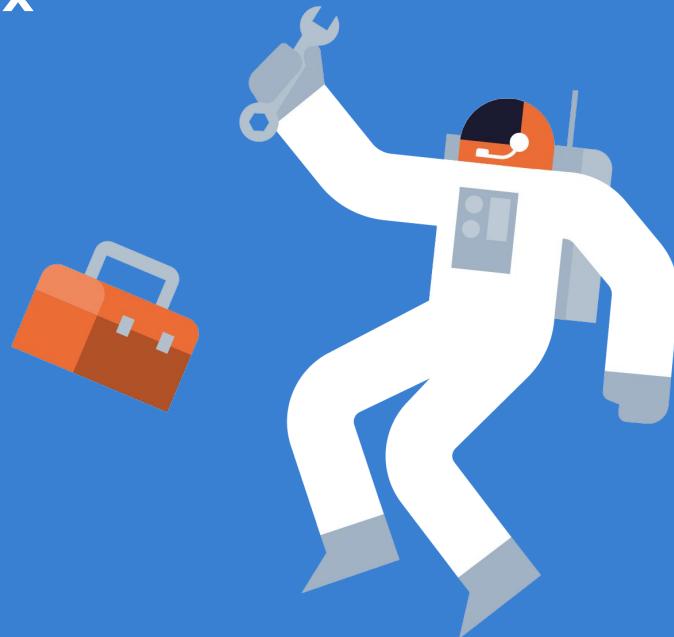


DataStax

Formation CapGemini 5 octobre

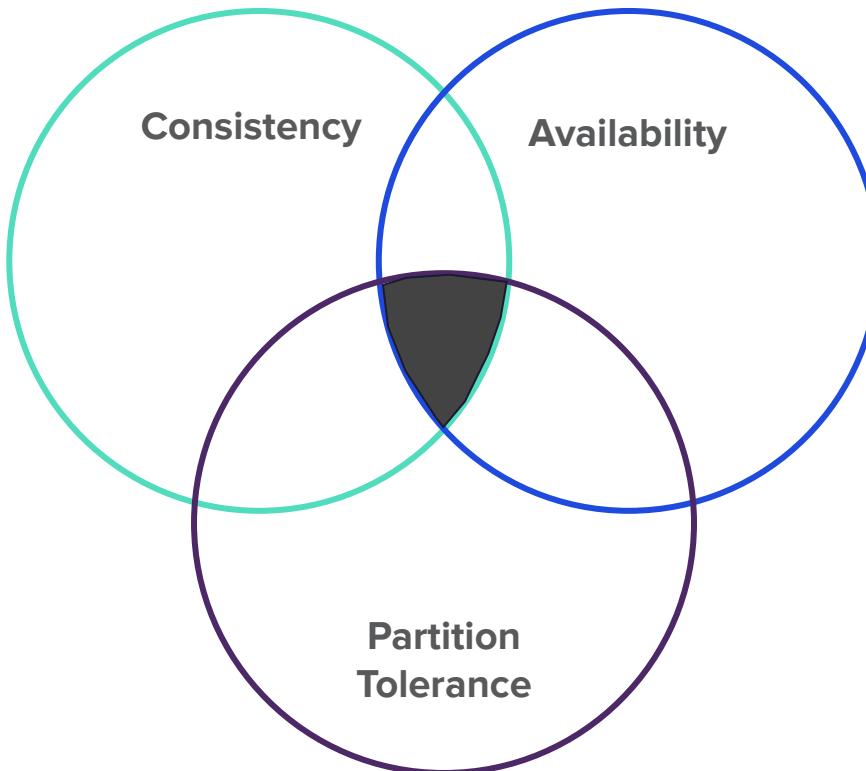
I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. Scalability
3. Distributed Data
4. **Consistency**
5. Use Cases
6. CQL



CAP Theorem

In the distributed environment, you can have only two guaranteed qualities out of three :(

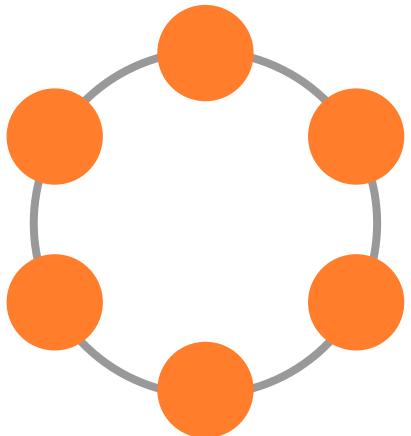


Is Cassandra AP or CP?

Cassandra is configurably consistent. In any moment, of the time, for any particular query you can set your the Consistency Level you require to have.

Cassandra Consistency Levels:

- ANY
- ONE
- TWO, THREE
- QUORUM
- ALL

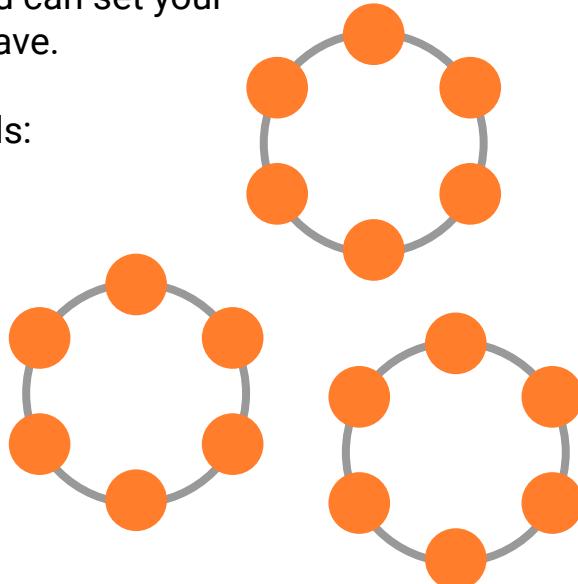


Is Cassandra AP or CP?

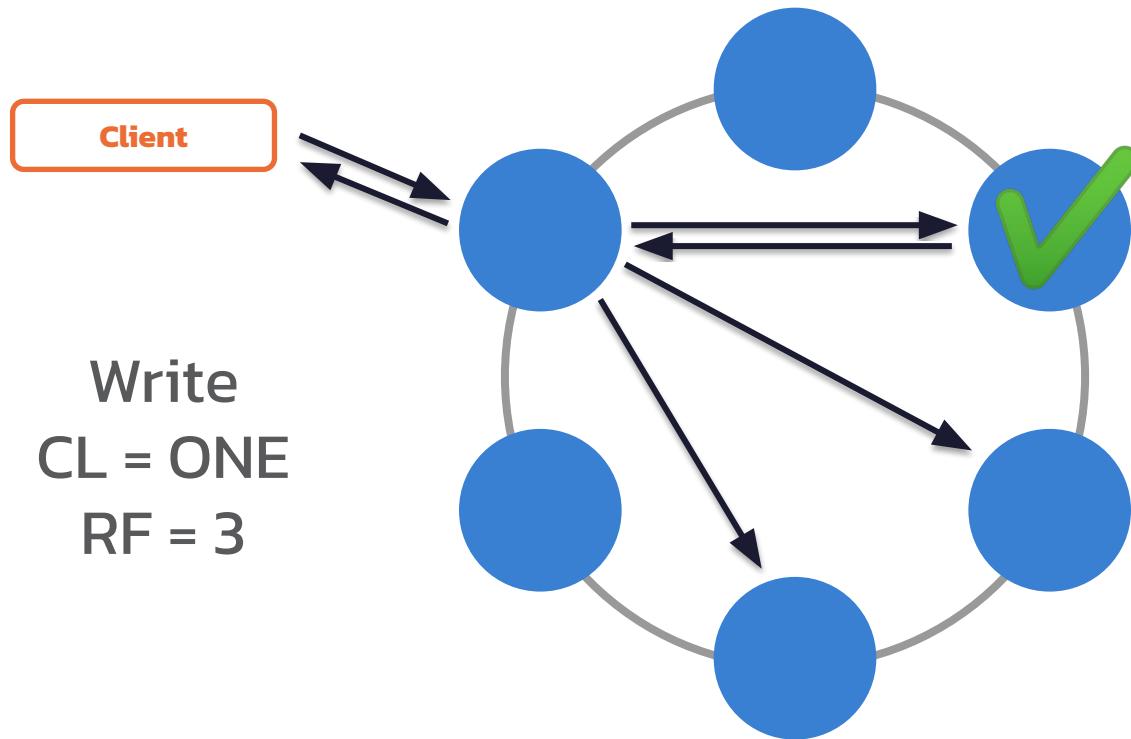
Cassandra is configurably consistent. In any moment, of the time, for any particular query you can set your the Consistency Level you require to have.

Cassandra Multi-DC Consistency Levels:

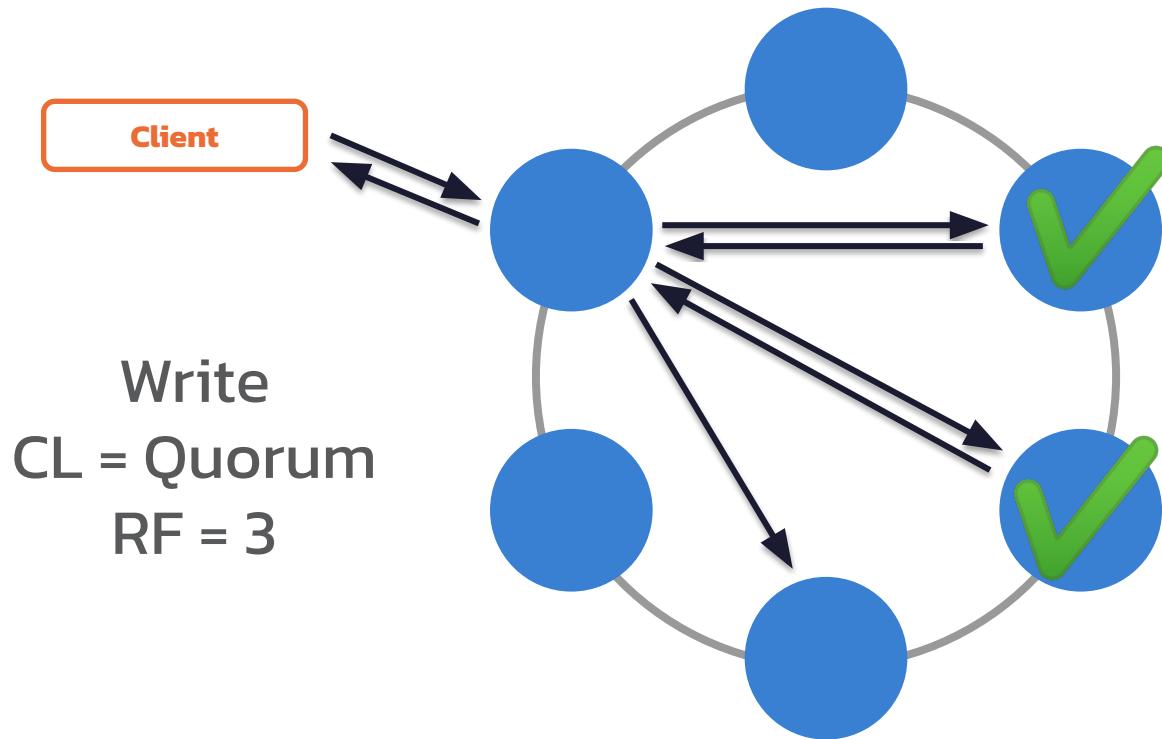
- LOCAL_ONE
- LOCAL_QUORUM
- EACH_QUORUM



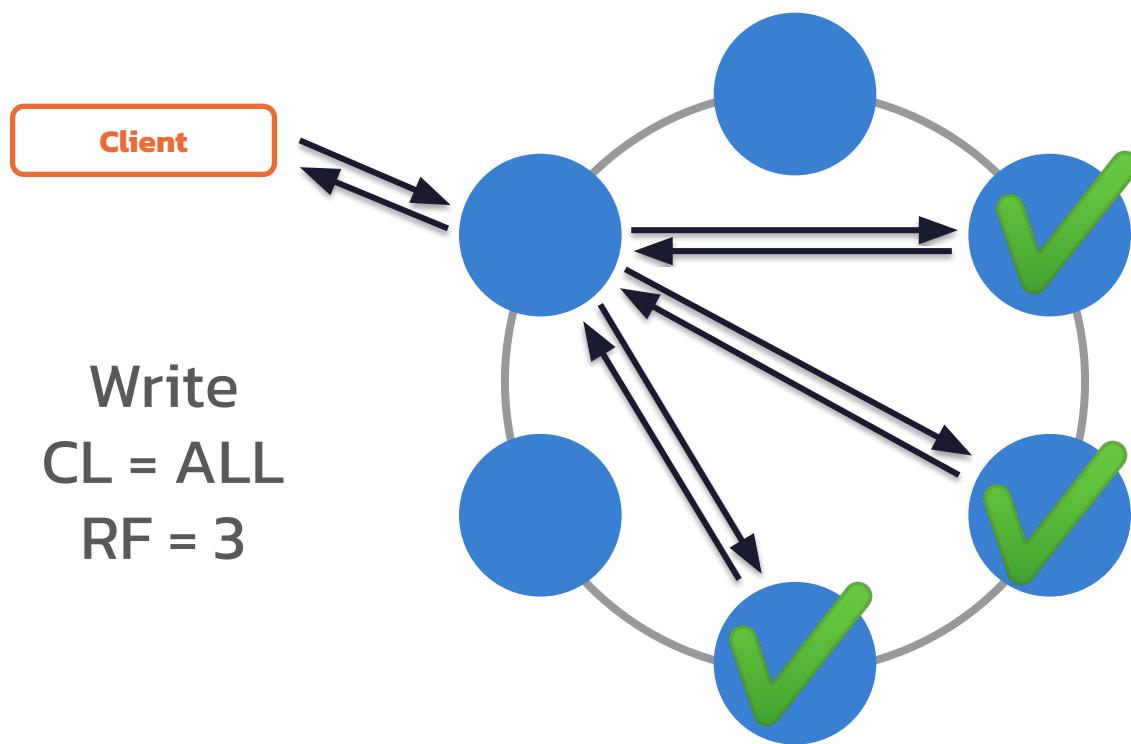
Consistency Level One



Consistency Level Quorum



Consistency Level ALL



Consistency Level ALL

Client

IT'S A TRAP!

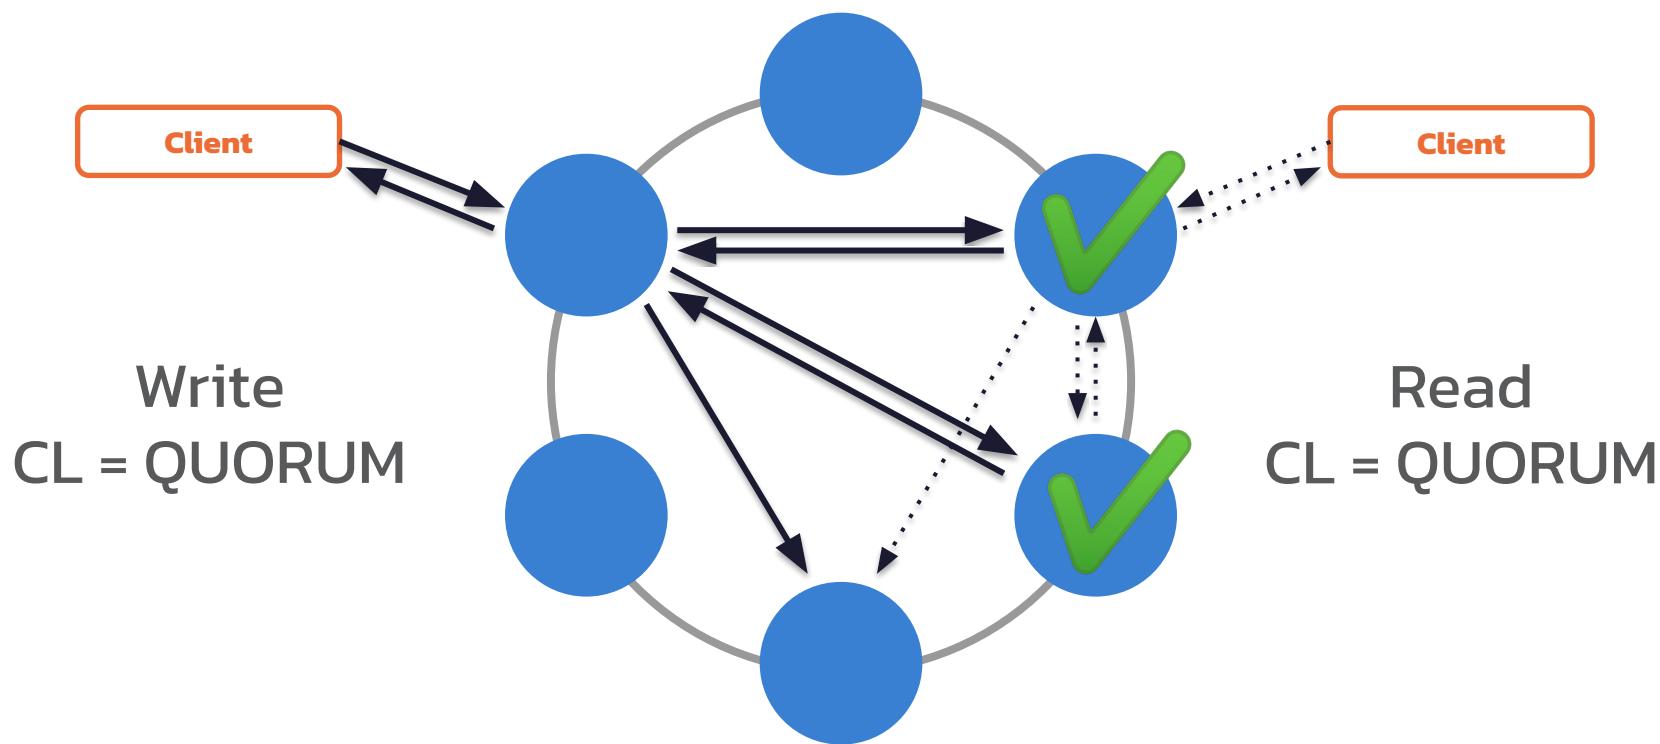
Write

CL = ALL

RF = 3

CAP THEOREM IS STILL HERE!

Immediate Consistency – The Right Way



Immediate Consistency – The Right Way

CL Write + CL Read > RF ← Immediate Consistency



DataStax

Formation CapGemini 5 octobre

I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. Scalability
3. Distributed Data
4. Consistency
- 5. Use Cases**
6. CQL



Apache Cassandra™



Open Source

Apache Licence



Flexible, Familiar interface

CQL

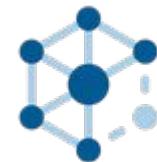


High Performance

Masterless, Replication

Active Everywhere

Zero Downtime



Scalability

Horizontal



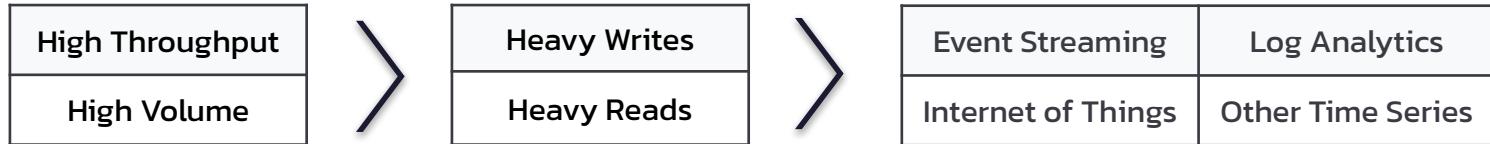
Seamless Replication

Asynchronous



Cas d'usages Apache Cassandra™ :

Scalabilité



Disponibilité



Distribuée



Cloud-native





Hands-on

<https://www.datastax.com/learn/cassandra-fundamentals#coursework>



Introduction to Apache Cassandra™

Learn about Cassandra benefits and install your first Cassandra NoSQL database instance

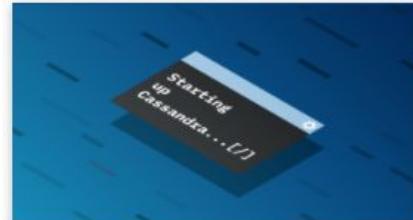
[START CODING](#)



Cassandra Query Language

Learn about the most essential data definition and data manipulation statements in Cassandra Query Language (CQL)

[START CODING](#)



Apache Cassandra™ Keyspaces and Data Replication Strategies

Learn how to create a keyspace and select a data replication strategy for a Cassandra NoSQL database

[START CODING](#)

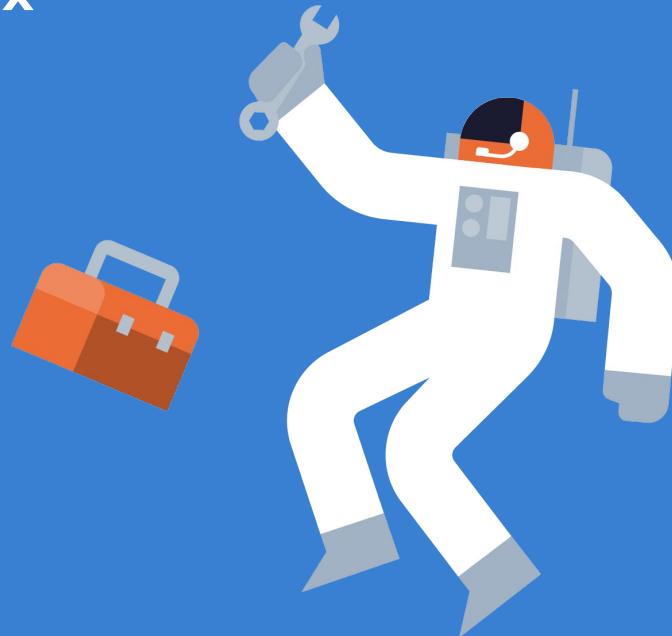


DataStax

Formation CapGemini 5 octobre

I - Apache Cassandra™ Les fondamentaux

1. Introduction
2. Scalability
3. Distributed Data
4. Consistency
5. Use Cases
6. CQL



Data Structure: a Cell



An intersection of a row
and a column, stores data.



Data Structure: a Row



A single, structured
data item in a table.

| | | | |
|---|------|-----|----------|
| 1 | John | Doe | Wizardry |
|---|------|-----|----------|

Data Structure: a Partition



| ID | First Name | Last Name | Department |
|-----|------------|-----------|------------|
| 1 | John | Doe | Wizardry |
| 399 | Marisha | Chavez | Wizardry |
| 415 | Maximus | Flavius | Wizardry |
| 999 | Anna | Chase | Dark Magic |
| 42 | Aleks | Volochnev | Dark Magic |

A group of rows having the same partition token, a base unit of access in Cassandra.

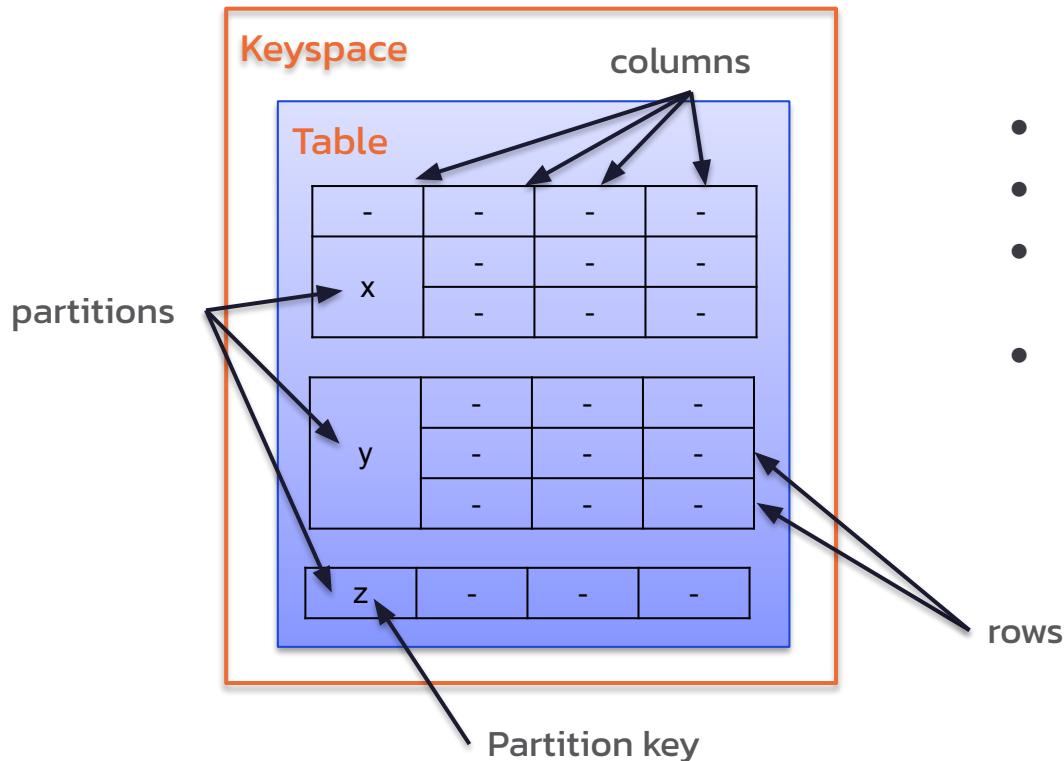
Data Structure: a Table



A group of columns and rows storing partitions.

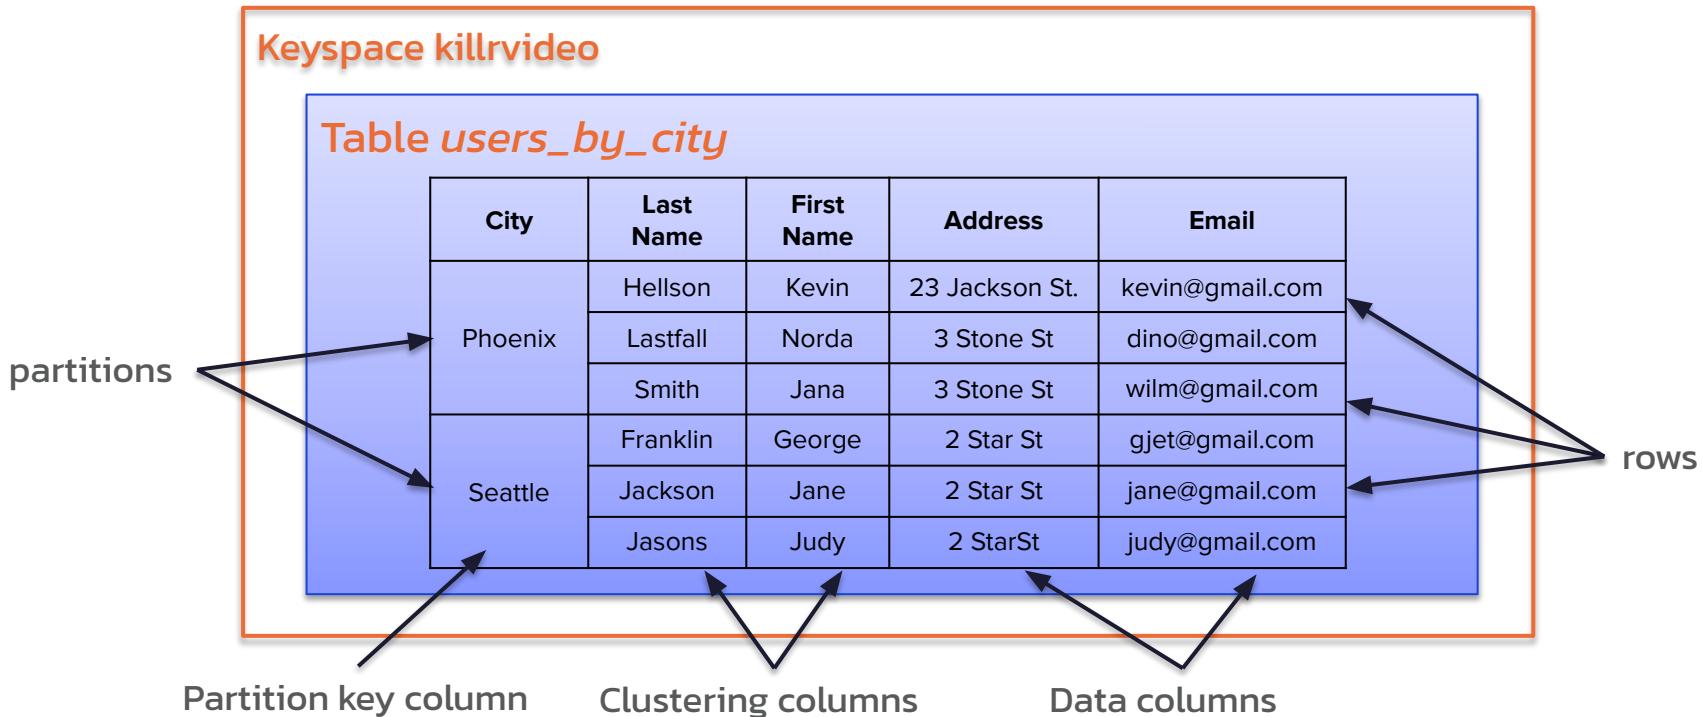
| ID | First Name | Last Name | Department |
|----|------------|-----------|------------|
| 1 | John | Doe | Wizardry |
| 2 | Mary | Smith | Dark Magic |
| 3 | Patrick | McFadin | DevRel |

Data Structure: Overall



- Tabular data model, with one twist
- *Tables* are organized in *rows* and *columns*
- Groups of related rows called *partitions* are stored together on the same node (or nodes)
- Each row contains a *partition key*
 - One or more columns that are hashed to determine which node(s) store that data

Example Data: Users organized by city



Keyspaces

```
CREATE KEYSPACE users
  WITH REPLICATION = {
    'class' : 'NetworkTopologyStrategy',
    'us-west-1' : 3,
    'eu-central-1' : 5
};
```

keyspace

replication strategy

Replication factor by data center

Creating a Table in CQL



Primary Key

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

MUST ENSURE UNIQUENESS.
MAY DEFINE SORTING.

```
CREATE TABLE killrvideo.users_by_city (
    city text,
    last_name text,
    first_name text,
    address text,
    email text,
    PRIMARY KEY ((city), last_name, first_name, email));
```



Good Examples:

```
PRIMARY KEY ((city), last_name, first_name, email);
```

```
PRIMARY KEY (user_id);
```

Bad Example:

```
PRIMARY KEY ((city), last_name, first_name);
```

Partition Key

An identifier for a partition.

Consists of at least one column,
may have more if needed

PARTITIONS ROWS.

```
CREATE TABLE killrvideo.users_by_city (
    city text,
    last_name text,
    first_name text,
    address text,
    email text,
    PRIMARY KEY ((city), last_name, first_name, email));
```



Good Examples:

```
PRIMARY KEY (user_id);
```

```
PRIMARY KEY ((video_id), comment_id);
```

Bad Example:

```
PRIMARY KEY ((sensor_id), logged_at);
```

Clustering Column(s)

Used to ensure uniqueness and sorting order. Optional.

```
CREATE TABLE killrvideo.users_by_city (
    city text,
    last_name text,
    first_name text,
    address text,
    email text,
    PRIMARY KEY ((city), last_name, first_name, email));
```

Partition key

Clustering columns

PRIMARY KEY ((city), last_name, first_name);



Not Unique

PRIMARY KEY ((city), last_name, first_name, email);



PRIMARY KEY ((video_id), comment_id);



Not Sorted

PRIMARY KEY ((video_id), created_at, comment_id);



Rules of a Good Partition

- **Store together what you retrieve together**
- Avoid big partitions
- Avoid hot partitions

Example: open a video? Get the comments in a single query!

```
PRIMARY KEY ((video_id), created_at, comment_id);
```



```
PRIMARY KEY ((comment_id), created_at);
```



Rules of a Good Partition

- Store together what you retrieve together
- **Avoid big partitions**
- Avoid hot partitions

```
PRIMARY KEY ((video_id), created_at, comment_id);
```



```
PRIMARY KEY ((country), user_id);
```



- Up to 2 billion cells per partition
- Up to ~100k rows in a partition
- Up to ~100MB in a Partition

Rules of a Good Partition

- Store together what you retrieve together
- **Avoid big and constantly growing partitions!**
- Avoid hot partitions

Example: a huge IoT infrastructure, hardware all over the world, different sensors reporting their state every 10 seconds. Every sensor reports its UUID, timestamp of the report, sensor's value.

- Sensor ID: UUID
- Timestamp: Timestamp
- Value: float

```
PRIMARY KEY ((sensor_id), reported_at);
```



Rules of a Good Partition

- Store together what you retrieve together
- **Avoid big and constantly growing partitions!**
- Avoid hot partitions

Example: a huge IoT infrastructure, hardware all over the world, different sensors reporting their state every 10 seconds. Every sensor reports its UUID, timestamp of the report, sensor's value.

```
PRIMARY KEY ((sensor_id), reported_at);
```



```
PRIMARY KEY ((sensor_id, month_year), reported_at);
```



BUCKETING

- Sensor ID: UUID
- **MonthYear:** Integer or String
- Timestamp: Timestamp
- Value: float

Rules of a Good Partition

- Store together what you retrieve together
- Avoid big partitions
- **Avoid hot partitions**

```
PRIMARY KEY (user_id);
```



```
PRIMARY KEY ((video_id), created_at, comment_id);
```



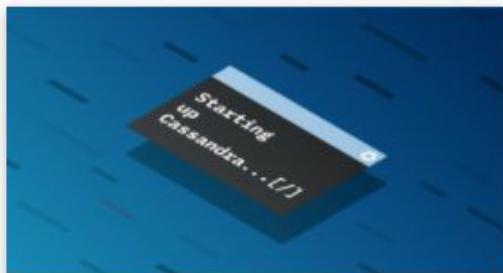
```
PRIMARY KEY ((country), user_id);
```





Hands-on

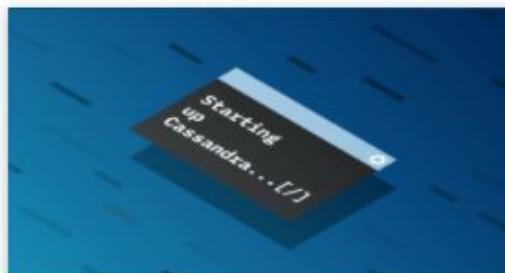
<https://www.datastax.com/learn/cassandra-fundamentals#coursework>



Inserts, Updates, Deletes and Upserts in Apache Cassandra™

Learn how to insert, update, delete, and upsert data into tables in a Cassandra NoSQL database

START CODING



Queries in Apache Cassandra™

Learn how to retrieve data from Cassandra NoSQL database tables

START CODING



Using Advanced Data Types in Apache Cassandra™

Learn about universally unique identifiers (UUIDs), collections, tuples, user-defined types (UDTs) and counters

START CODING

Next Step: Academy.datastax.com



*Designed for professionals that
use Apache Cassandra clusters
to manage data*

**application developers
data architects
database designers
database administrators**

Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

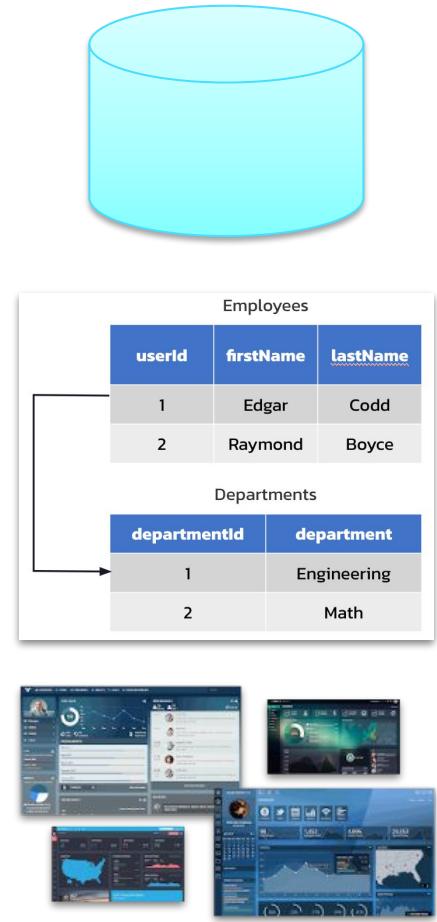
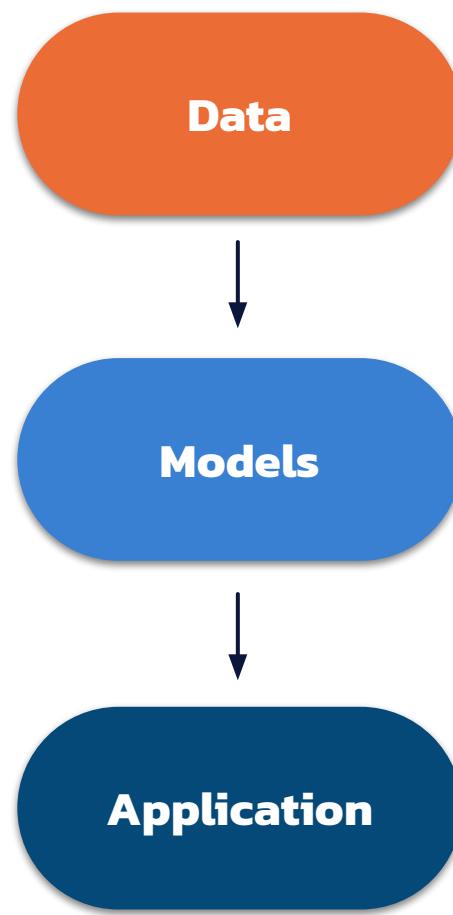
05



K8ssandra
Tour d'horizon

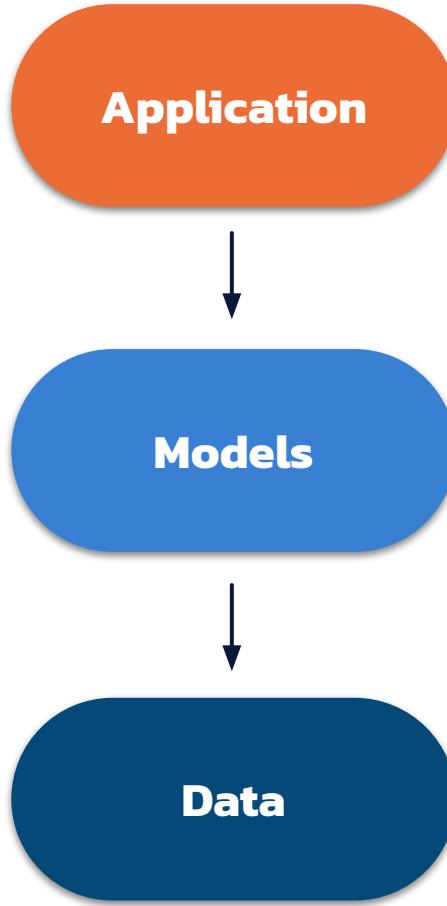
Relational Data Modelling

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using **normalization** and foreign keys.
4. Use JOIN when doing queries to join normalized data from multiple tables

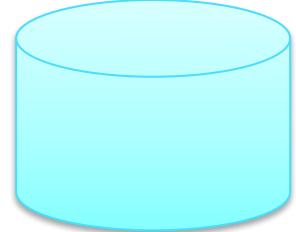


NoSQL Data Modelling

1. Analyze user behaviour
(customer first!)
2. Identify workflows, their dependencies and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables, using **denormalization**.
5. Use BATCH when inserting or updating denormalized data of multiple tables

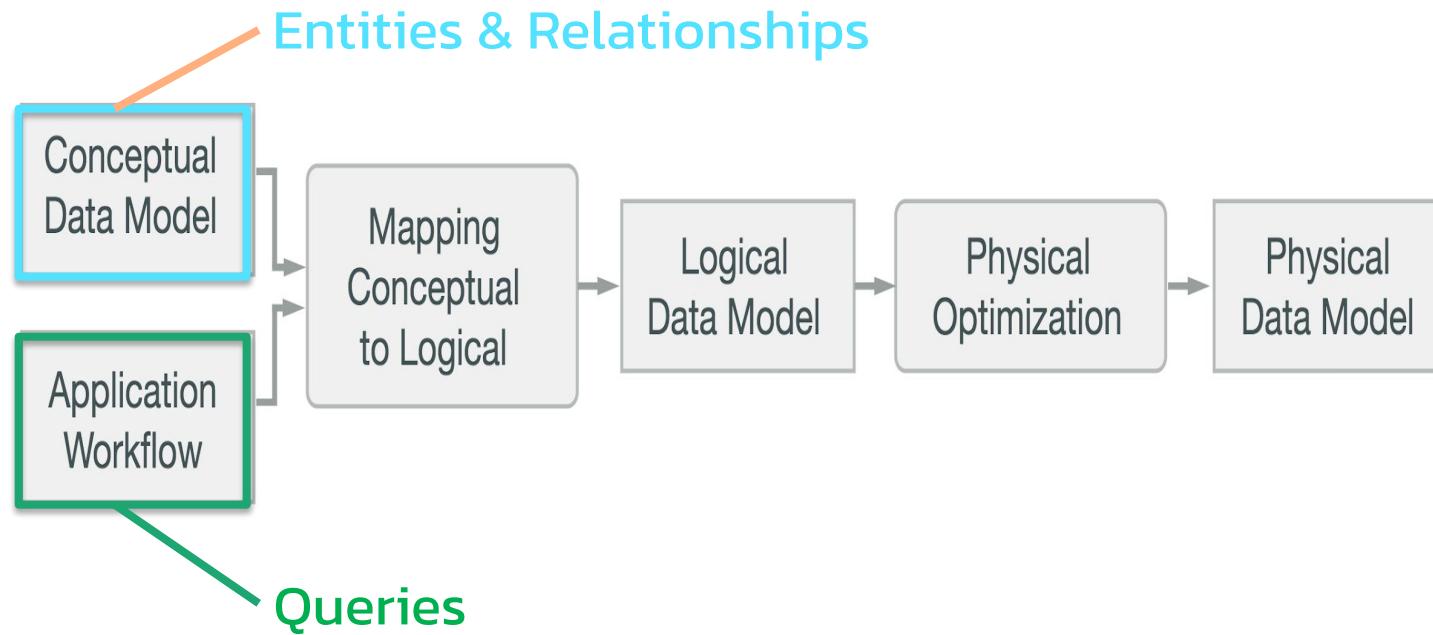
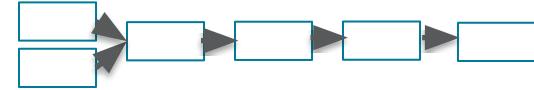


| Employees | | | |
|-----------|-----------|----------|-------------|
| userId | firstName | lastName | department |
| 1 | Edgar | Codd | Engineering |
| 2 | Raymond | Boyce | Math |
| 3 | Sage | Lahja | Math |
| 4 | Juniper | Jones | Botany |



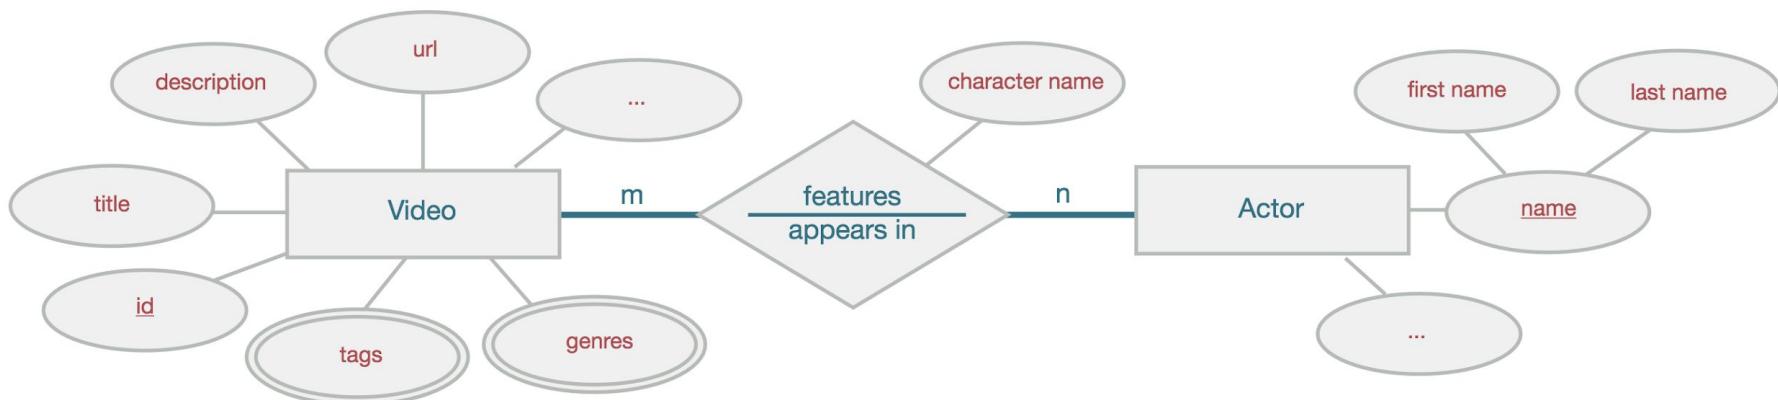
DataStax Developers

Designing Process: Step by Step

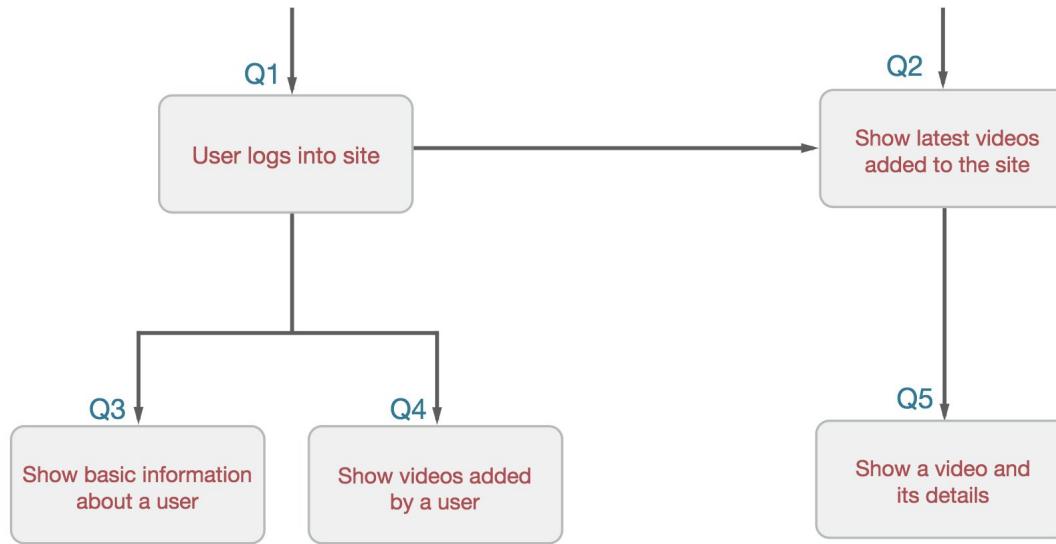


Conceptual Data Model

- Abstract view of the domain
- Technology independent
- Not specific to any database system



Application Workflow



ACCESS PATTERNS

Q1: Find a user with a **specified email**

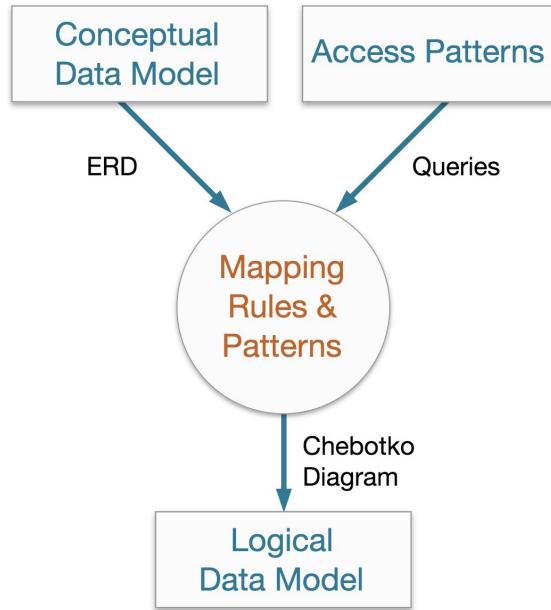
Q2: Find most recently uploaded **videos**

Q3: Find a **user with a specified id**

Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)

Q5: Find a video with a **specified video id**

Mapping Conceptual to logical

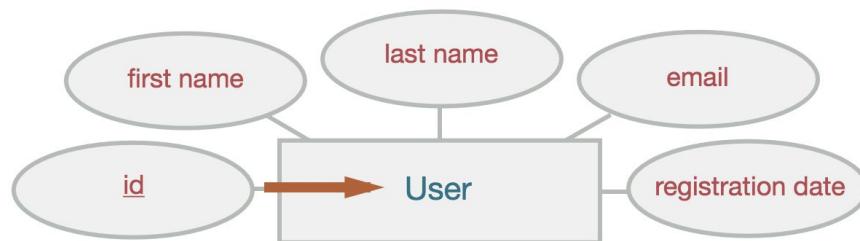


- Mapping Rule 1: Entities and relationships
- Mapping Rule 2: Equality search attributes
- Mapping Rule 3: Inequality search attributes
- Mapping Rule 4: Ordering attributes
- Mapping Rule 5: Key attributes

MR1 : Entities and Relationships

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



Logical Model

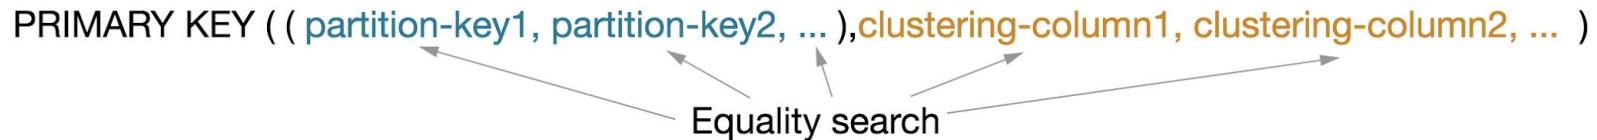
The Logical Model diagram shows a table structure. The top row is labeled "users" in blue. Below it, the columns are listed in orange: "user_id", "email", "first_name", "last_name", and "registration_date". To the right of the table, the letter "K" is written in orange. A brown arrow points from the "users" label in the Conceptual Model to the "users" label in the Logical Model.

| |
|-------------------|
| users |
| user_id |
| email |
| first_name |
| last_name |
| registration_date |

MR2: Equality Search Attributes

Attributes to query on must be in the front of the primary key

- Primary key is an ordered set of columns, made up of partition key and clustering columns
- A partition key is formed by one or more of the leading primary key columns
- Supported queries must include all partition key columns in the query

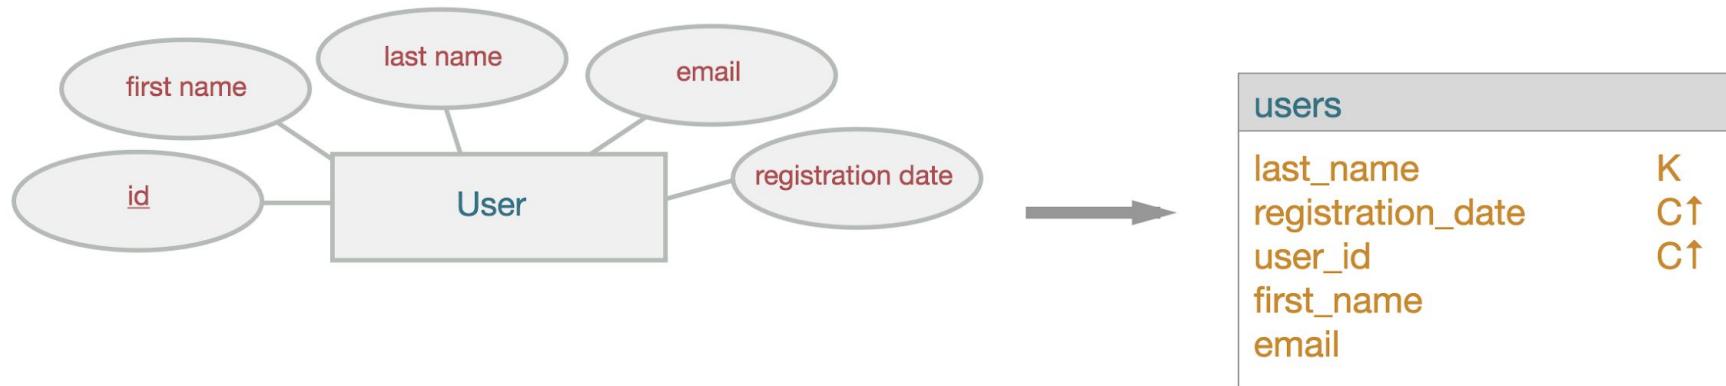


- ⚠️ Violating this rule means that you may not be able to query on a specific column

MR3: Inequality Search Attributes

Entity type example

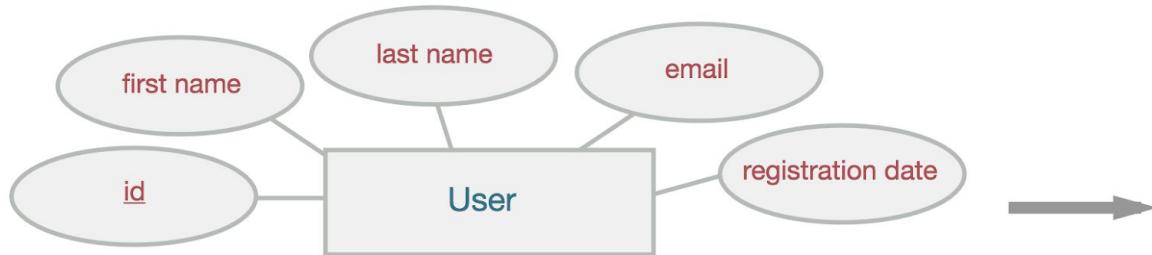
- Inequality search attributes become clustering columns
- Querying on: `last_name = ? and registration_date > ?`



MR4: Ordering Attributes

Entity type example

- Ordering attributes become clustering columns
- Querying on: *last_name = ? and registration_date > ?*
- Ordering attributes: *registration_date* (ASC)

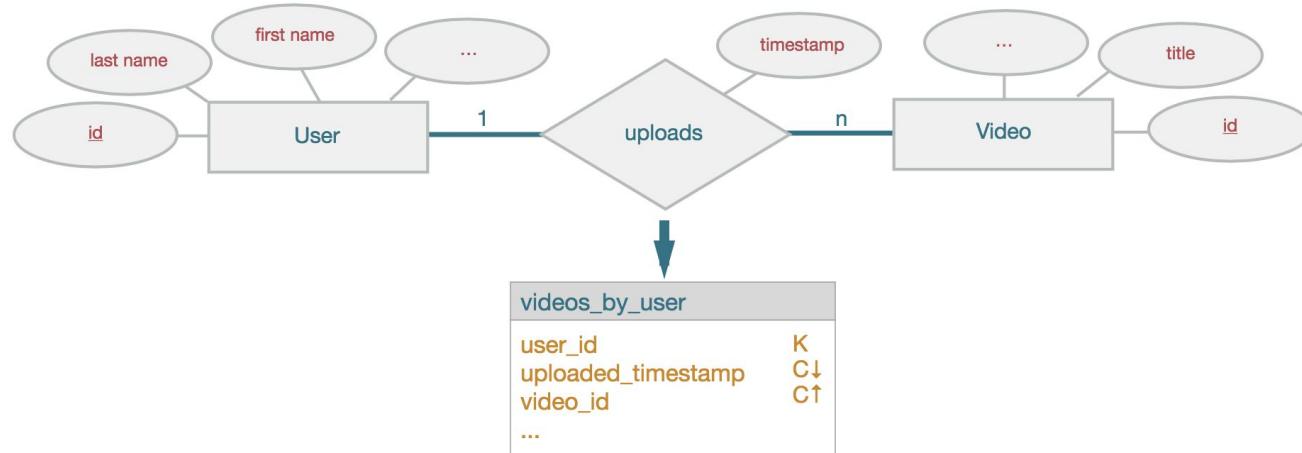


| users_registration_date | | |
|-------------------------|----|--|
| last_name | K | |
| registration_date | C↑ | |
| user_id | C↑ | |
| first_name | | |
| email | | |

MR4: Ordering Attributes

Relationship type example

- Ordering attributes become clustering columns
- Querying on: $user_id = ?$ and $uploaded_timestamp > ?$
- Ordering attributes: $uploaded_timestamp (DESC)$



MR5: Key Attributes

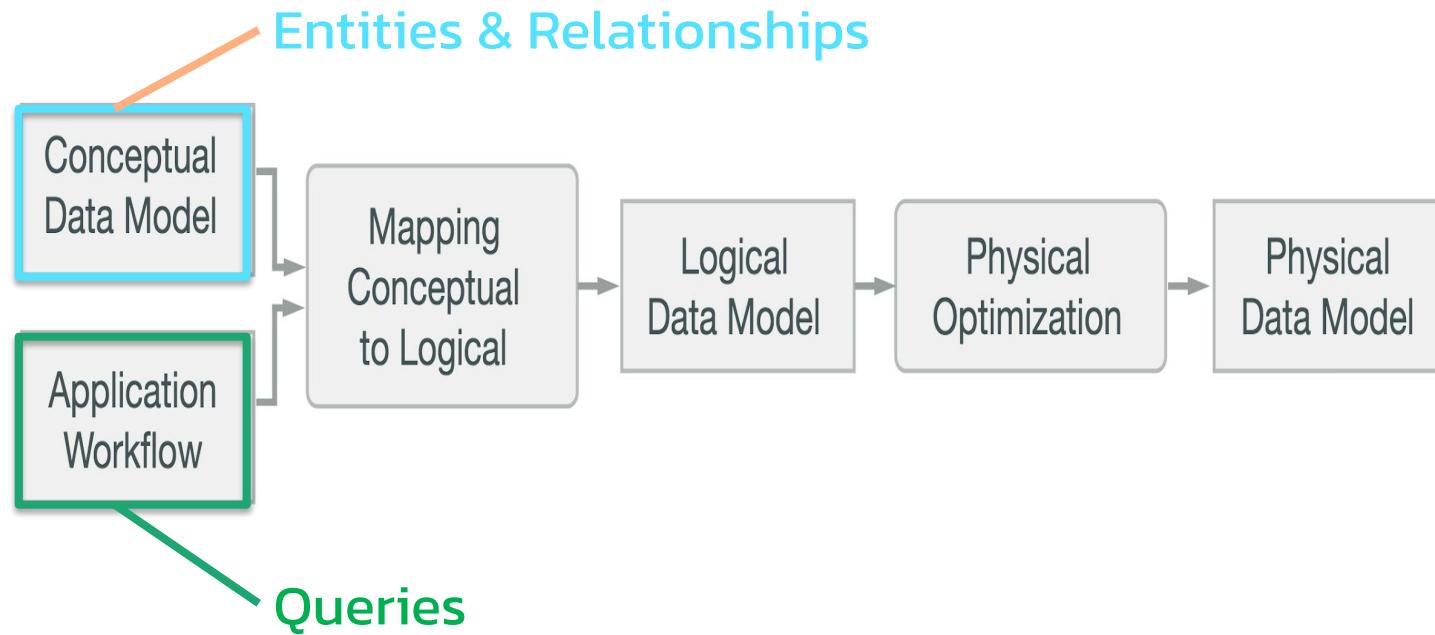
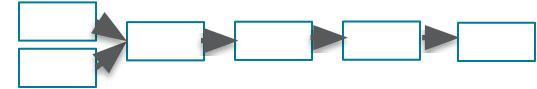
Key attributes map to primary key columns

- Primary key must include the columns that represent key attributes
- Position and order of such columns may vary
- Primary key may have additional columns to support specific queries

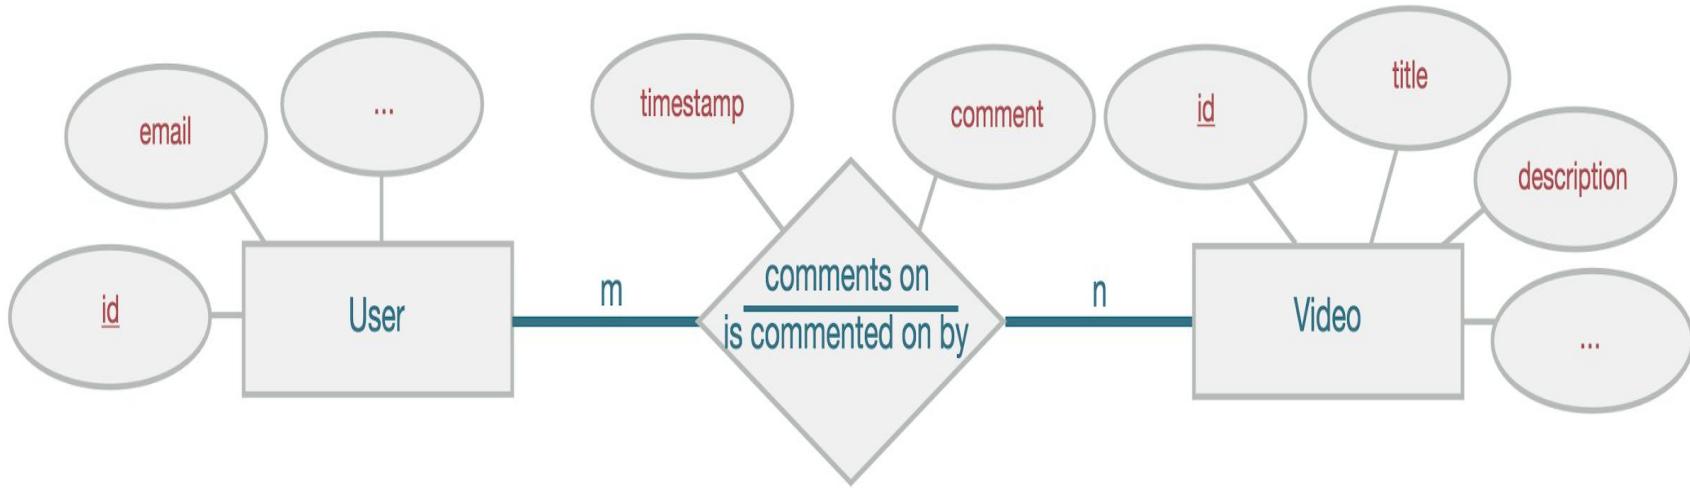
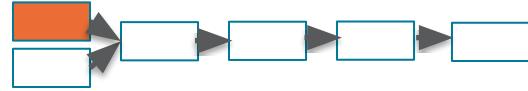


Violating this rule may result in upsert operations—loss of data

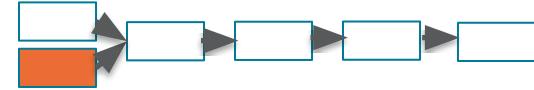
Designing Process: Step by Step



Designing Process: Conceptual Data Model



Designing Process: Application Workflow



Use-Case I:

- A User opens a Profile

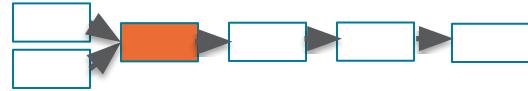
WF2: Find **comments** related to target **user** using its identifier, get most recent first

Use-Case II:

- A User opens a Video Page

WF1: Find **comments** related to target **video** using its identifier, most recent first

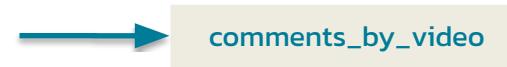
Designing Process: Mapping



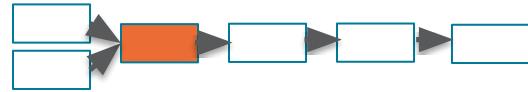
Query I: Find comments posted for a user with a known id (show most recent first)



Query II: Find comments for a video with a known id (show most recent first)



Designing Process: Mapping



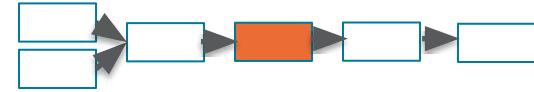
```
SELECT * FROM comments_by_user  
WHERE userid = <some UUID>
```



```
SELECT * FROM comments_by_video  
WHERE videoid = <some UUID>
```



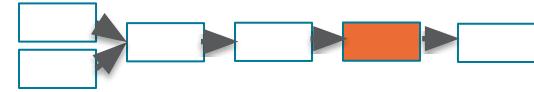
Designing Process: Logical Data Model



| comments_by_user | |
|------------------|-----|
| userid | K |
| creationdate | C ↓ |
| commentid | C ↑ |
| videoid | |
| comment | |

| comments_by_video | |
|-------------------|-----|
| videoid | K |
| creationdate | C ↓ |
| commentid | C ↑ |
| userid | |
| comment | |

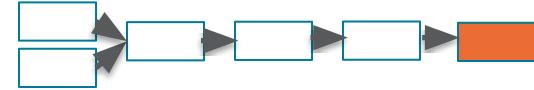
Designing Process: Physical Data Model



| comments_by_user | | |
|------------------|----------|-----|
| userid | UUID | K |
| commentid | TIMEUUID | C ↓ |
| videoid | UUID | |
| comment | TEXT | |

| comments_by_video | | |
|-------------------|----------|-----|
| videoid | UUID | K |
| commentid | TIMEUUID | C ↓ |
| userid | UUID | |
| comment | TEXT | |

Designing Process: Schema DDL



```
CREATE TABLE IF NOT EXISTS comments_by_user (
    userid uuid,
    commentid timeuuid,
    videoid uuid,
    comment text,
    PRIMARY KEY ((userid), commentid)
) WITH CLUSTERING ORDER BY (commentid DESC);
```

```
CREATE TABLE IF NOT EXISTS comments_by_video (
    videoid uuid,
    commentid timeuuid,
    userid uuid,
    comment text,
    PRIMARY KEY ((videoid), commentid)
) WITH CLUSTERING ORDER BY (commentid DESC);
```

Hands-on

<https://www.datastax.com/dev/modeling>

Try It Out: Cassandra Data Modeling

Learn how to create basic Cassandra data models

So, you want to create a Cassandra schema? Cassandra's schema development methodology is different from the relational world's approach.

In this scenario, we'll learn how to create a Cassandra schema that deals with:

- How to create tables
- How to handle table joins
- How to handle queries on non-primary key columns

If you are coming from a relational world, you create a schema by thinking about your data, creating a normalized model and then figuring out how to use the model in your app. Cassandra reverses this process by having you focus on queries within the app and using those queries to drive table design. We'll show you how!

CONTRIBUTORS



START SCENARIO

TIME TO COMPLETE
10-15 minutes

DIFFICULTY
Beginner

SHARE



<https://www.datastax.com/learn/data-modeling-by-example>



Time Series Data Modeling

Learn how to create a data model for time series data

START MODELING

Next Step: Academy.datastax.com



*Designed for professionals that
use Apache Cassandra clusters
to manage data*

**application developers
data architects
database designers
database administrators**

Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

05



K8ssandra
Tour d'horizon



DataStax

Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. **Connectivity**
2. Build your queries
3. Execute Statements
4. Object Mapping
5. Asynchronous APIs
6. Reactive APIs



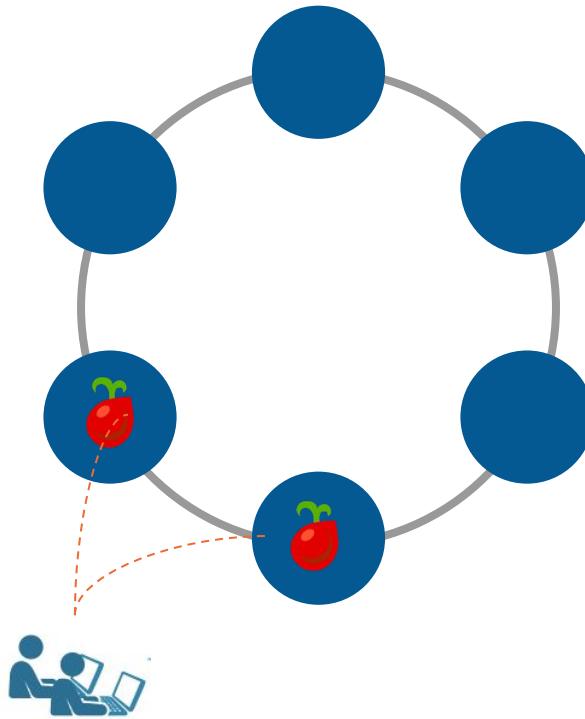
DataStax Drivers

- OSS Driver Features
 - CQL Support
 - Sync / Async API
 - Load Balancing Policies
 - Retry Policies
 - Reconnection Policies
 - Connection Pooling
 - SSL
 - Compression
 - Query Builder
 - Object Mapper
- Enterprise Driver Features
 - OSS Driver features, plus...
 - DSE Advanced Security, Unified Authentication
 - DSE Graph Fluent API
 - DSE Geometric Types



Contact Points

- Only one necessary
- Unless that node is down
- More are good



File-based Configuration

- Based on Typesafe Config
- Attributes are grouped into basic and advanced categories
- A reference file (`reference.conf`) provide default values embedded in the jar file. Can be override with key in `application.conf`.
- Driver searches `application.conf` in the classpath

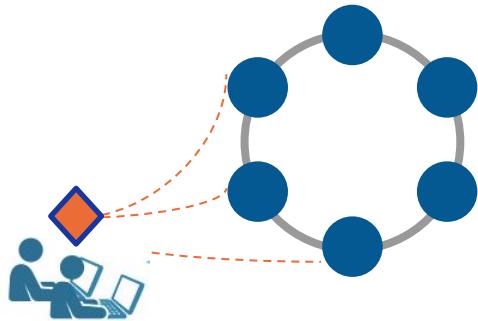


`application.conf`

```
datastax-java-driver {  
    basic {  
        request.timeout      = 5 seconds  
        request.consistency = LOCAL_QUORUM  
    }  
}
```

Load-Balancing

- Used to create query plans for each statement executed
- Default policy is token-aware, round robin
- Requests are routed to nodes in the “local” data center only



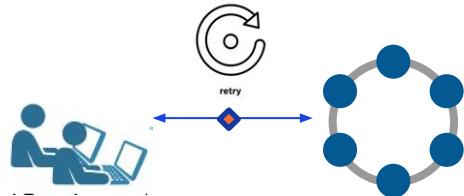
application.conf

```
datastax-java-driver {  
    basic {  
        load-balancing-policy {  
  
            # The class of the policy.  
            class = DefaultLoadBalancingPolicy  
  
            # The datacenter that is considered "local"  
            # The default policy will only include nodes from  
            # this datacenter in its query plans.  
            local-datacenter = datacenter1  
  
            # A custom filter to include/exclude nodes  
            // filter.class=  
        }  
    }  
}
```

Retry Policy

Determines when queries are retried on failure

- **DefaultRetryPolicy**
 - Default
 - Retries once onReadTimeout or onWriteTimeout
 - Enough replicas for your consistency level must be online
 - Only retries idempotent mutations



application.conf

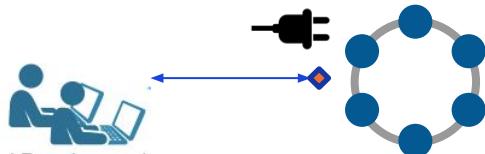
```
datastax-java-driver {  
    # The policy that controls if the driver retries  
    # requests that have failed on one node.  
    advanced.retry-policy {  
  
        # The class of the policy  
        class = DefaultRetryPolicy  
  
    }  
}
```

Reconnection Policy

Reconnects driver to a downed node

Two options:

- **ConstantReconnectionPolicy**
 - Check every N milliseconds
- **ExponentialReconnectionPolicy**
 - Increases every interval
 - Caps out at a max



```
datastax-java-driver {
```

```
    # Whether to schedule reconnection attempts  
    # if all contact points are unreachable at init  
    advanced.reconnect-on-init = false
```

```
    advanced.reconnection-policy {
```

```
        # The class of the policy  
        class = ExponentialReconnectionPolicy
```

```
        # Parameters  
        base-delay = 1 second  
        max-delay = 60 seconds
```

```
}
```

Important to know about CqlSession

- **CqlSession** is a stateful object handling communications with each node
- **CqlSession** should be unique in the Application (*Singleton*)
- **CqlSession** should be closed at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

```
@PreDestroy  
public void cleanup() {  
    if (null != cqlSession) {  
        cqlSession.close();  
    }  
}
```



DataStax

Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. Connectivity
2. **Build your queries**
3. Execute Statements
4. Object Mapping
5. Asynchronous APIs
6. Reactive APIs



How to execute queries ?

- First job of **CqlSession** is to execute queries using, well, execute method.

```
cqlSession.execute("SELECT * FROM killrvideo.users");
```

Statement

SimpleStatement

```
Statement statement = ...  
  
// (1) Explicit SimpleStatement Definition  
SimpleStatement.newInstance("select * from t1 where c1 = 5");  
  
// (2) Externalize Parameters (no name)  
SimpleStatement.builder("select * from t1 where c1 = ?")  
    .addPositionalValue(5);  
  
// (3) Externalize Parameters (name)  
SimpleStatement.builder("select * from t1 where c1 = :myVal")  
    .addNamedValue("myVal", 5);  
  
cqlSession.execute(statement);
```

Prepared and Bound Statements

- Compiled once on each node automatically as needed
- Prepare each statement only once per application
- Use one of the many bind variations to create a BoundStatement

```
PreparedStatement ps = cqlSession.prepare("SELECT * from t1 where c1 = ?");  
  
BoundStatement bound = ps.bind(5);  
  
cqlSession.execute(bound);
```

Query Builder



Query Builder

- Fluent API for building CQL string queries programmatically
- Contains methods to build SELECT, UPDATE, INSERT and DELETE statements
- Generates a Statement as per the earlier techniques

OSS Driver (current version 4.2.0)

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>
    java-driver-query-builder
  </artifactId>
</dependency>
```

DSE Driver (current version 2.2.0)

```
<><dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>
    dse-java-driver-query-builder
  </artifactId>
</dependency>
```



QueryBuilder

```
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.bindMarker;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.deleteFrom;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.selectFrom;
import static com.datastax.oss.driver.api.querybuilder.relation.Relation.column;

// Simple SELECT using QueryBuilder
Statement stmtSelect = selectFrom("killrvideo", "videos_by_users")
    .column("userid").column("commentid")
    .function("toTimestamp", Selector.column("commentid")).as("comment_timestamp")
    .where(column("userid").isEqualTo(bindMarker("userid"))))
    .build()

// Simple DELETE using QueryBuilder
Statement stmtDelete = deleteFrom("killrvideo", "videos_by_users")
    .where(column("userid").isEqualTo(bindMarker("userid"))))
    .build()
```

Query Builder

- Can also use **QueryBuilder** to create **PreparedStatements** and later execute at runtime
- Note use of **bindMarker()** to designate parameters that will be provided later

```
// Prepared QueryBuilder statements as any statement
PreparedStatement psStmt = cqlSession.prepare(
    deleteFrom("killrvideo", "videos_by_users")
        .where(column("userid").isEqualTo(bindMarker("userid"))))
        .build()));

// Binding
BoundStatement bsStmt = psStmt.bind("e7a8ac9f-c12d-415c-a526-4137815df573");

// Execute
cqlSession.execute(bsStmt);
```



DataStax

Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. Connectivity
2. Build your queries
3. **Execute Statements**
4. Object Mapping
5. Asynchronous APIs
6. Reactive APIs



ResultSet

- **ResultSet** is the object returned for executing query. It contains **ROWS** (data) and **EXECUTION INFO**.
- **ResultSet** is **iterable** and as such you can navigate from row to row.
- Results are **always paged** for you (avoiding memory and response time issues)

```
ResultSet rs = cqlSession.execute(myStatement);

// Plumbery
ExecutionInfo info = rs.getExecutionInfo();
int executionTime = info.getQueryTrace().getDurationMicros();

// Data: NOT ALL DATA RETRIEVED IMMEDIATELY (only when needed .next())
Iterator<Row> iterRow = rs.iterator();
int itemsFirstCall = rs.getAvailableWithoutFetching();
```

Parsing ResultSet

```
// We know there is a single row (eg: count)
Row singleRow = resultSet.one();

// We know there are not so many results we can get all (fetch all pages)
List<Row> allRows = resultSet.all();

// Browse iterable
for(Row myRow : resultSet.iterator()) {
    // .. Parsing rows
}

// Use Lambda
rs.forEach(row -> { row.getColumnDefinitions(); });

// Use for LWT
boolean isQueryExecuted = rs.wasApplied();
```

Parsing Rows

```
// Sample row
Row row = resultSet.one();

// Check null before read
Boolean isUserNameNull = row.isNull("username");

// Reading Values from row
String userName1 = row.get("username", String.class);
String userName2 = row.getString("username");
String userName3 = row.getString(CqlIdentifier.fromCql("username"));

// Tons of types available
row.getUuid("userid");
row.getBoolean("register");
row.getCqlDuration("elapsed");
...
```

Paging

- ResultSet contains up to “**pageSize**” items. When browsing records you may hit this number that will trigger fetching next “pageSize” items.
- To fetch anything else than first page you must provide a **PagingState**.

```
// Enforce few items per page (often = UI requirements)
myStatement = myStatement.setPageSize(10);
ResultSet page1 = cqlSession.execute(myStatement);

// Paging State
ByteBuffer pagingState = page1.getExecutionInfo().getPagingState();
myStatement = myStatement.setPageState(pagingState);

// Very same statement with pagingState provided
ResultSet page2 = cqlSession.execute(myStatement);
```

Batches – What you need to know

- Batches about **data integrity** between tables
- Not Atomic & not used for mass query optimization
- Used to keep denormalized data in sync
- There is no guarantee that a batch will complete all operations.
- There are still edge cases where things can fail out.
- There is no rollback if something fails
- This is where upserts come into play as you can simply re-fire the batch.

Batch Example

```
// Sample statements (insert same data in multiple tables)
Statement stmt1 = SimpleStatement
    .builder("INSERT INTO users_by_group(groupid,userid) values(?,?)")
    .addPositionalValue(groupname, username);
Statement stmt2 = SimpleStatement
    .builder("INSERT INTO groups_by_user(userid,groupid) values(?,?)")
    .addPositionalValue(username, groupname);

// Group as a Batch
BatchStatement batchStmt = BatchStatement
    .builder(DefaultBatchType.LOGGED)
    .addStatement(stmt1).addStatement(stmt2).build();

// Execute
cqlSession.execute(batchStmt);
```

Profiles

Override parameters for dedicated request

```
SimpleStatement  
.newInstance("select...")  
.setPageSize(10)  
// here is the magic ☐  
.setExecutionProfileName("dse_search");
```



application.conf

```
datastax-java-driver {  
    profiles{  
  
        # DSE Search type of queries  
        dse_search {  
            basic {  
                request.consistency = LOCAL_ONE  
                request.timeout      = 5 seconds  
            }  
        } # The class of the policy  
        fast_query {  
            basic.request.consistency = ONE  
            basic.request.timeout     = 1 second  
        }  
    }  
}
```



DataStax

Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. Connectivity
2. Build your queries
3. Execute Statements
4. **Object Mapping**
5. Asynchronous APIs
6. Reactive APIs



Object Mapper

- WHAT ?
 - Abstracts details of mapping Java attributes to/from CQL types and UDTs
 - Packaged separately from the driver – pom.xml update required
 - This slide shows the runtime dependency – will show compile-time shortly
- HOW ?
 - Some annotation processors will GENERATE Mapper, Dao, and Entity implementations for you
 - At each update in the files, the IDE (eclipse, intelliJ) will use annotation processor
 - Compiler plugin must be updated to define the annotation processor

Object Mapper

OSS Driver

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.oss</groupId>
      <artifactId>java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
    </configuration>
</plugin>
```

DSE Driver

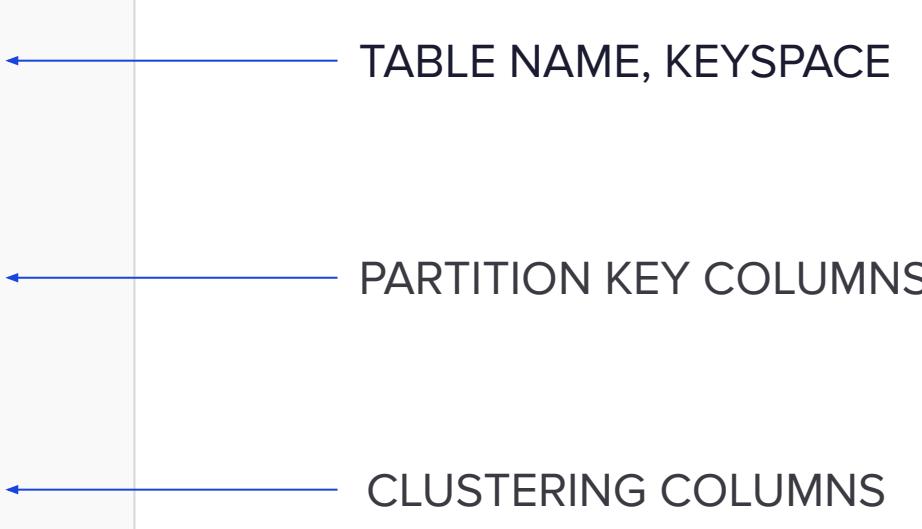
```
<dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>dse-java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.dse</groupId>
      <artifactId>dse-java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
    </configuration>
</plugin>
```

Annotate Entities

```
@Entity  
@CqlName("user_v")  
public class UserVideo {  
  
    @PartitionKey  
    @CqlName("userid")  
    private UUID userid;  
  
    @ClusteringColumn(1)  
    @CqlName("added")  
    private UUID videoid;  
  
}
```



Annotate DAO Interface (1/2)

```
@Dao
public interface VideoDao {

    @Select
    Optional<UserVideo> findUserById(UUID userid);

    @Query("SELECT * FROM ${tableId}")
    PagingIterable<UserVideo> findAll();

    @Select(customWhereClause = "videoid = : vid")
    PagingIterable<UserVideo>
        findUserByVideoId(@CqlName("videoid") UUID vid);
}
```

Annotate DAO Interface (2/2)

```
// Save a bean
@Insert
void save(UserVideo userVideo);

// Userid id is PK
@Delete
void delete(UUID userid);

// Custom implementations
@QueryProvider(
    providerClass = MySampleQueryProvider.class,
    entityHelpers = { UserVideo.class })
String doSomething(String abc);
```

Annotate Mapper Interface

```
@Mapper
public interface MyApplicationMapper {
    @DaoFactory
    VideoDao videoDao(@DaoKeyspace CqlIdentifier keyspace);
}
```

Sample QueryProvider

```
public class MySampleQueryProvider {  
  
    // Constructor, getting session  
    public MySampleQueryProvider(  
        MapperContext context,  
        EntityHelper<UserVideo> helperUser) {}  
  
    // Custom implementation method  
    public String doSomething(String abc) {  
    }  
}
```

Object Mapper

- WHAT ?
 - Abstracts details of mapping Java attributes to/from CQL types and UDTs
 - Packaged separately from the driver – pom.xml update required
 - This slide shows the runtime dependency – will show compile-time shortly
- HOW ?
 - Some annotation processors will GENERATE Mapper, Dao, and Entity implementations for you
 - At each update in the files, the IDE (eclipse, intelliJ) will use annotation processor
 - Compiler plugin must be updated to define the annotation processor

Object Mapper

OSS Driver

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.oss</groupId>
      <artifactId>java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```

DSE Driver

```
<dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>dse-java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.dse</groupId>
      <artifactId>dse-java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```

Annotate Entities

```
@Entity  
@CqlName("user_v")  
public class UserVideo {  
  
    @PartitionKey  
    @CqlName("userid")  
    private UUID userid;  
  
    @ClusteringColumn(1)  
    @CqlName("added")  
    private UUID videoid;  
  
}
```



TABLE NAME, KEYSPACE

PARTITION KEY COLUMNS

CLUSTERING COLUMNS

Annotate DAO Interface (1/2)

```
@Dao
public interface VideoDao {

    @Select
    Optional<UserVideo> findUserById(UUID userid);

    @Query("SELECT * FROM ${tableId}")
    PagingIterable<UserVideo> findAll();

    @Select(customWhereClause = "videoid = : vid")
    PagingIterable<UserVideo>
        findUserByVideoId(@CqlName("videoid") UUID vid);
}
```

Annotate DAO Interface (2/2)

```
// Save a bean
@Insert
void save(UserVideo userVideo);

// Userid id is PK
@Delete
void delete(UUID userid);

// Custom implementations
@QueryProvider(
    providerClass = MySampleQueryProvider.class,
    entityHelpers = { UserVideo.class })
String doSomething(String abc);
```

Annotate Mapper Interface

```
@Mapper
public interface MyApplicationMapper {
    @DaoFactory
    VideoDao videoDao(@DaoKeyspace CqlIdentifier keyspace);
}
```

Sample QueryProvider

```
public class MySampleQueryProvider {  
  
    // Constructor, getting session  
    public MySampleQueryProvider(  
        MapperContext context,  
        EntityHelper<UserVideo> helperUser) {}  
  
    // Custom implementation method  
    public String doSomething(String abc) {  
    }  
}
```



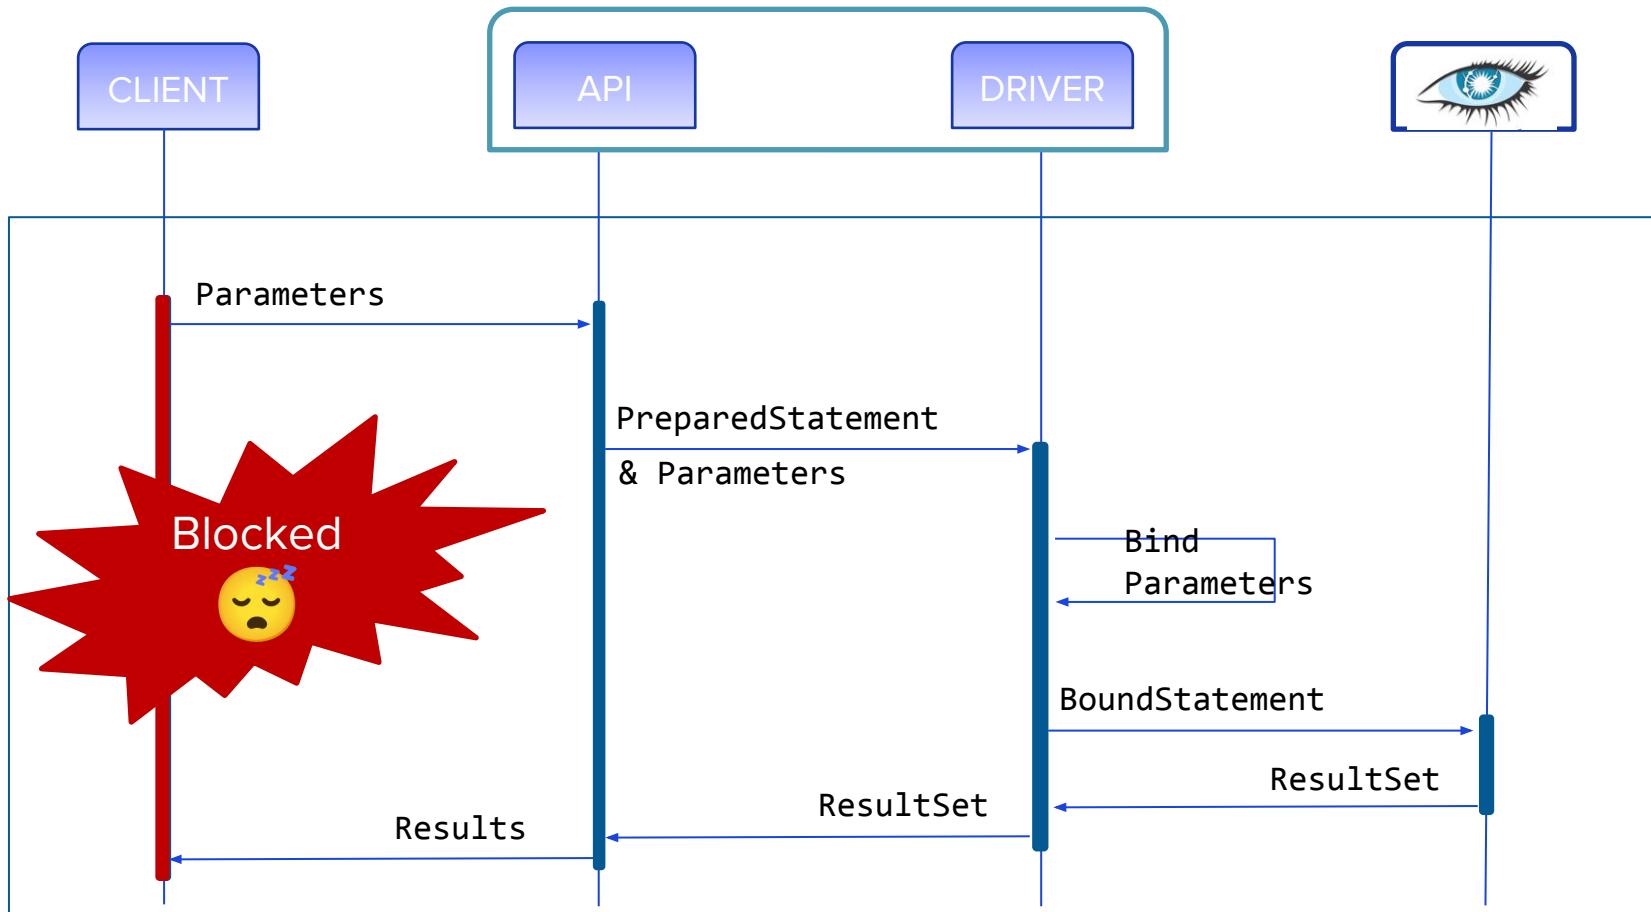
DataStax

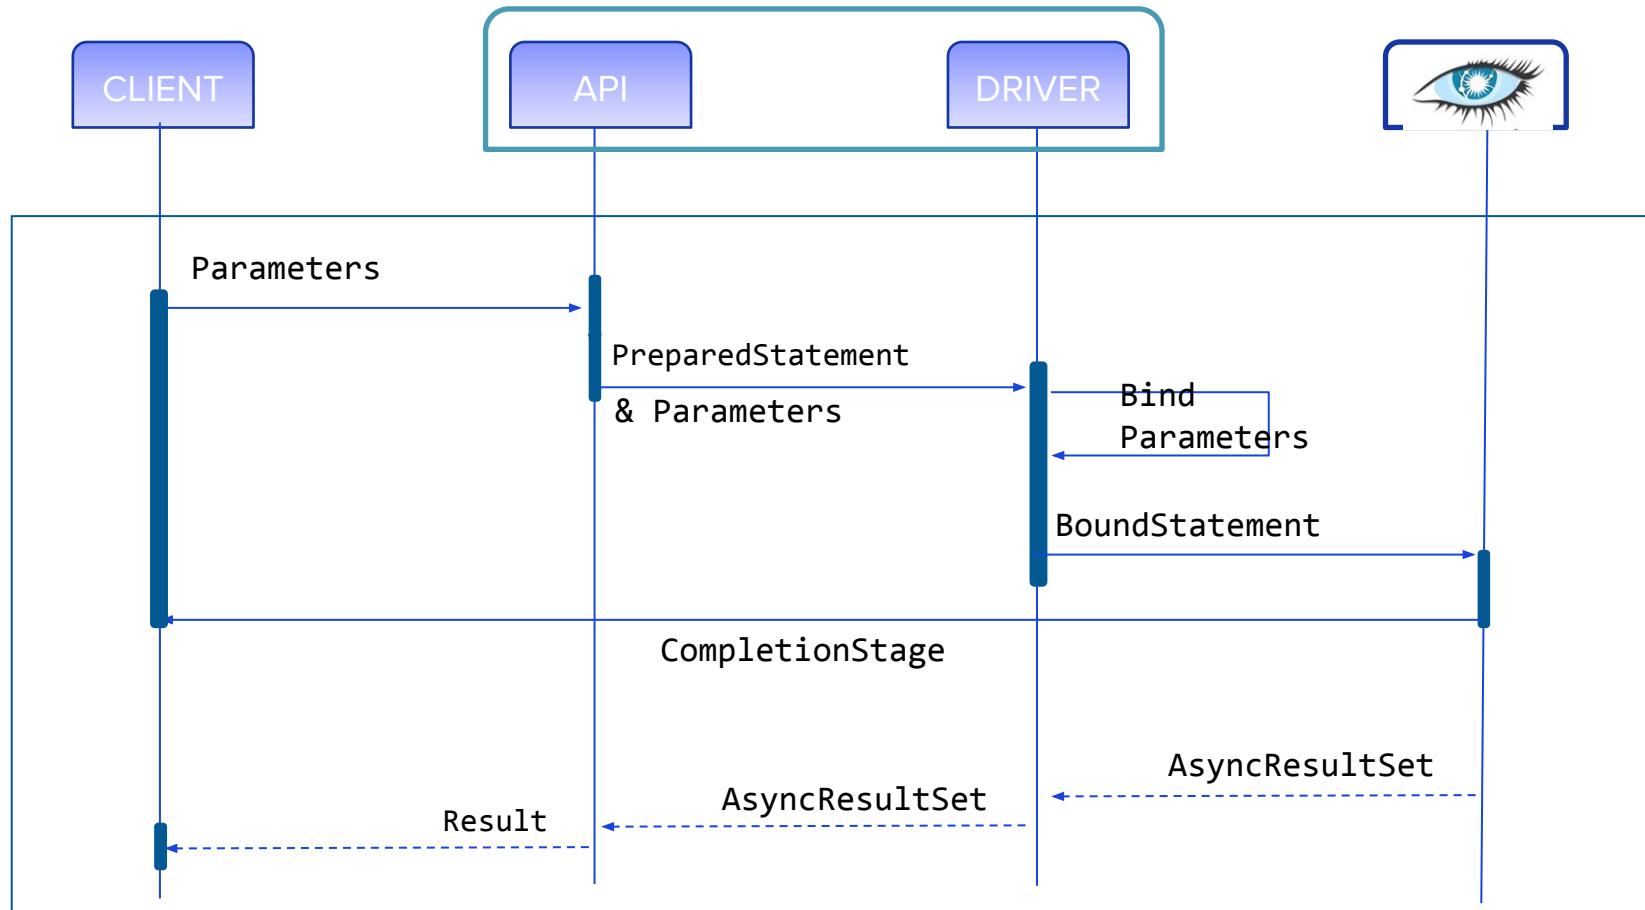
Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. Connectivity
2. Build your queries
3. Execute Statements
4. Object Mapping
5. **Asynchronous APIs**
6. Reactive APIs







Asynchronous Queries

```
// From Synchronous
ResultSet resSync = cqlSession.execute(myStatement);

// to Asynchronous
CompletionStage<AsyncResultSet> resAsync =
    cqlSession.executeAsync(myStatement);

resAsync.thenApply(AsyncPagingIterable::one)
    .thenApply(Optional::ofNullable)
    .thenApply(optional -> optional.map(rowMapper))
    .then..
```



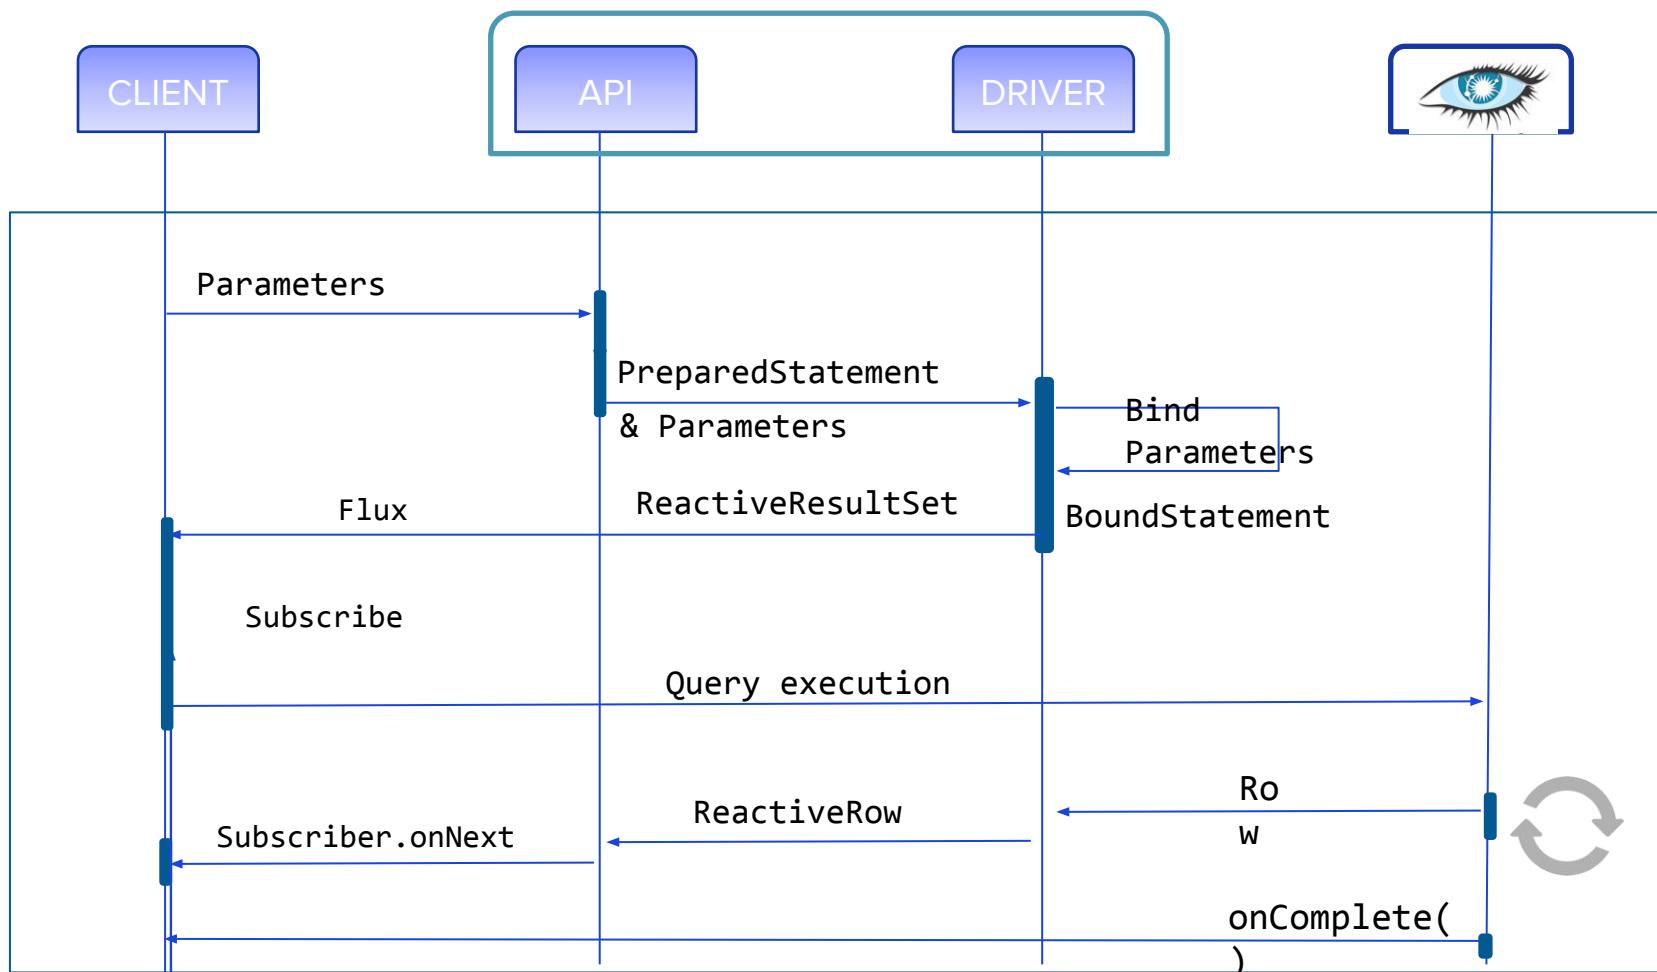
DataStax

Formation CapGemini 5 octobre

III – Apache Cassandra™ Développement

1. Connectivity
2. Build your queries
3. Execute Statements
4. Object Mapping
5. Asynchronous APIs
6. **Reactive APIs**





Reactive Queries

```
private final Function<Row, MyBean> rowMapper = ..;

// Execute in reactive way
ReactiveResultSet rs = session.executeReactive(myStatement);

// Return Flux for lists
Flux<ReactiveRow> flux = Flux.from(rs);
Flux<MyBean> res1= flux.skip(offset).take(limit).map(rowMapper);

// Return Mono for single bean
Mono<MyBean> res2 = Mono.fromDirect(rs).map(rowMapper);
```

Cassandra and Spring

<https://www.datastax.com/dev/spring>

Build a Spring Boot REST Service

This scenario is a brief introduction to Spring Boot so you can see how to create an initial project.

Do you want to set up your first Spring Boot project? Then, you're at the right place!

In this scenario, we'll learn how to:

- >Create a Spring Boot project for a simple REST server
- Create some service endpoints
- Pass parameters to an endpoint
- Create a REST endpoint that uses Cassandra

This scenario is a brief introduction to Spring Boot so you can see how to create an initial project.

ver 0.001

CONTRIBUTORS



START SCENARIO

TIME TO COMPLETE

15 minutes

DIFFICULTY

Beginner

SHARE



Hands-on

Sample Codes

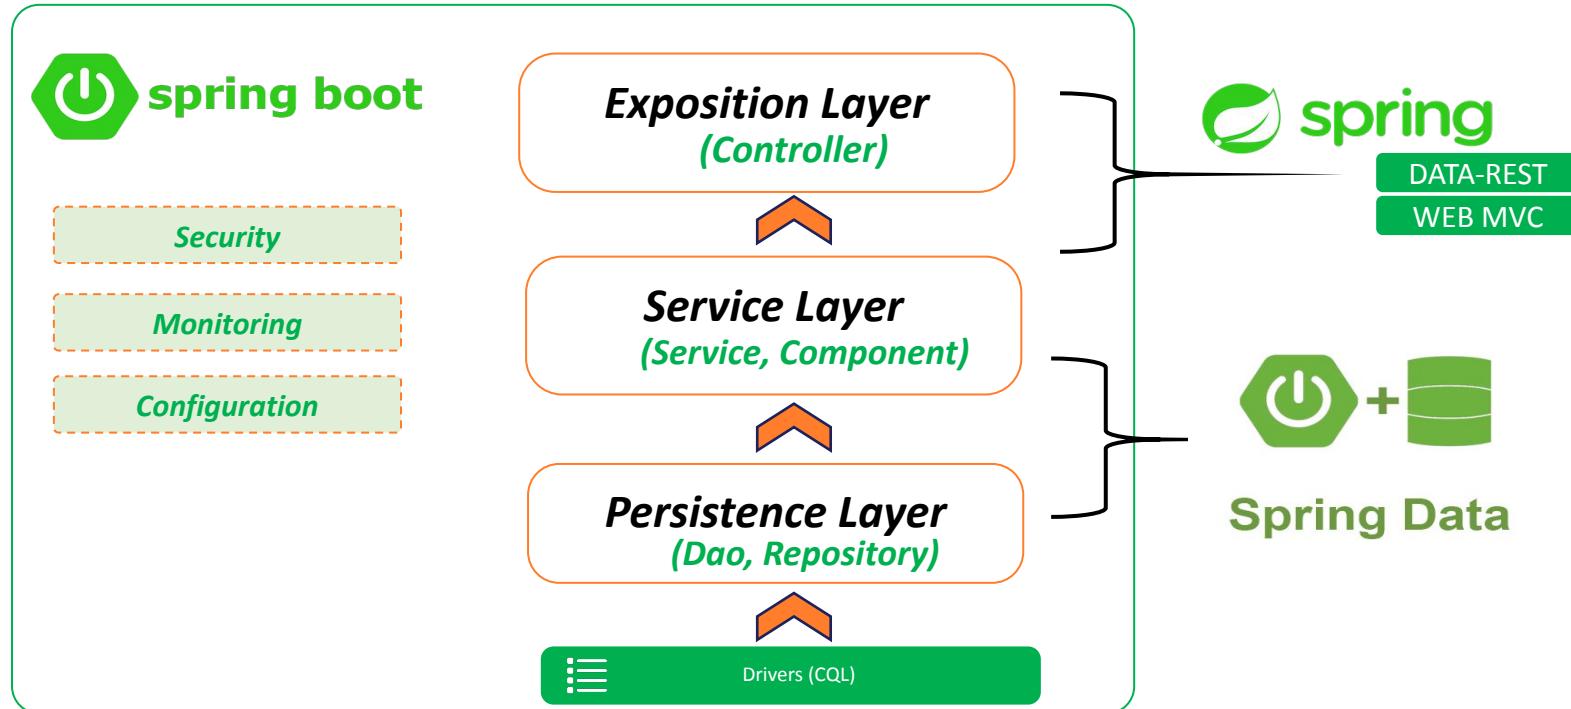
<https://github.com/DataStax-Examples/java-cassandra-driver-from3x-to4x>

<https://github.com/datastaxdevs/conference-2020-javafest-api>

<https://github.com/DataStax-Academy/microservices-java-workshop-online>

Cassandra and Spring Data

<https://github.com/datastaxdevs/workshop-spring-stargate>



Cassandra and Spring Data

Documents

[Spring Data Cassandra & DataStax.pdf](#)

[Spring Data Cassandra 3.x and Drivers 4.x](#)

Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

05

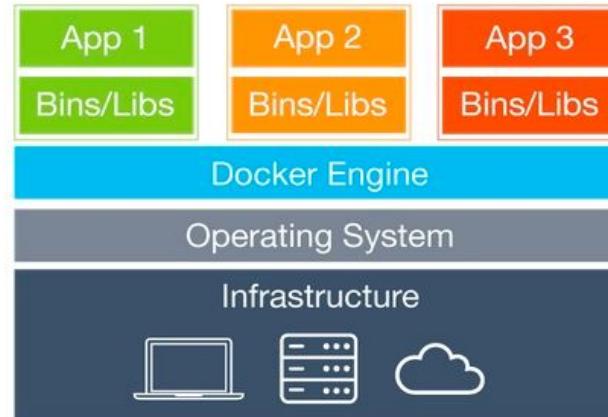


K8ssandra
Tour d'horizon

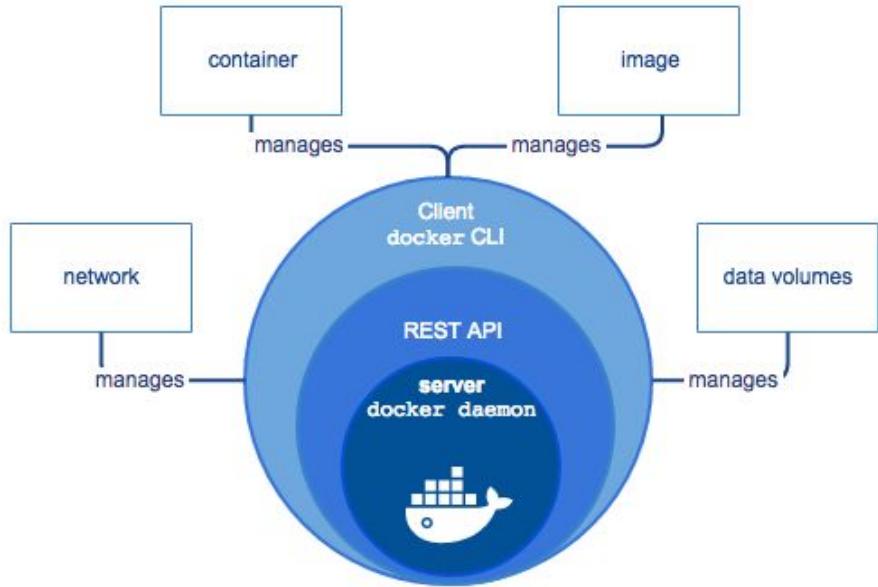
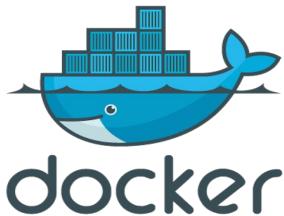
Why containers?



Virtual Machines



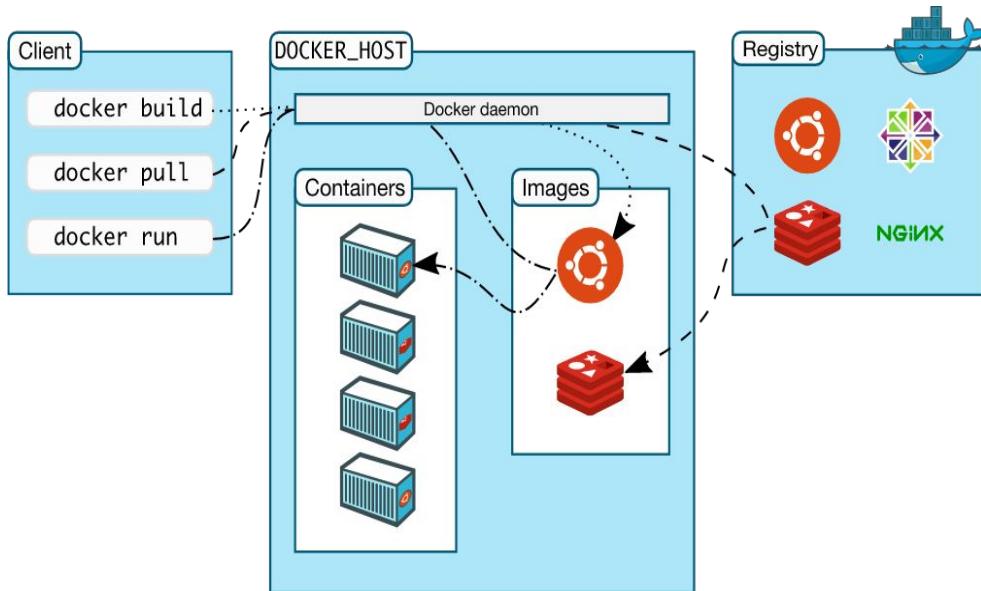
Containers



Docker Engine is Client-server application with :

- A server which is a type of long-running program called a daemon process (**the `dockerd` command**).
- A **REST API** which specifies interfaces that programs can use to talk to the daemon and instruct it what to do.¹³⁸
- A **command line interface (CLI)** client (the `docker` command).

Docker Architecture



Some Docker client commands

- We move immutable **images** (build/pull/push)
- Docker Daemon start containers based on images (run)
 - *Adding Env Var*
 - *Adding Storage (volumes)*
 - *Adding Network*
 - *Exposing port*
- We interact with **running containers** (ps, exec...)

¹³⁹



cassandra ☆

Docker Official Images

Apache Cassandra is an open-source distributed storage system.

⬇ 100M+

Container

Linux

PowerPC 64 LE

ARM

ARM 64

386

x86-64

Databases

Official Image

Linux - x86 (latest)

Copy and paste to pull this image

`docker pull cassandra`

▼

[View Available Tags](#)

Running Cassandra in Docker

- Define a proper [network](#)
- [Env variables](#) can be defined to override keys in [cassandra.yaml](#).
- Export ports [7000, 9042, ...](#)
- Define volumes to stores data
 - **/var/lib/cassandra**

```
$ docker run
--name some-cassandra -d \
-e CASSANDRA_BROADCAST_ADDRESS=10.42.42.42 \
-p 7000:7000,9042:9042
-v /my/own/datadir:/var/lib/cassandra \
cassandra:tag
```

Hands-ON

<https://github.com/datastaxdevs/conference-2021-apachecon-stargate>

Running Cassandra in Docker

Learn to deploy a Cassandra Docker app

Docker is a containerization platform that packages applications and their dependencies together so that they can be easily deployed across multiple environments.

How are containers different from virtual machines?

In this scenario we will deploy a Web application, a Microservice and a Cassandra cluster in Docker.

We will:

- Create a Docker network
- Pull the images from Docker Hub
- Deploy the Cassandra cluster
- Deploy the Microservice
- Test the Microservice to database connection
- Deploy the Web application
- Test the Web application
- Shut down the containers
- Re-create the application using Docker compose

Let's get started!

CONTRIBUTORS



START SCENARIO

TIME TO COMPLETE

15 minutes

DIFFICULTY

Intermediate

SHARE



Docker-Compose

```
docker-compose -f docker-compose.yml up -d --scale cassandra-node=2
```

Define and run multi-container Docker applications through the use of a **YAML** file to configure your applications

```
version: '2'
services:

  cassandra-seed:
    container_name: cassandra-seed-node
    image: cassandra: 3.11.6
    ports:
      - "9042:9042"    # Native transport
      - "7199:7199"    # JMX
      - "9160:9160"    # Thrift clients

  cassandra-node:
    image: cassandra: 3.11.6
    command: /bin/bash -c "echo 'Waiting for seed node' && sleep 30 && /docker-entrypoint.sh cassandra -f"
    environment:
      - "CASSANDRA_SEEDS=cassandra-seed-node"
    depends_on:
      - "cassandra-seed"
```

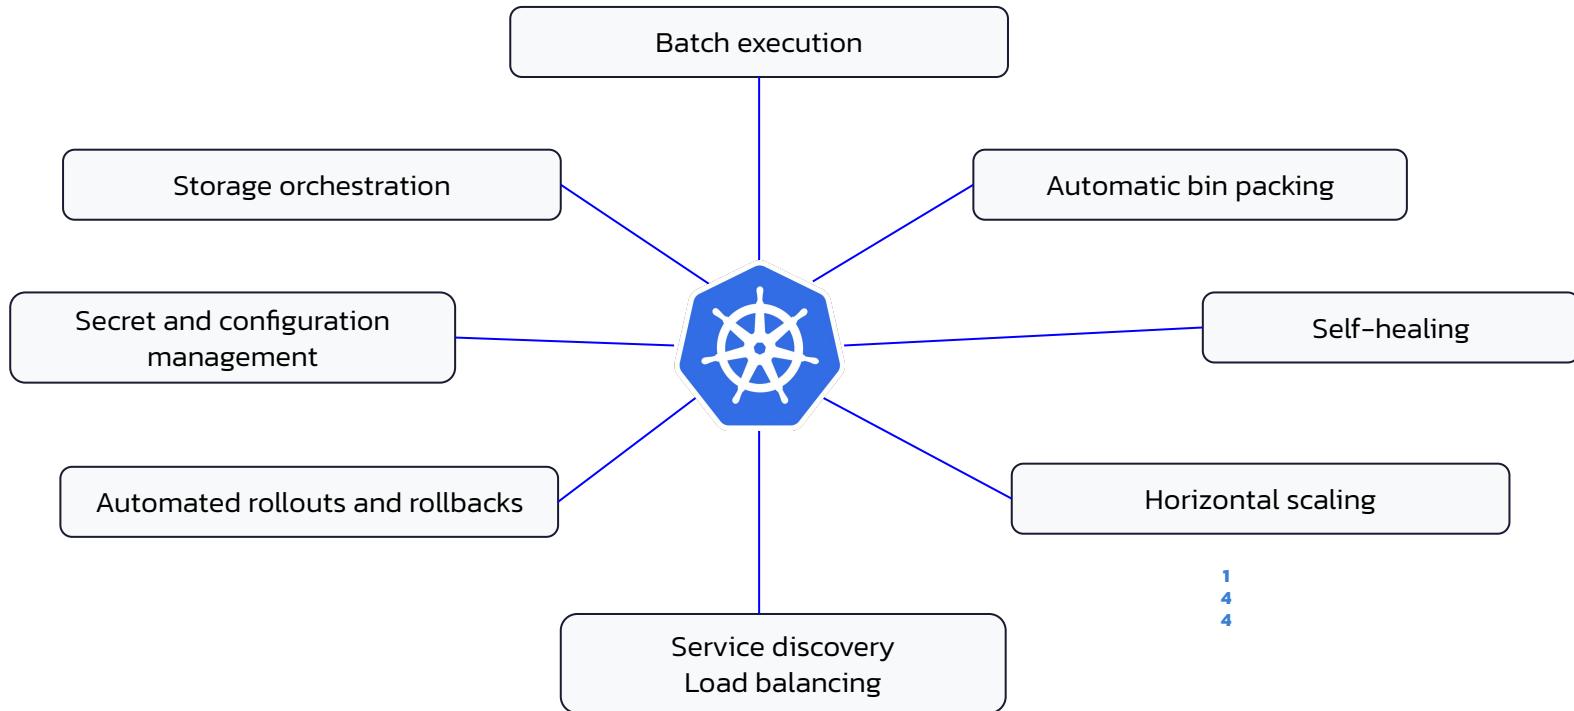


kubernetes

"Kubernetes is an open-source system for **automating** deployment, scaling, and **management** of containerized applications."

1
4
3

Why Kubernetes ?



K8s Infrastructure



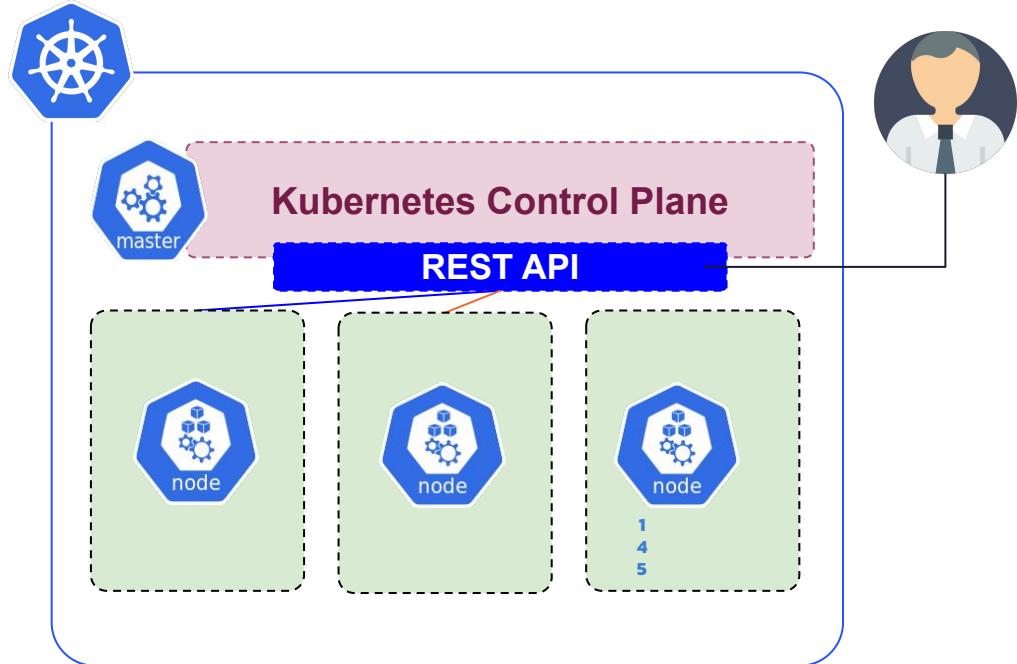
Cluster:
Kubernetes cluster.



Master:
Kubernetes Control Plane.



Node:
Worker machine



K8s Infrastructure – Control Plane



K8s API Server

Kubernetes API



Controller Manager

Kubernetes controller manager



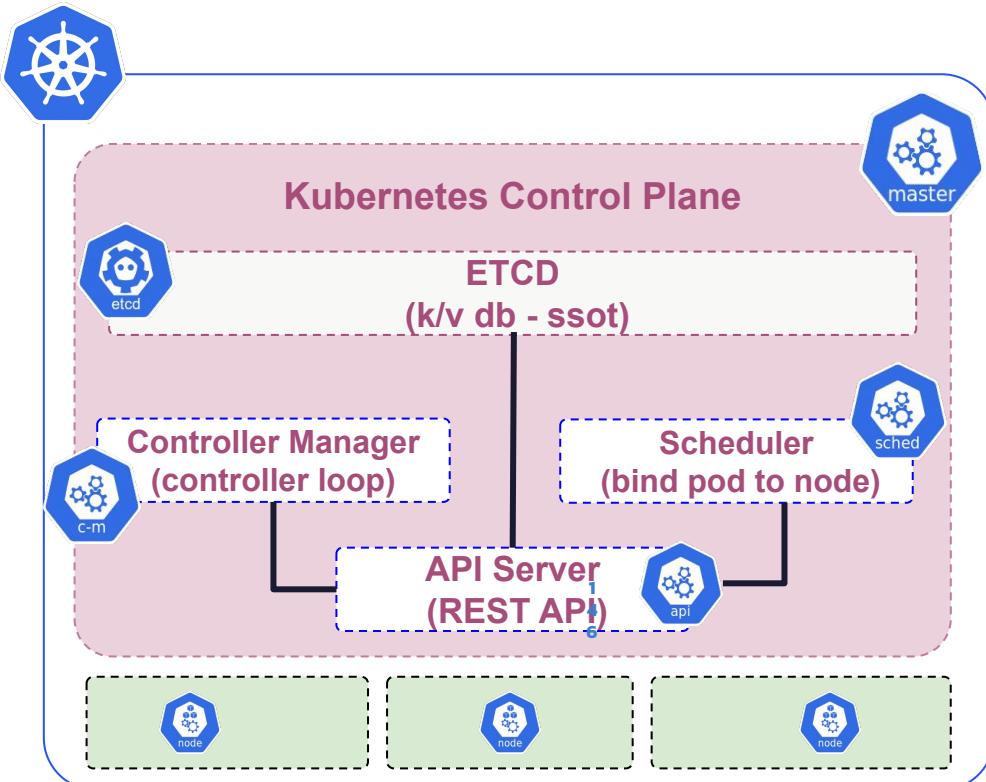
Scheduler

In charge of placing pods



ETCD

Kubernetes' backing store.



K8s Infrastructure – Worker Node



Kubelet:

The kubelet is the primary “node agent” that runs on each node.



Kube-proxy

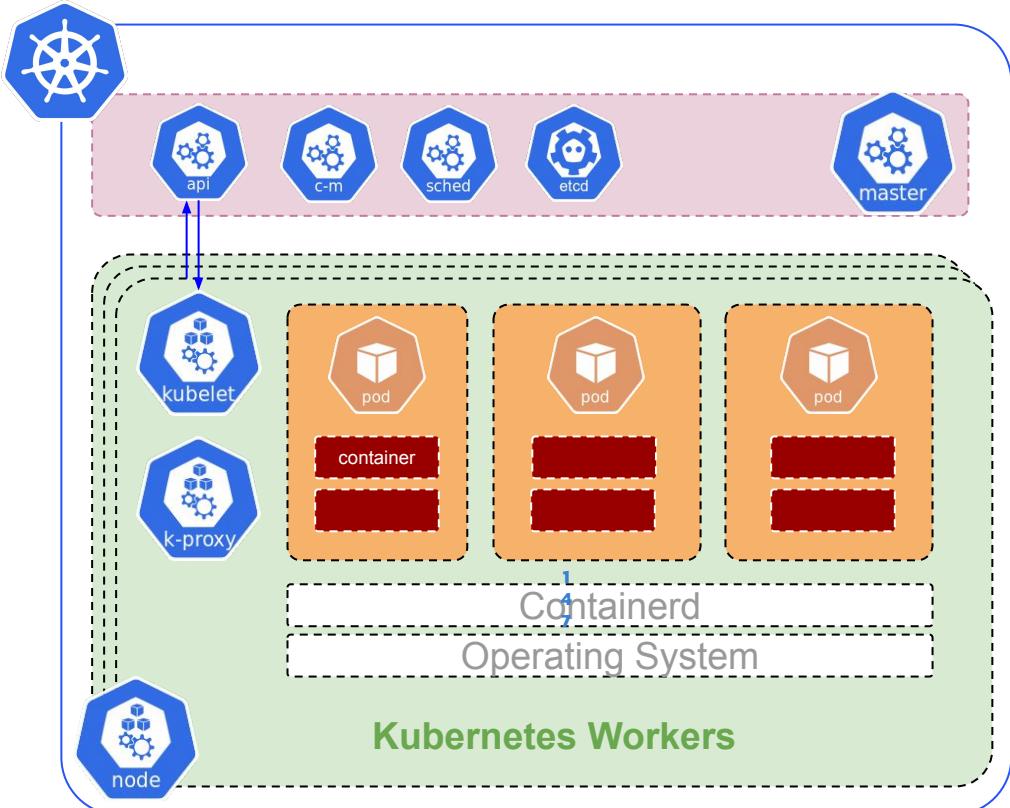
The Kubernetes network proxy runs on each node. This reflects services as defined in the Kubernetes API on each node.



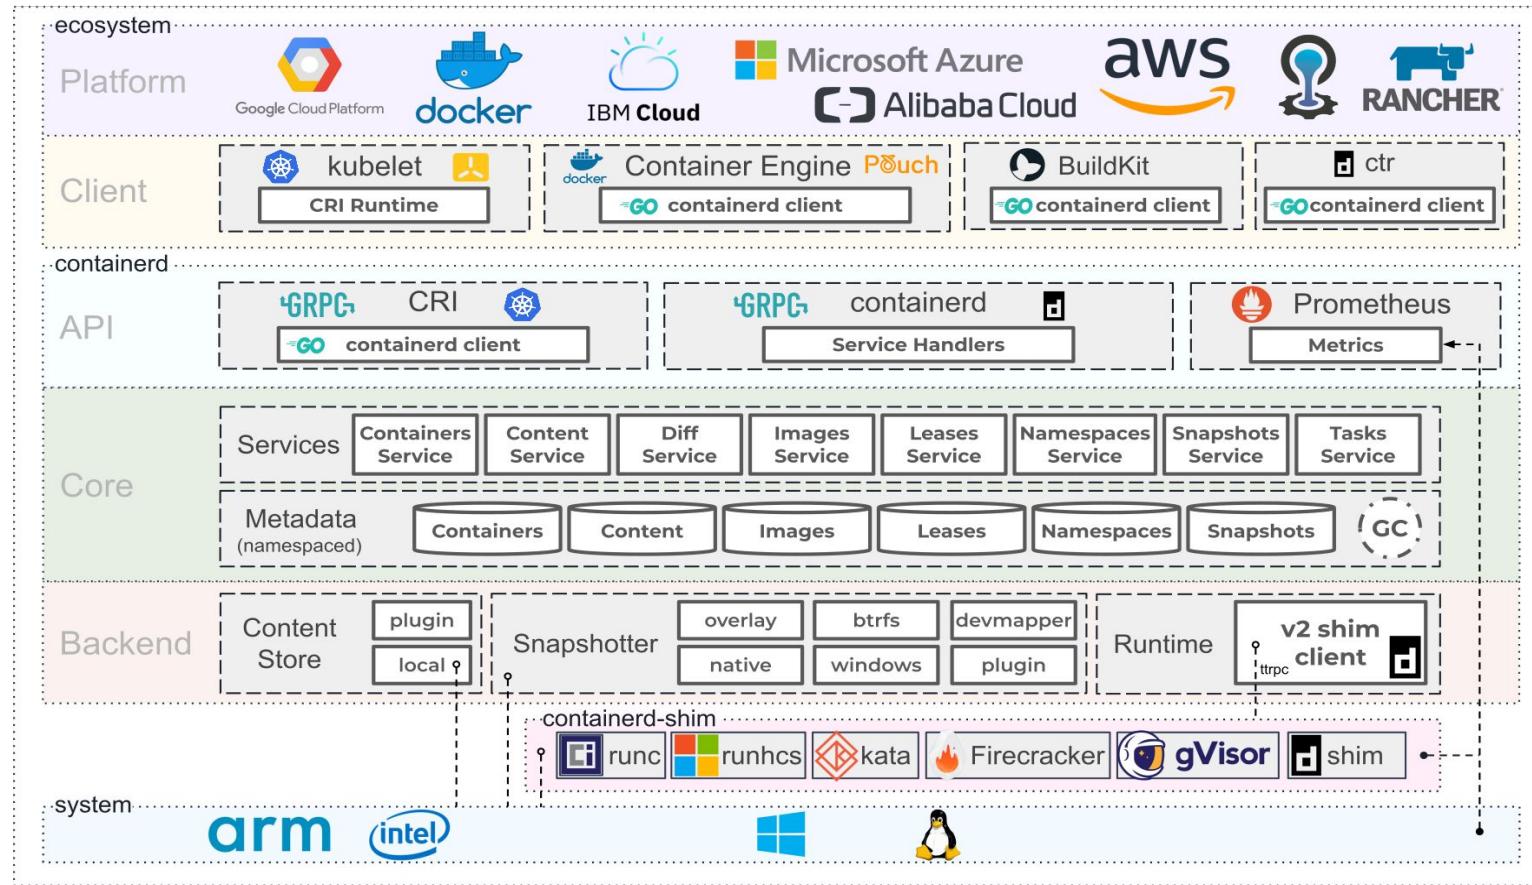
POD

Collection of containers that can run on a host.

This resource is created by clients and scheduled onto hosts.



Open specification they said...



Kubernetes Distributions

Clouds



- Azure Kubernetes
 - Service (AKS)

vmware®



Amazon
EKS

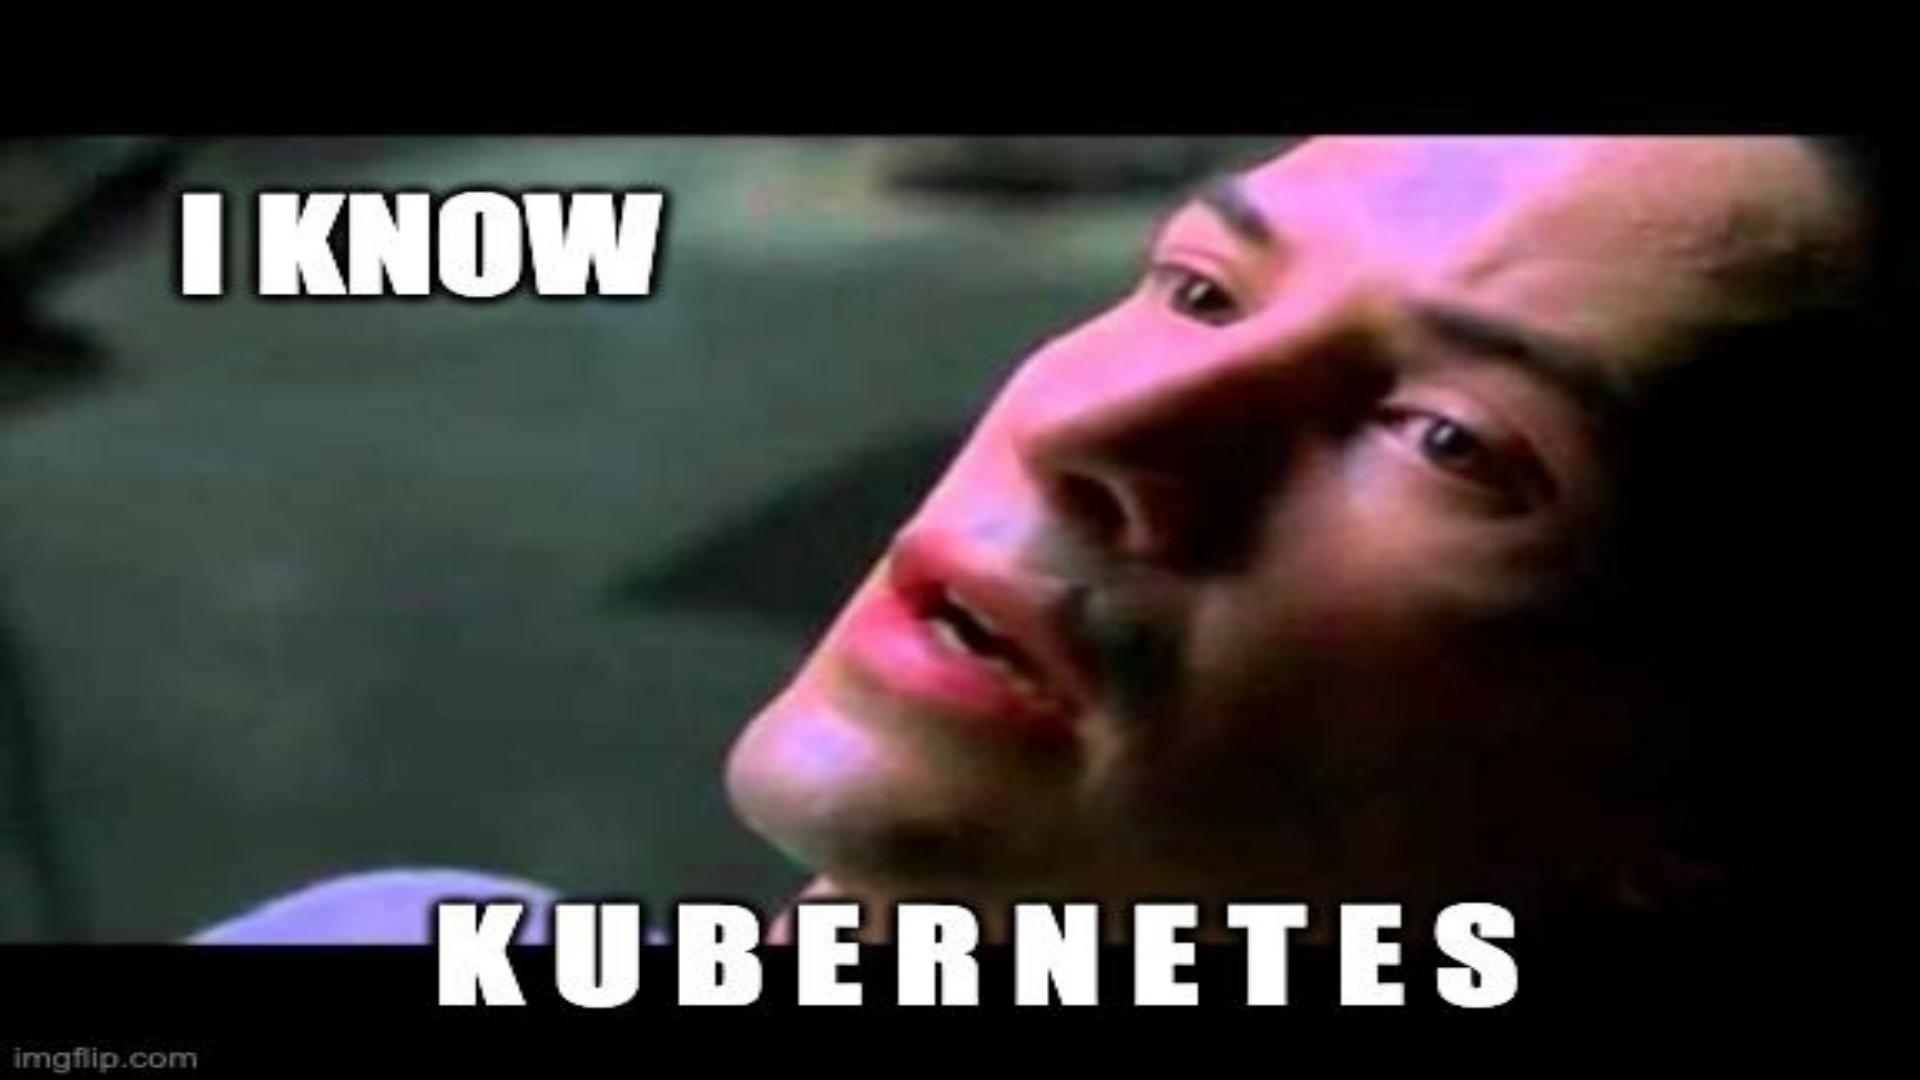


- Google Kubernetes
 - Engine (GKE)



IBM





I KNOW

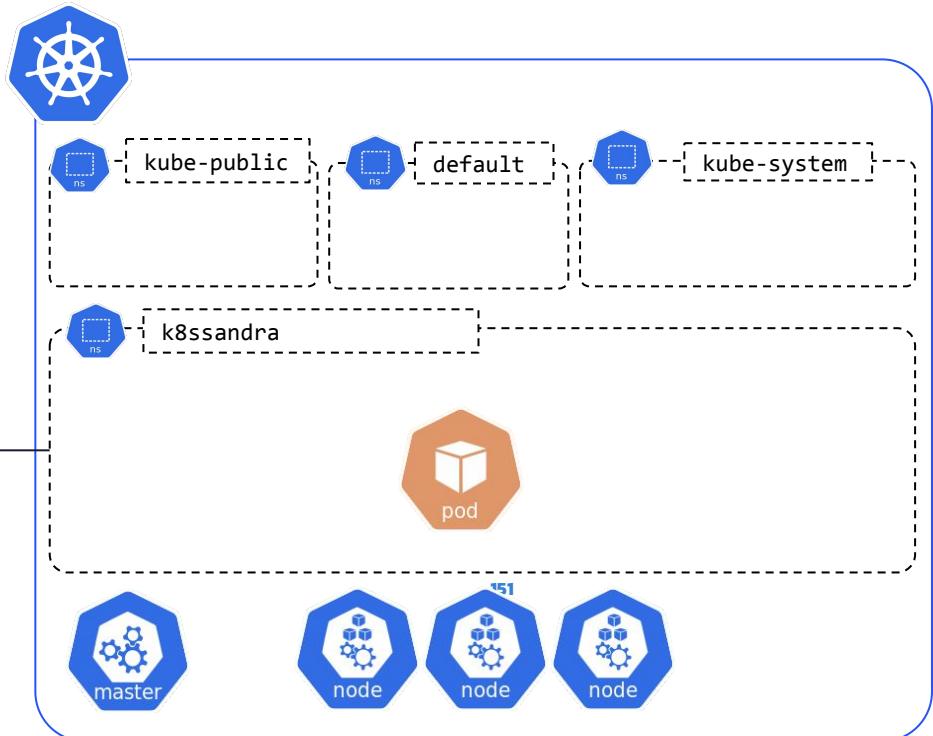
KUBERNETES

Kubernetes Namespaces

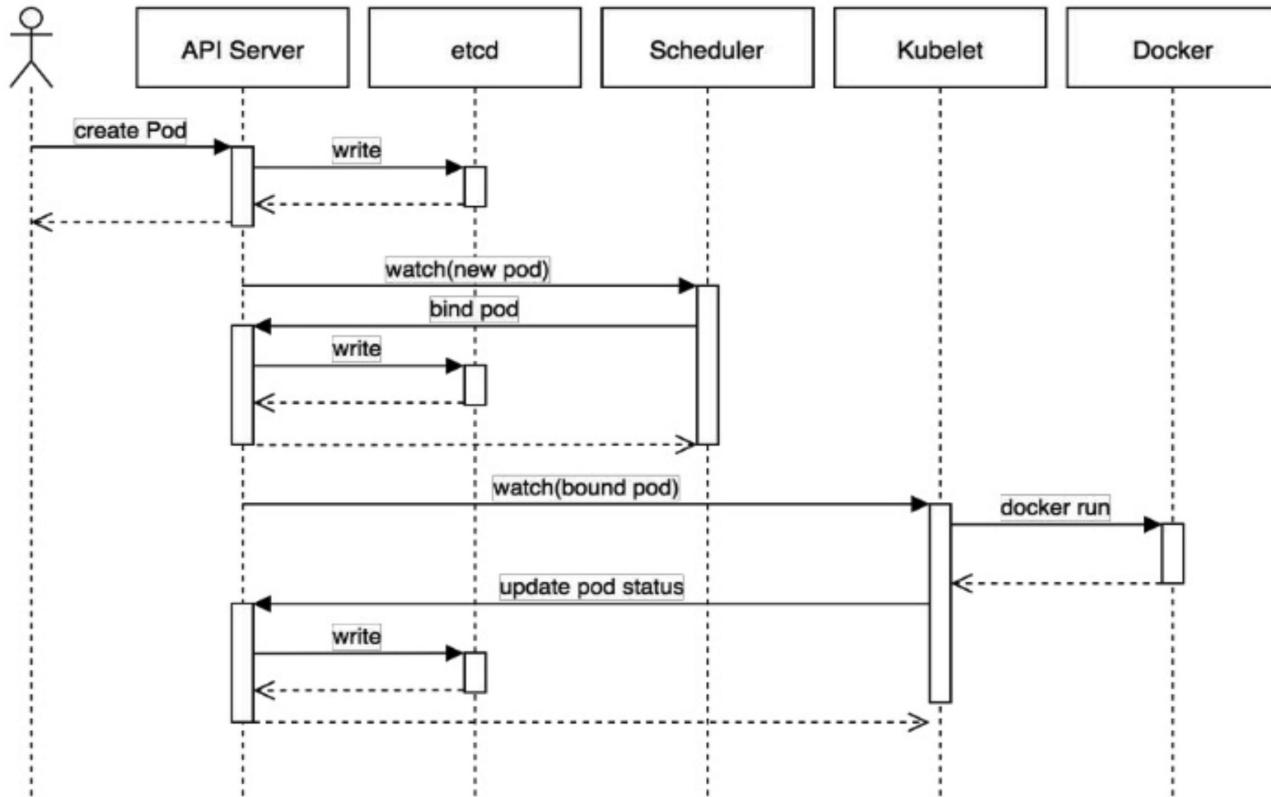


Namespace: Namespace provides a scope for Names. Use of multiple namespaces is optional.

We create resources in
namespaces that span across
nodes.



Pod Lifecycle



K8s Primitives : Storage



PersistentVolume: is a storage resource provisioned by an administrator or dynamic provisioner.

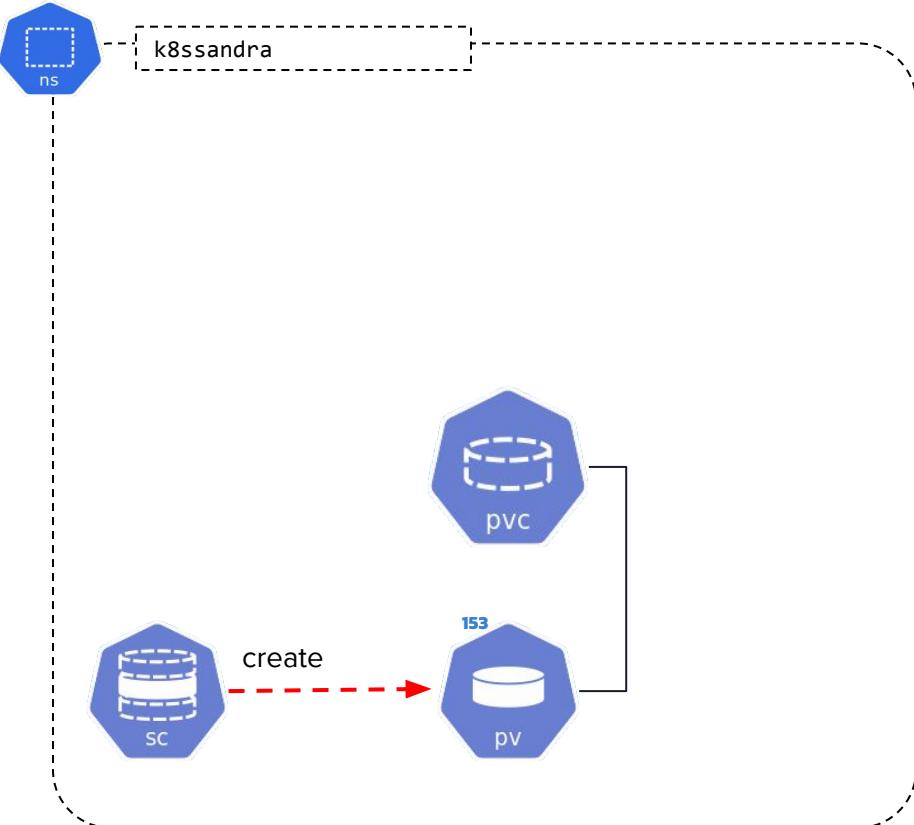


PersistentVolumeClaim:

PersistentVolumeClaim is a user's request for and claim to a persistent volume.



StorageClass: StorageClass describes the parameters for a class of storage for which *PersistentVolumes* can be dynamically provisioned.



K8s Primitives : Service



Ingress is a collection of rules that allow inbound connections to reach the endpoints defined by a backend.



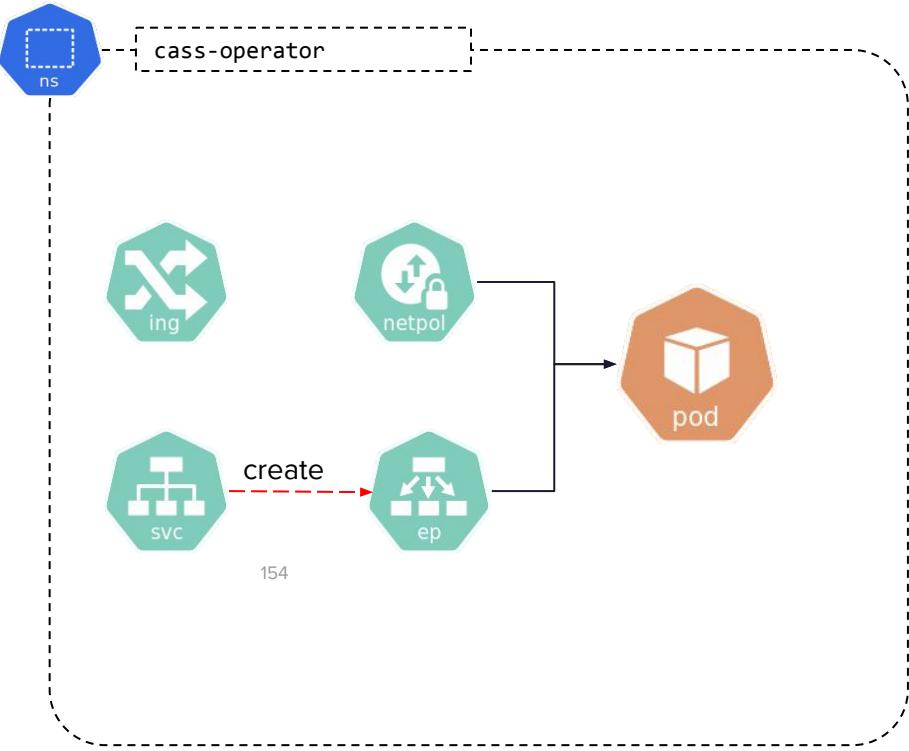
Service is a named abstraction of software service with ports to listen on and selector to determine which pods will answer requests.



EndPoint is a collection of endpoints that implement the actual service..



NetworkPolicy: Describes what network traffic is allowed for a set of Pods.



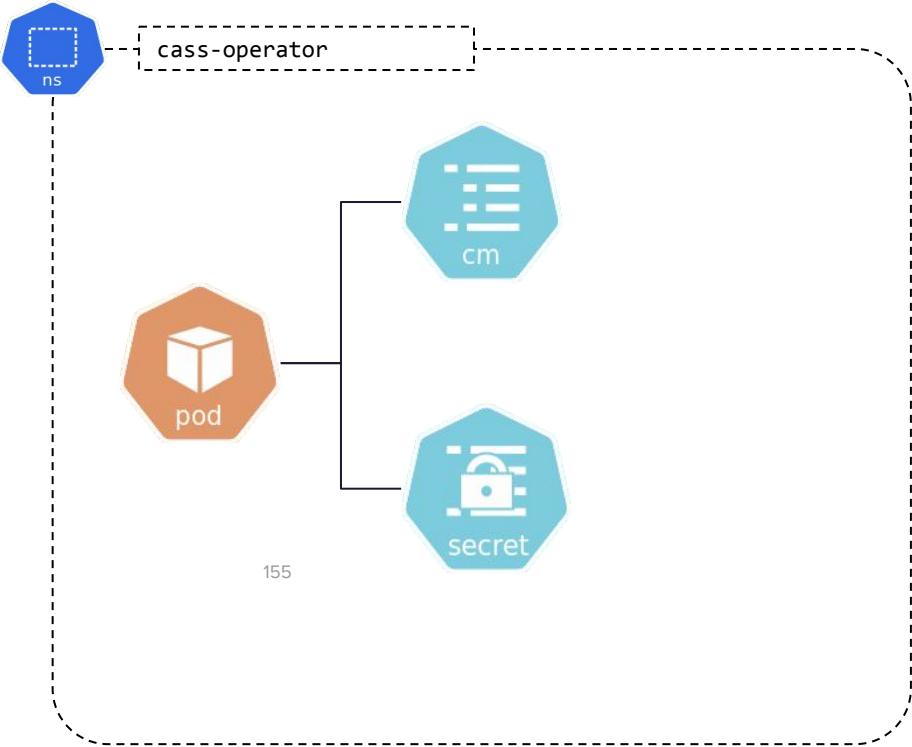
K8s Primitives : Configuration



ConfigMap: ConfigMap holds configuration data for pods to consume..



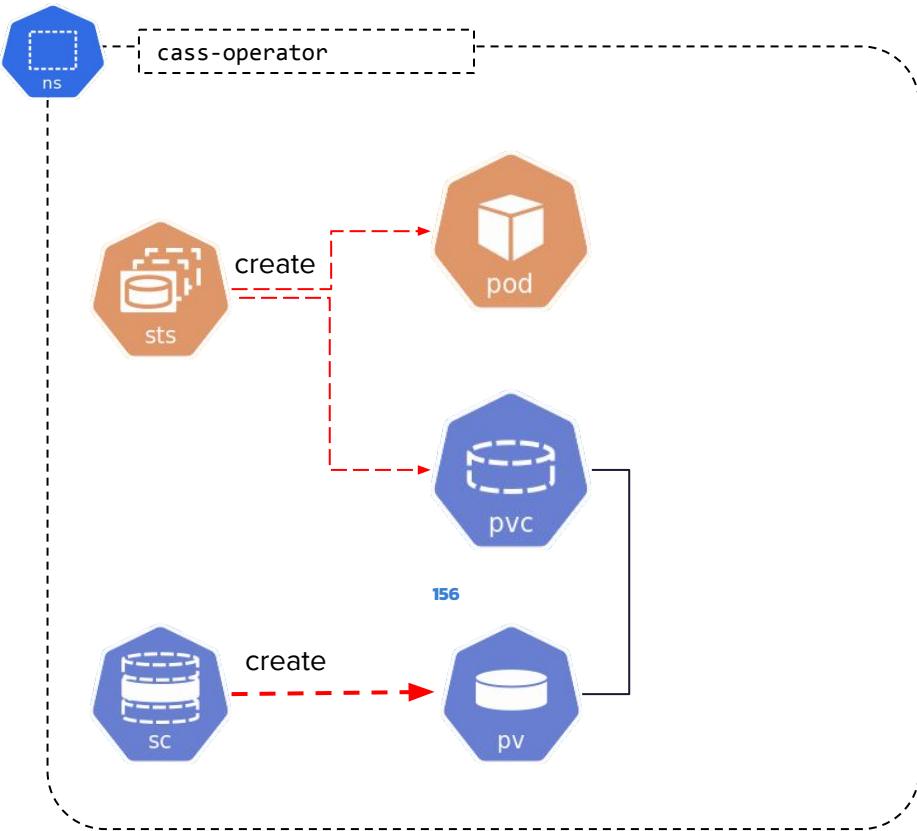
Secret: Secret holds secret data of a certain type..



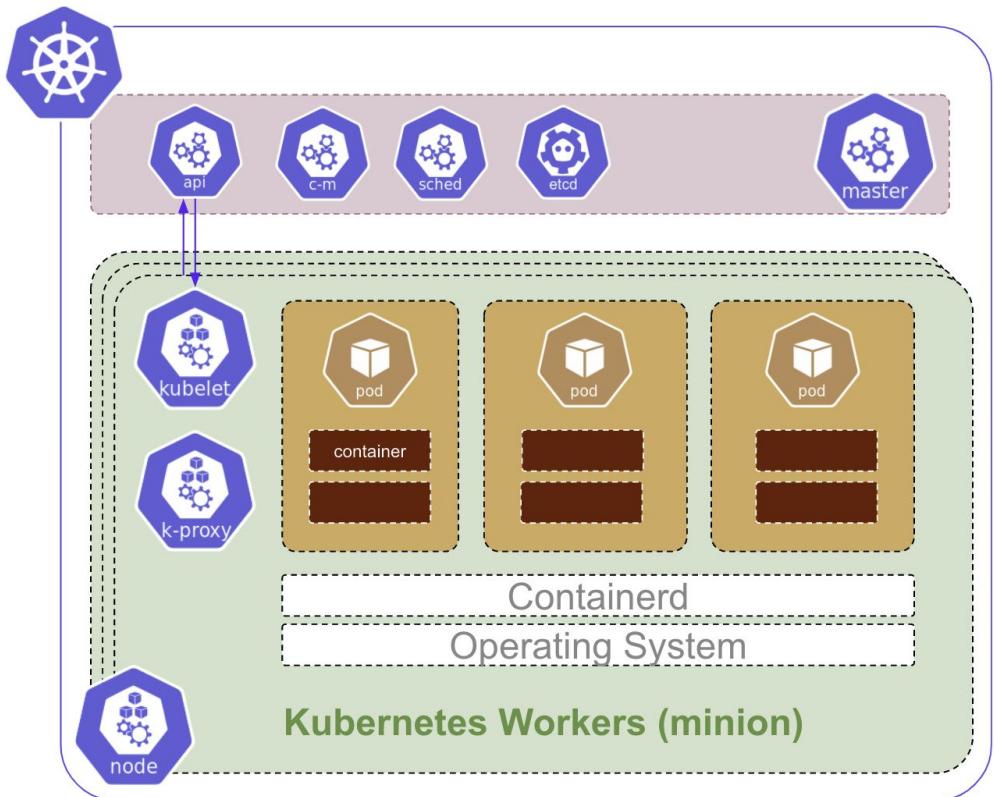
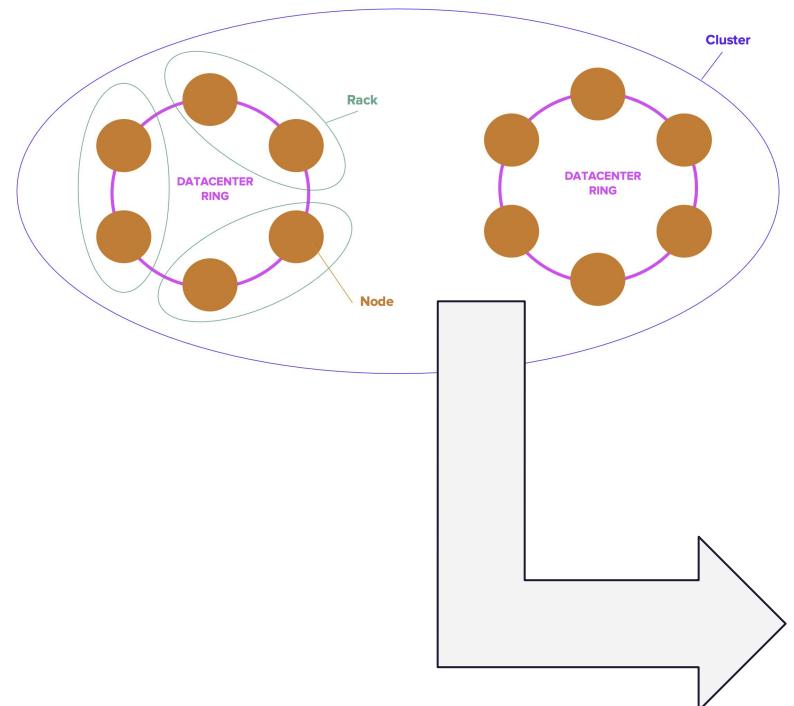
K8s Primitives : StatefulSet

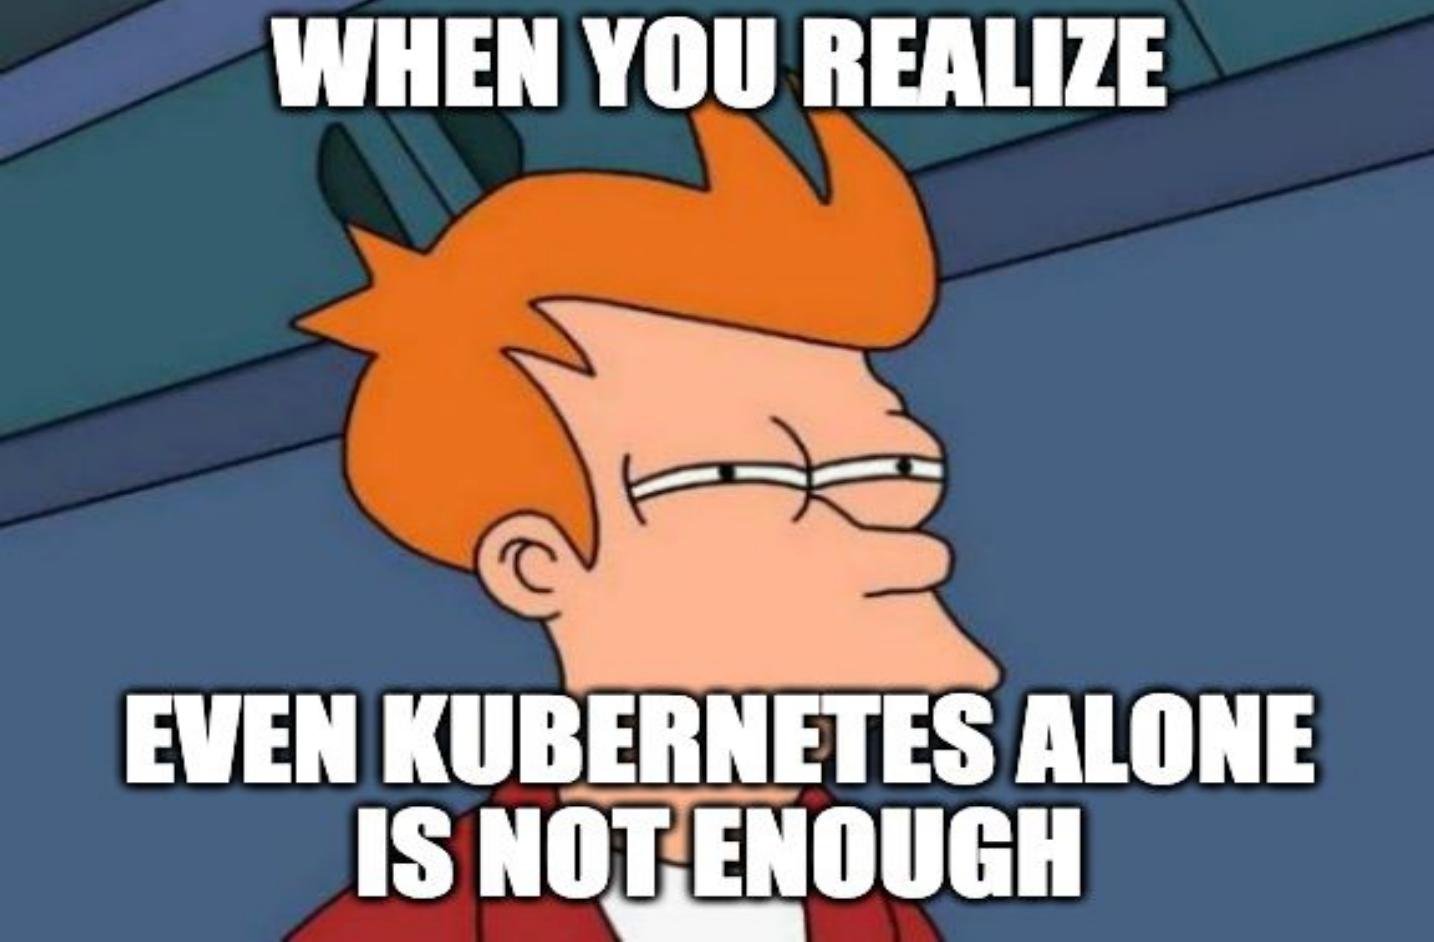


StatefulSet: StatefulSet represents a set of pods with consistent identities. Identities are defined as: network, storage.



Vous mission si vous décidez de l'accepter....





WHEN YOU REALIZE

**EVEN KUBERNETES ALONE
IS NOT ENOUGH**

Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

05



K8ssandra
Tour d'horizon

K8s Primitives : Custom Resources



Custom Resource Definition :

Extensions of the Kubernetes API.
Customization making K8s more modular

- **Spec** declares the desired state of a resource
 - Configuration settings provided by the user
 - Default values expanded by the system
 - Other properties initialized by other internal components after resource creation.
- **Status** : describes the object's current, observed state.
 - Kubernetes API server provides a REST API to clients. A Kubernetes object or resource is a REST resource.
 - The status of a Kubernetes resource is typically implemented as a **REST subresource** that can only be modified by internal, system components

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: cassandradatacenters.cassandra.datastax.com
spec:
  group: cassandra.datastax.com
  names:
    kind: CassandraDatacenter
    listKind: CassandraDatacenterList
    plural: CassandraDatacenters
    shortNames:
      - cassdc
      - cassdcs
    singular: CassandraDatacenter
  scope: Namespaced
  subresources:
    status: {}
...
...
```

Examples used in K8ssandra:

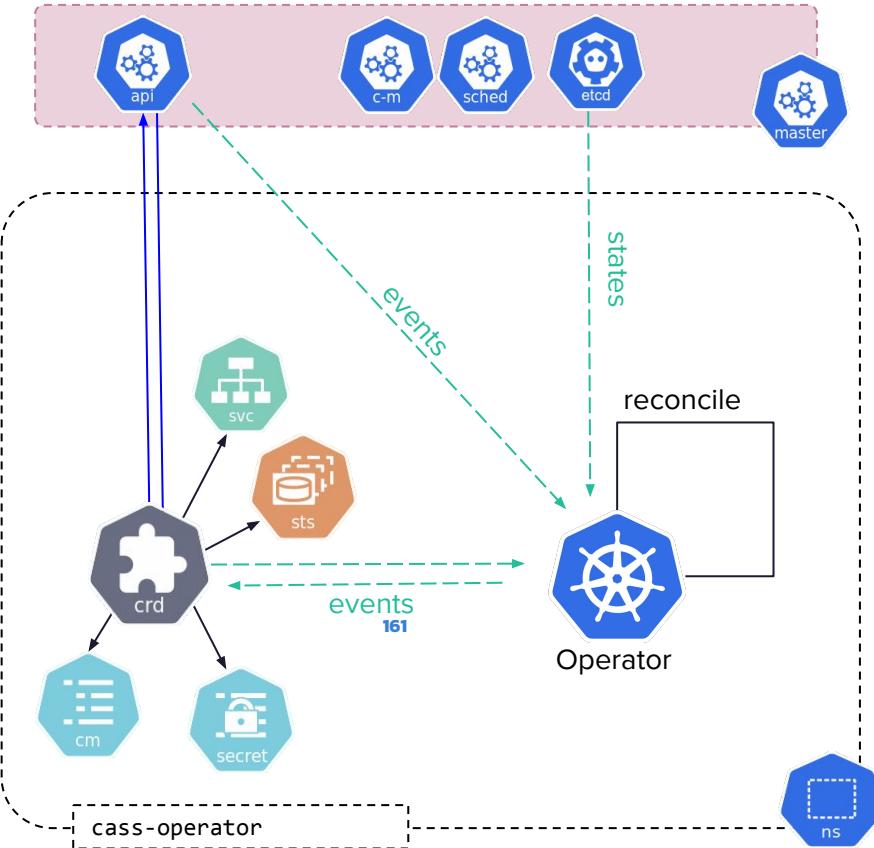
- CassandraDataCenter (**cass-operator**)
- CassandraBackup, CassandraRestore (**cassandra-medusa**)
- Reaper (**reaper**)
- GrafanaDataSource (**grafana**)

K8s Primitives : Operator

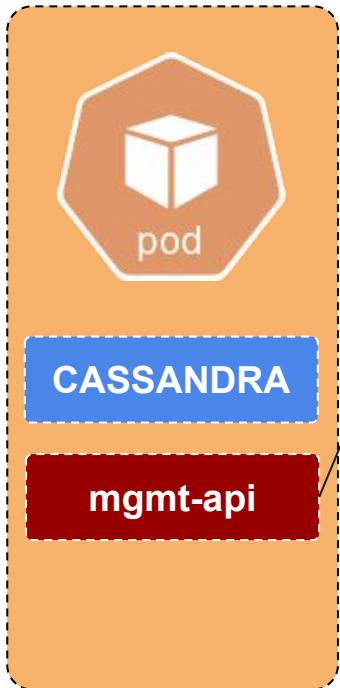
Building an application and driving an application on top of Kubernetes, behind Kubernetes APIs

A Kubernetes Operator helps extend the types of applications that can run on Kubernetes by allowing developers to provide additional knowledge to applications that need to maintain state." —*Jonathan S. Katz*

- **Reconcile** CRD instances which states defined within the "**spec**" attribute.
- Listen for **events** and **status evolution** to react accordingly.
- **Examples used in K8ssandra:**
 - cass-operator (`cass-operator`)
 - reaper-operator (`cassandra-reaper`)
 - PrometheusOperator (`Prometheus`)



Kubernetes Probes (readiness, liveness)



Cassandra Management API Service

<https://github.com/datastax/management-api-for-apache-cassandra>

Management API for Apache Cassandra 0.1 OAS3

<https://raw.githubusercontent.com/datastax/management-api-for-apache-cassandra/master/management-api-server/doc/openapi.json>

This is a Restful service for operating Apache Cassandra. You can find out more about the Management API on [Github](#)

Apache 2.0

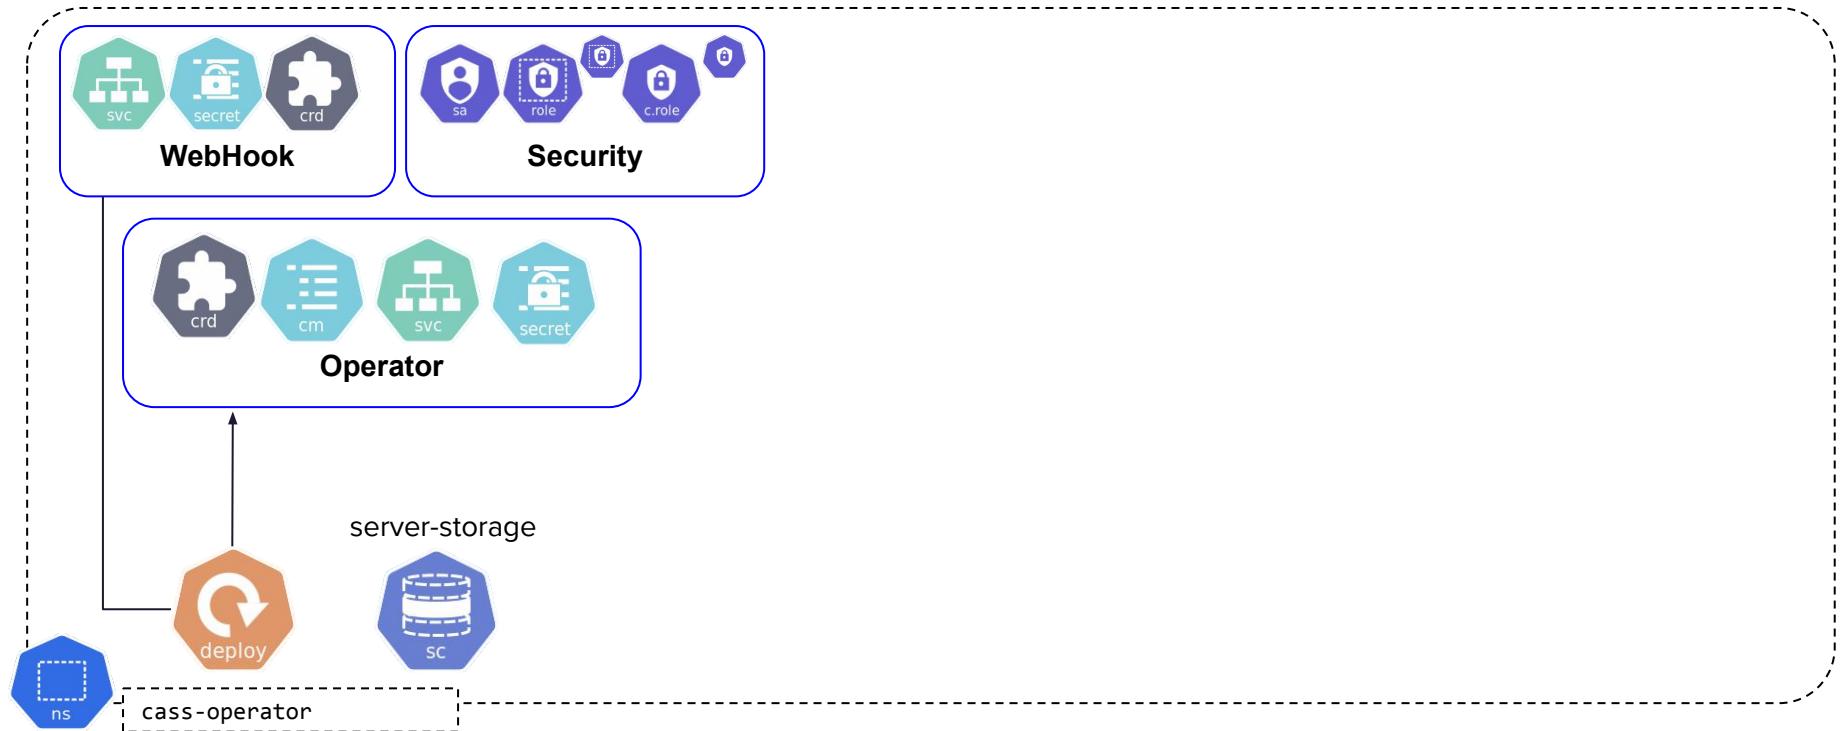
default

| | | |
|------|-----------------------------------|---|
| POST | /api/v0/ops/auth/role | Creates a new user role |
| GET | /api/v0/probes/liveness | Indicates whether this service is running |
| GET | /api/v0/probes/readiness | Indicates whether the Cassandra service is ready to service requests |
| GET | /api/v0/probes/cluster | Indicated whether the Cassandra cluster is able to achieve the specified consistency |
| POST | /api/v0/ops/seeds/reload | |
| POST | /api/v0/ops/keyspace/refresh | Load newly placed SSTables to the system without restart |
| POST | /api/v0/ops/keyspace/cleanup | Triggers the immediate cleanup of keys no longer belonging to a node. By default, clean all keyspaces |
| POST | /api/v0/lifecycle/start | |
| POST | /api/v0/lifecycle/stop | |
| POST | /api/v0/lifecycle/configure | |
| GET | /api/v0/lifecycle/pid | |
| GET | /api/v0/metadata/versions/release | Returns the Cassandra release version |
| GET | /api/v0/metadata/endpoints | Returns this nodes view of the endpoint states of nodes |
| POST | /api/v0/ops/node/drain | Drain the node (stop accepting writes and flush all tables) |

Cass Operator : Features

- Proper token **ring initialization**, with only one node bootstrapping at a time
- **Seed node management** -
 - one per rack, or three per datacenter, whichever is more
- Server configuration integrated into the **CassandraDatacenter CRD**
 - Rolling reboot nodes by changing the CRD
 - Store data in a rack-safe way - one replica per cloud AZ
 - Scale up racks evenly with new nodes
 - Replace dead/unrecoverable nodes
- Multi DC clusters (limited to one Kubernetes namespace)

Installing the Cass Operator Manifest

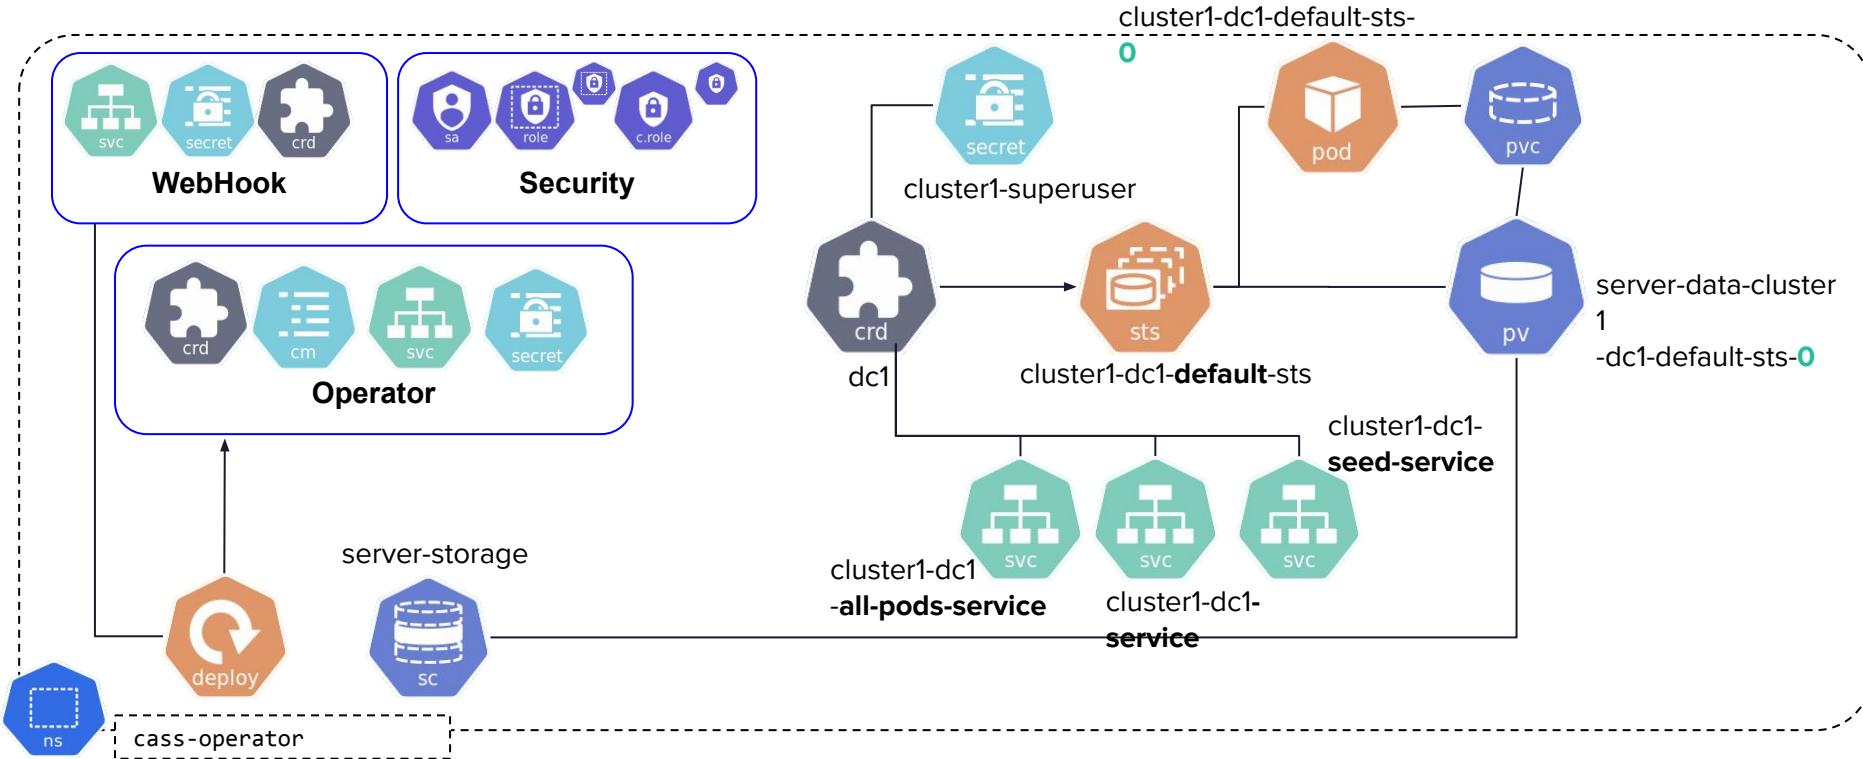


CRD CassandraDataCenter

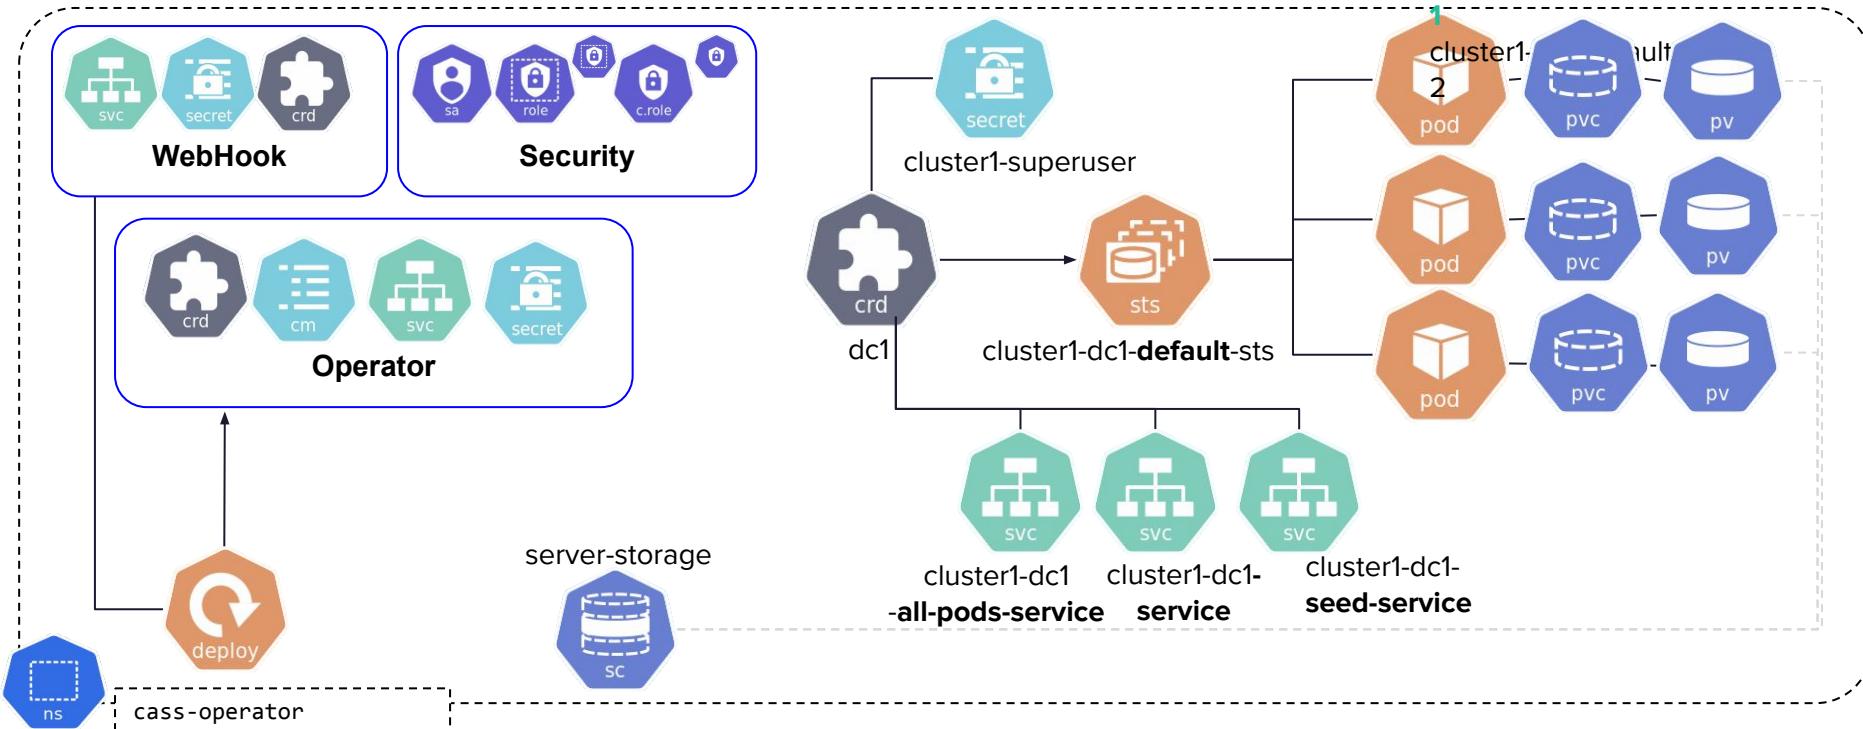
```
apiVersion: cassandra.datastax.com/v1beta1
kind: CassandraDatacenter
metadata:
  name: dc1
spec:
  clusterName: cluster1
  serverType: cassandra
  serverVersion: "3.11.6"
  managementApiAuth:
    insecure: {}
  size: 1
  storageConfig:
    cassandraDataVolumeClaimSpec:
      storageClassName: server-storage
      accessModes:
        - ReadWriteOnce
      resources:
        requests:
          storage: 5Gi
  config:
    cassandra-yaml:
      authenticator: org.apache.cassandra.auth.PasswordAuthenticator
      authorizer: org.apache.cassandra.auth.CassandraAuthorizer
      role_manager: org.apache.cassandra.auth.CassandraRoleManager
  jvm-options:
    initial_heap_size: "800M"
    max_heap_size: "800M"
```

```
apiVersion: cassandra.datastax.com/v1beta1
kind: CassandraDatacenter
metadata:
  name: multi-rack
spec:
  clusterName: multi-rack
  serverType: cassandra
  serverVersion: 3.11.6
  managementApiAuth:
    insecure: {}
  size: 9
  racks:
    - name: us-east1-b
      zone: us-east1-b
    - name: us-east1-c
      zone: us-east1-c
    - name: us-east1-d
      zone: us-east1-d
  storageConfig:
    cassandraDataVolumeClaimSpec:
      storageClassName: standard
      accessModes:
        - ReadWriteOnce
      resources:
        requests:
          storage: 5Gi
```

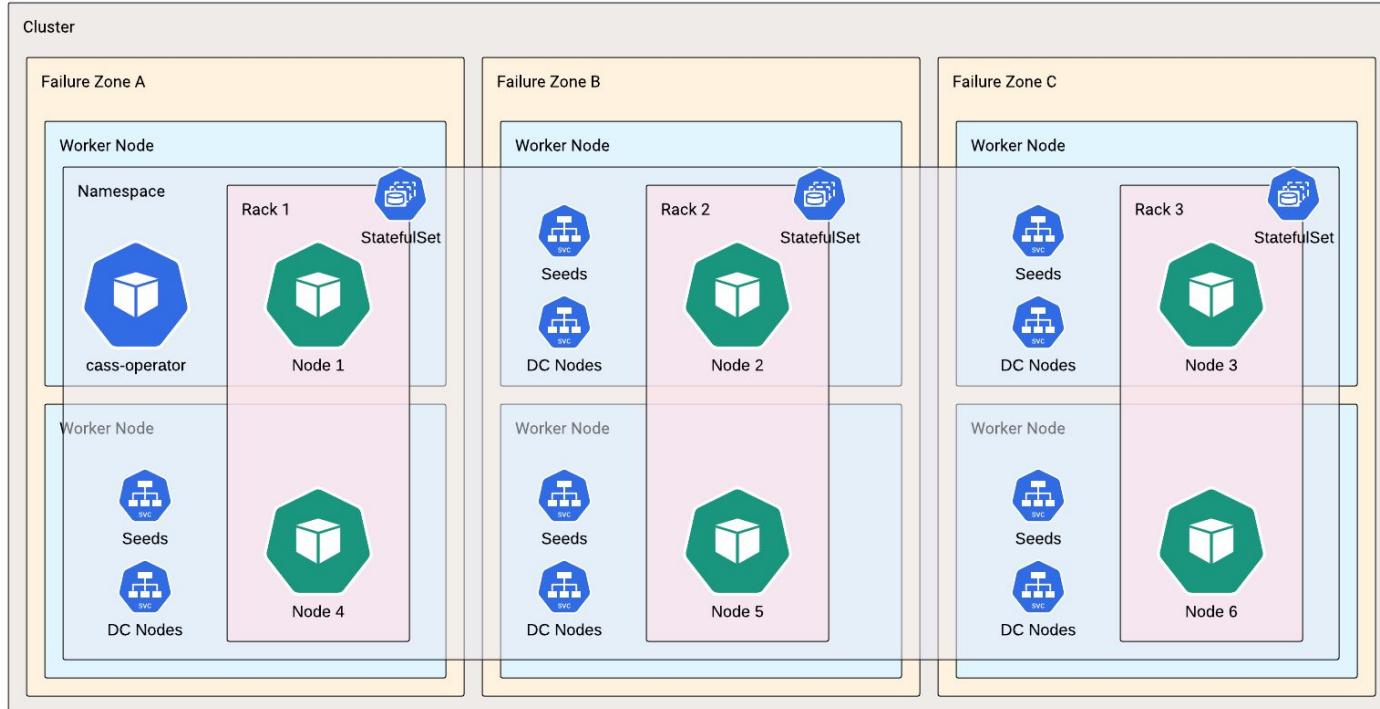
Creating DataCenter dc1



Scale up DataCenter dc1



Example Deployment Across Multiple K8s Workers



Hands-ON

<https://www.datastax.com/dev/kubernetes>

DataStax Kubernetes Operator for Apache Cassandra™

Deploy Cassandra using the Kubernetes Operator

Simplify your ops-life by using the Cassandra Kubernetes Operator to run your cluster!

In this scenario, we'll learn how to:

- >Create a Kubernetes cluster
- Install the Cassandra Kubernetes operator
- Create a single node cluster
- Create a keyspace and table with data in the cluster
- Scale up the Cassandra cluster to two nodes using Kubernetes

With Kubernetes and this operator, your Cassandra life just got a lot easier!

CONTRIBUTORS



START SCENARIO

TIME TO COMPLETE

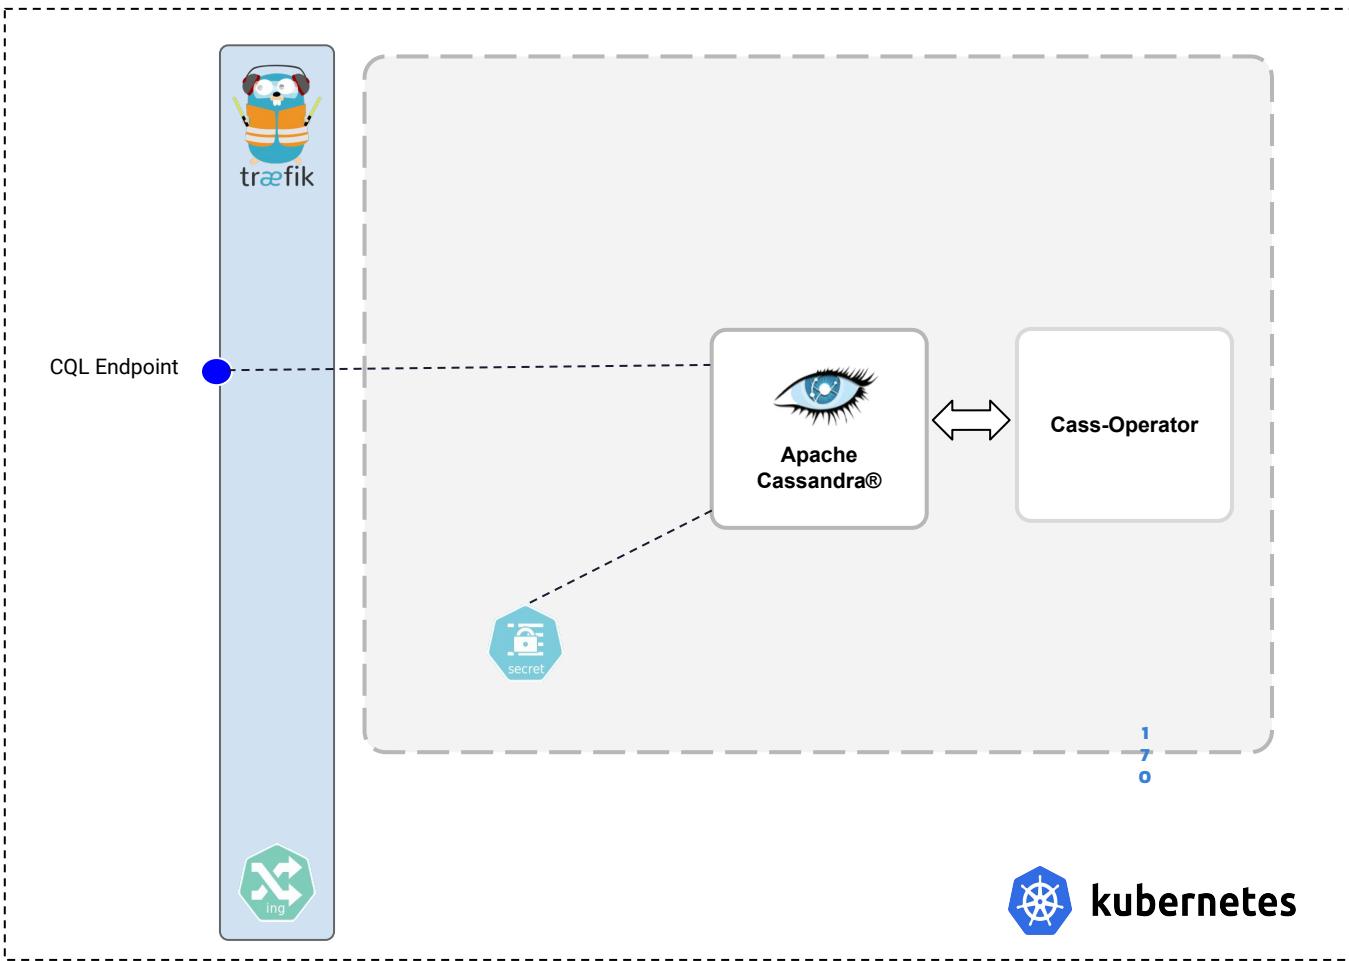
20 minutes

DIFFICULTY

Intermediate

SHARE





MCAC = Cassandra Management API SideCar

The Management API is a sidecar service layer that attempts to build a well supported set of operational actions on Cassandra® nodes that can be administered centrally. It currently works with official Apache Cassandra® 3.11.x and 4.0 via a drop-in Java agent.

- Lifecycle Management
 - Start Node
 - Stop Node
- Configuration Management (alpha)
 - Change YAML
 - Change jvm-opts
- Health Checks
 - Kubernetes liveness/readiness checks
 - Consistency level checks
- Per node actions
 - All nodetool commands

Management API for Apache Cassandra 0.1.0 AS3
<https://www.githubusercontent.com/datasstax/management-api-for-apache-cassandra/master/management-api-server/doc/openapi.json>

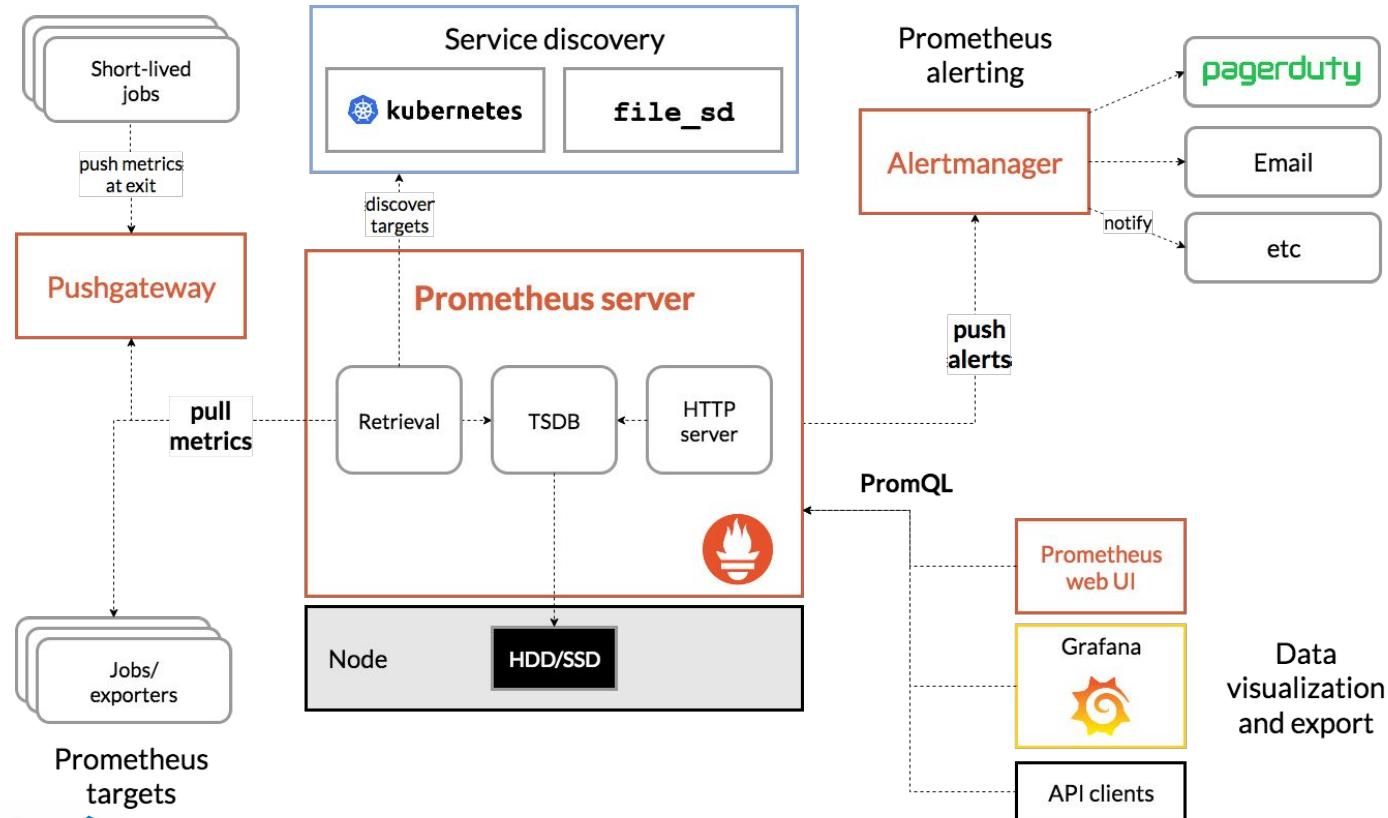
This is a RESTful service for operating Apache Cassandra. You can find out more about the Management API on [GitHub](#).

Apache 2.0

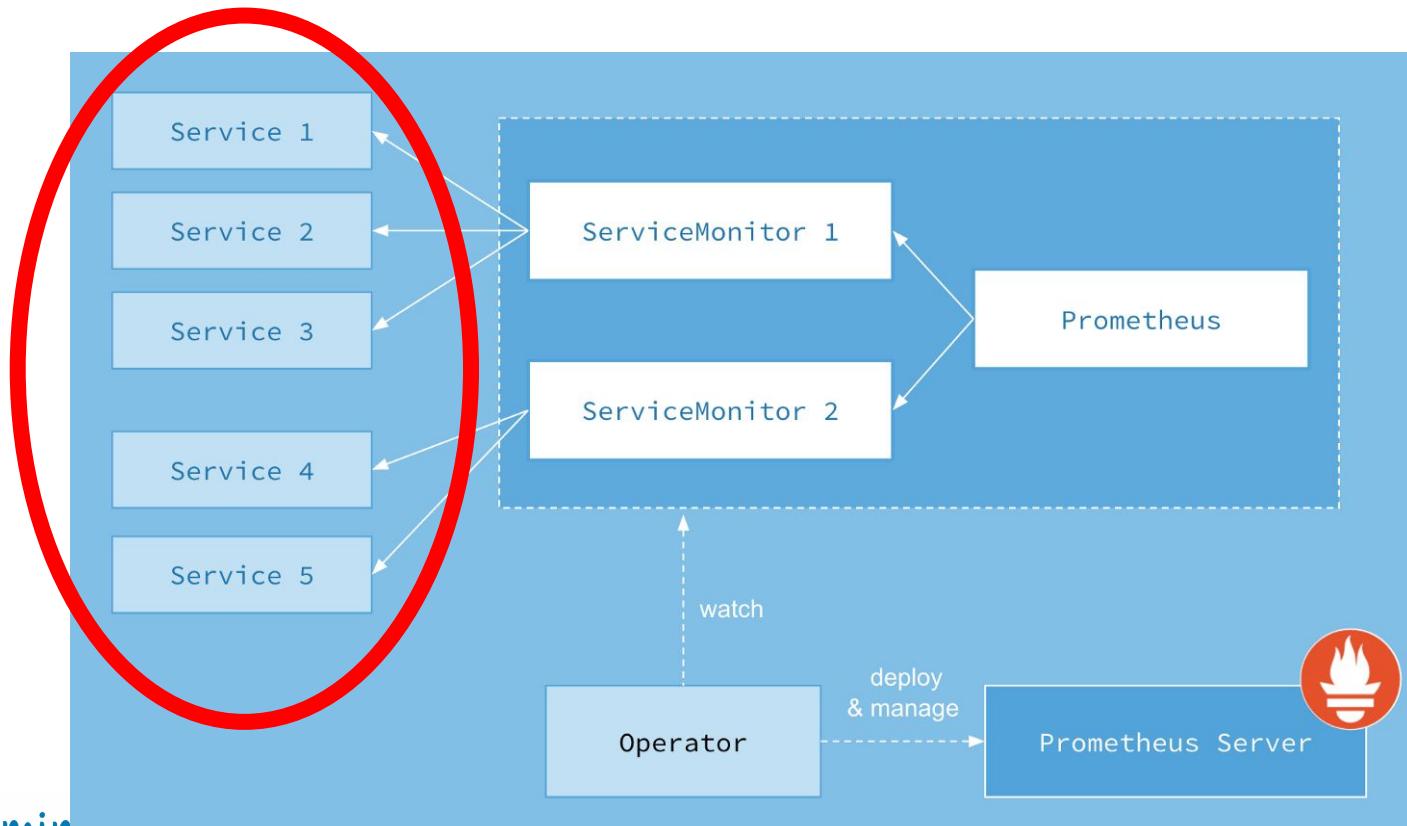
default

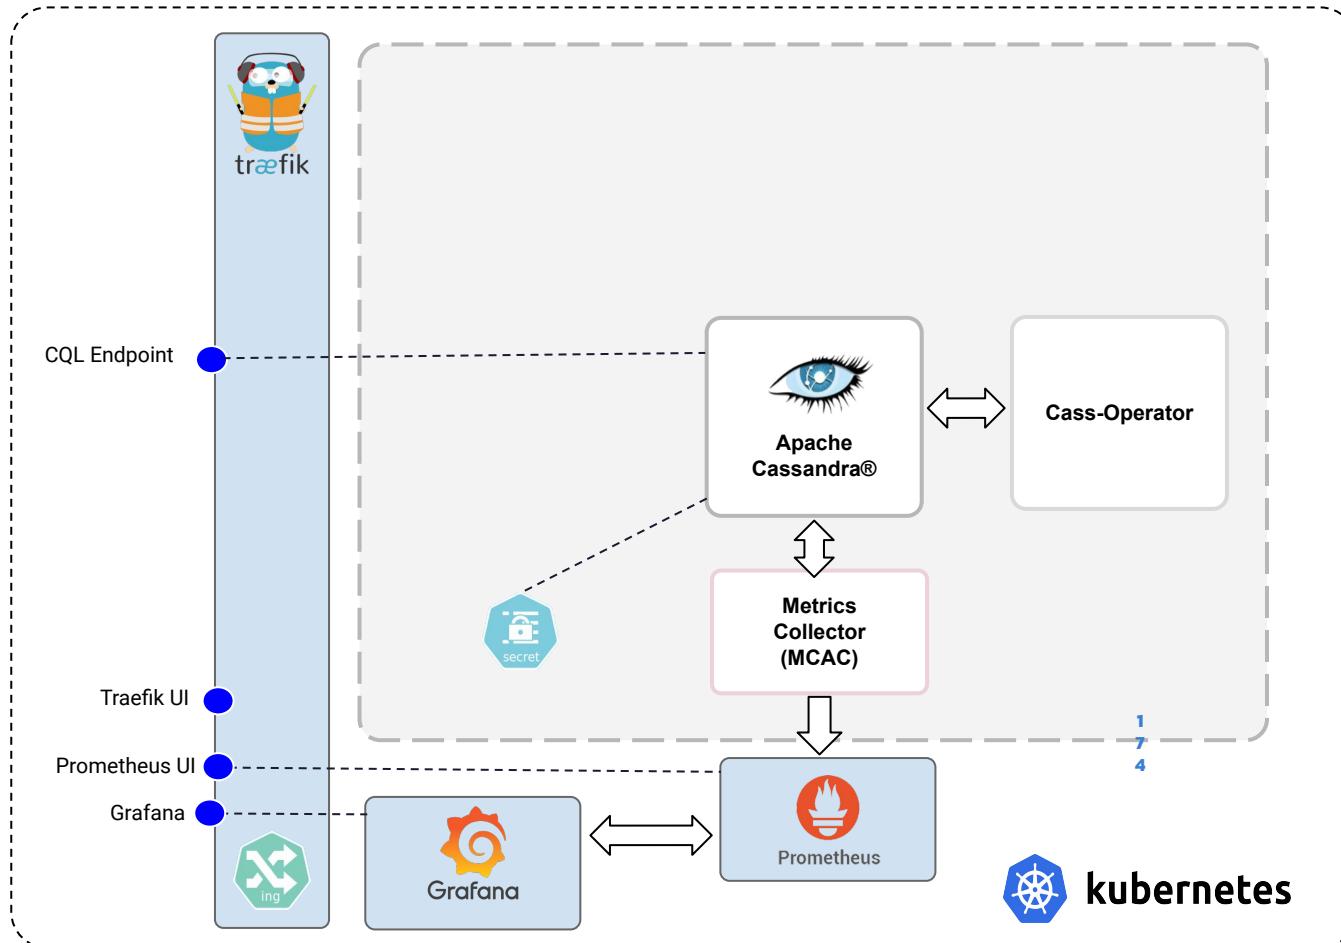
| | |
|-------------|--|
| POST | /api/v0/ops/auth/role Creates a new user role |
| GET | /api/v0/probes/liveness Indicates whether this service is running |
| GET | /api/v0/probes/readiness Indicates whether the Cassandra service is ready to service requests |
| GET | /api/v0/probes/cluster Indicated whether the Cassandra cluster is able to achieve the specified consistency |
| POST | /api/v0/ops/seeds/reload |
| POST | /api/v0/ops/keyspace/refresh Load newly placed SSTables to the system without restart |
| POST | /api/v0/ops/keyspace/cleanup Triggers the immediate cleanup of keys no longer belonging to a node. By default, clean all keyspaces |
| POST | /api/v0/lifecycle/start |
| POST | /api/v0/lifecycle/stop |
| POST | /api/v0/lifecycle/configure |
| GET | /api/v0/lifecycle/pid |
| GET | /api/v0/metadata/versions/release Returns the Cassandra release version |
| GET | /api/v0/metadata/endpoints Returns this nodes view of the endpoint states of nodes |
| POST | /api/v0/ops/node/drain Drain the node (stop accepting writes and flush all tables) |

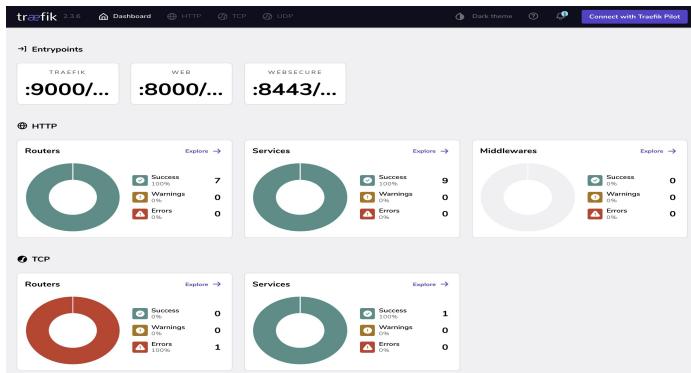
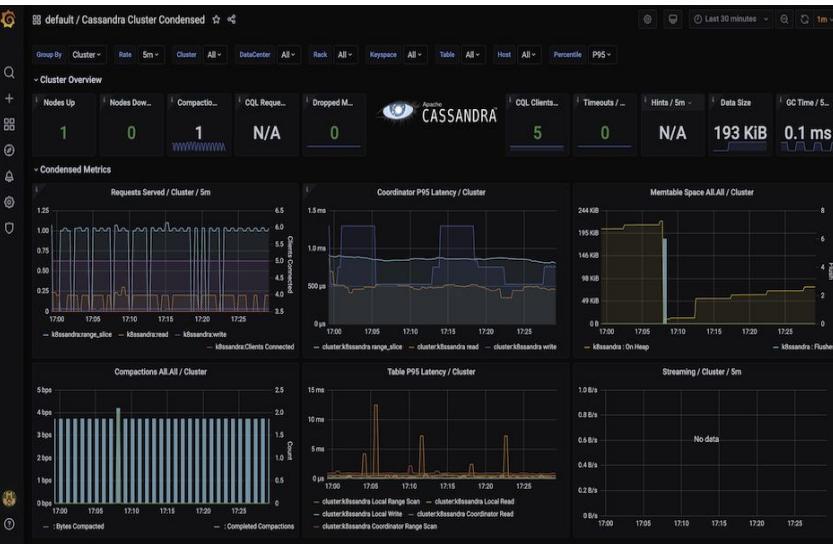
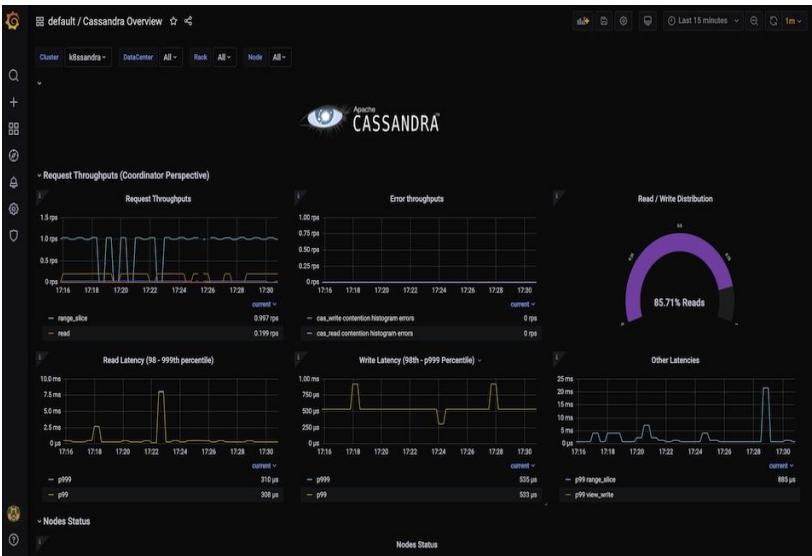
Prometheus + Grafana



Prometheus metrics collection







Agenda (6H)

01



Apache Cassandra™
Les fondamentaux

02



Apache Cassandra™
Modélisation

03



Apache Cassandra™
Développements

04



Kubernetes
Les fondamentaux

05



cassandra

Kubernetes
Stateful & Cassandra

05



K8ssandra
Tour d'horizon



+

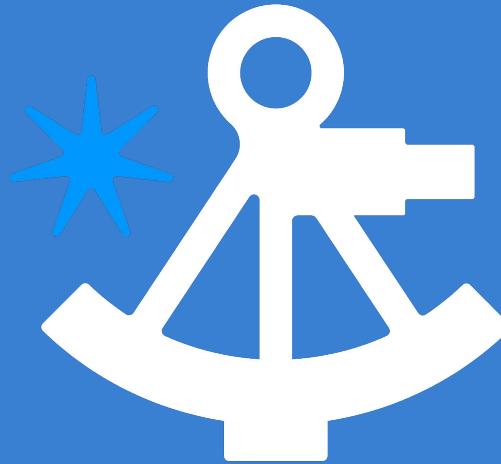


=



K8SSANDRA

Cloud native, scalable data tier with administration tools and easy data access



The screenshot shows the K8ssandra website homepage. At the top right is a navigation bar with links: About, Get Started, Docs, Community, Forum, Blog, and Events. A search bar is located at the top right corner. The main title "K8ssandra" is prominently displayed in the center, with a subtitle below it: "Built on the rock-solid Apache Cassandra® NoSQL database, K8ssandra brings together a complete operational data platform for Kubernetes including APIs, monitoring, and backups." Below this are two buttons: "Get Started" and "Read the Docs". To the right is a sidebar titled "Latest Articles" featuring three recent posts: "K8ssandra and the Community" by Sean Boggard, "KubeCon / CloudNativeCon – North America 2021" by Kate Sandra, and "Running a Cloud-Native Data Tier with K8ssandra hosted by DataStax" by Jeff Carpenter.

K8ssandra

Built on the rock-solid Apache Cassandra® NoSQL database, K8ssandra brings together a complete operational data platform for Kubernetes including APIs, monitoring, and backups.

[Get Started](#) [Read the Docs](#)

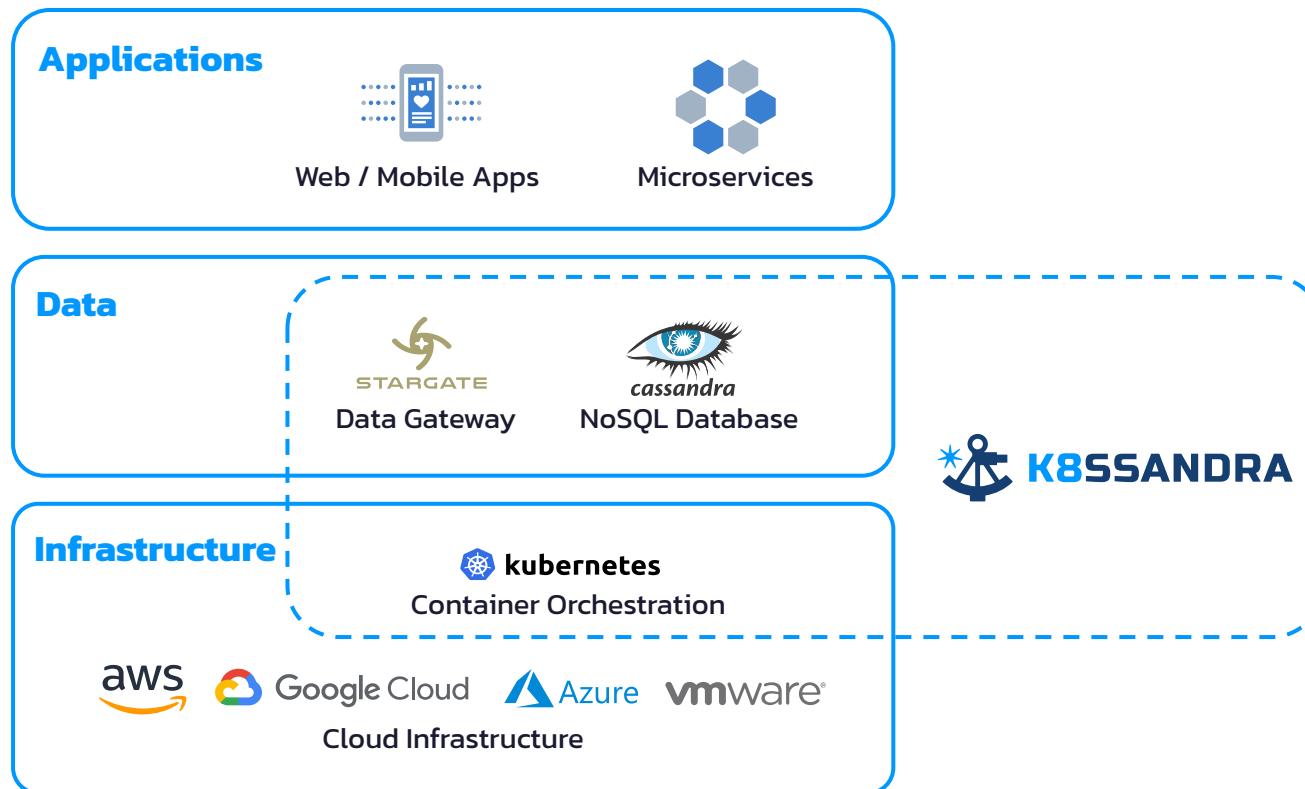
Latest Articles

K8ssandra and the Community
Sean Boggard — 04.28.2021 — ARTICLES

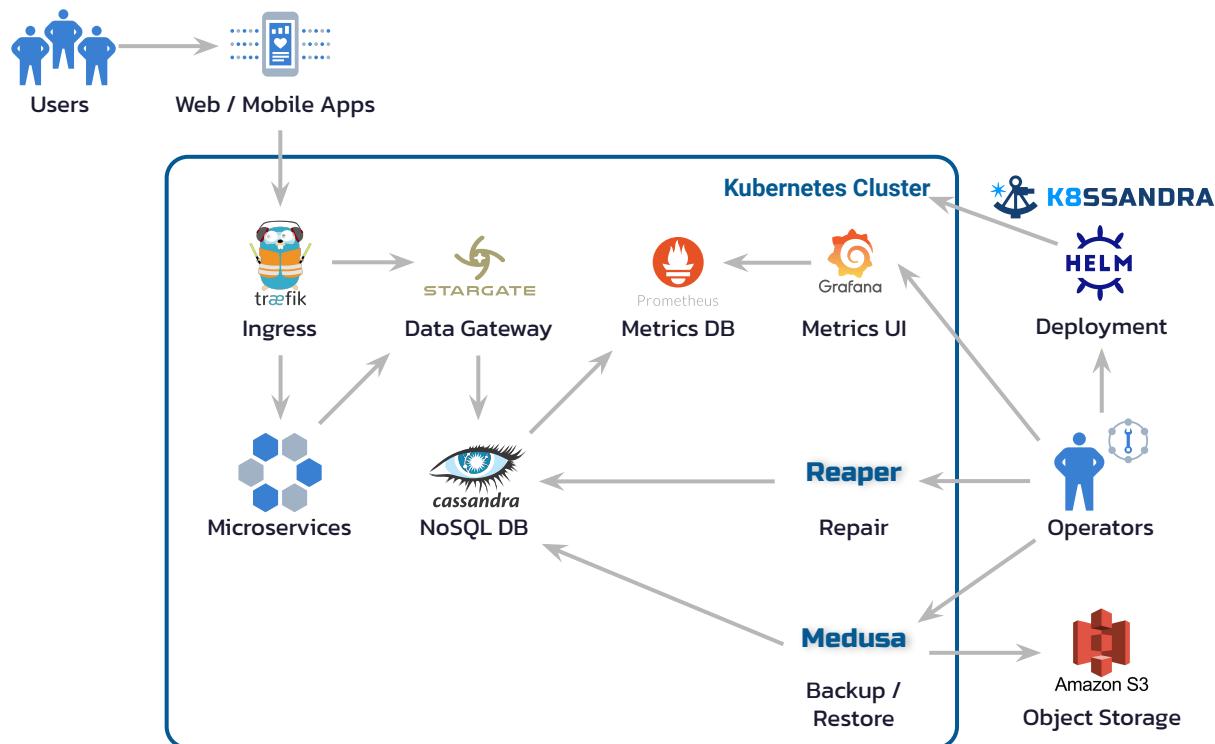
KubeCon / CloudNativeCon – North America 2021
Kate Sandra — 04.23.2021 — CONFERENCE, EVENTS

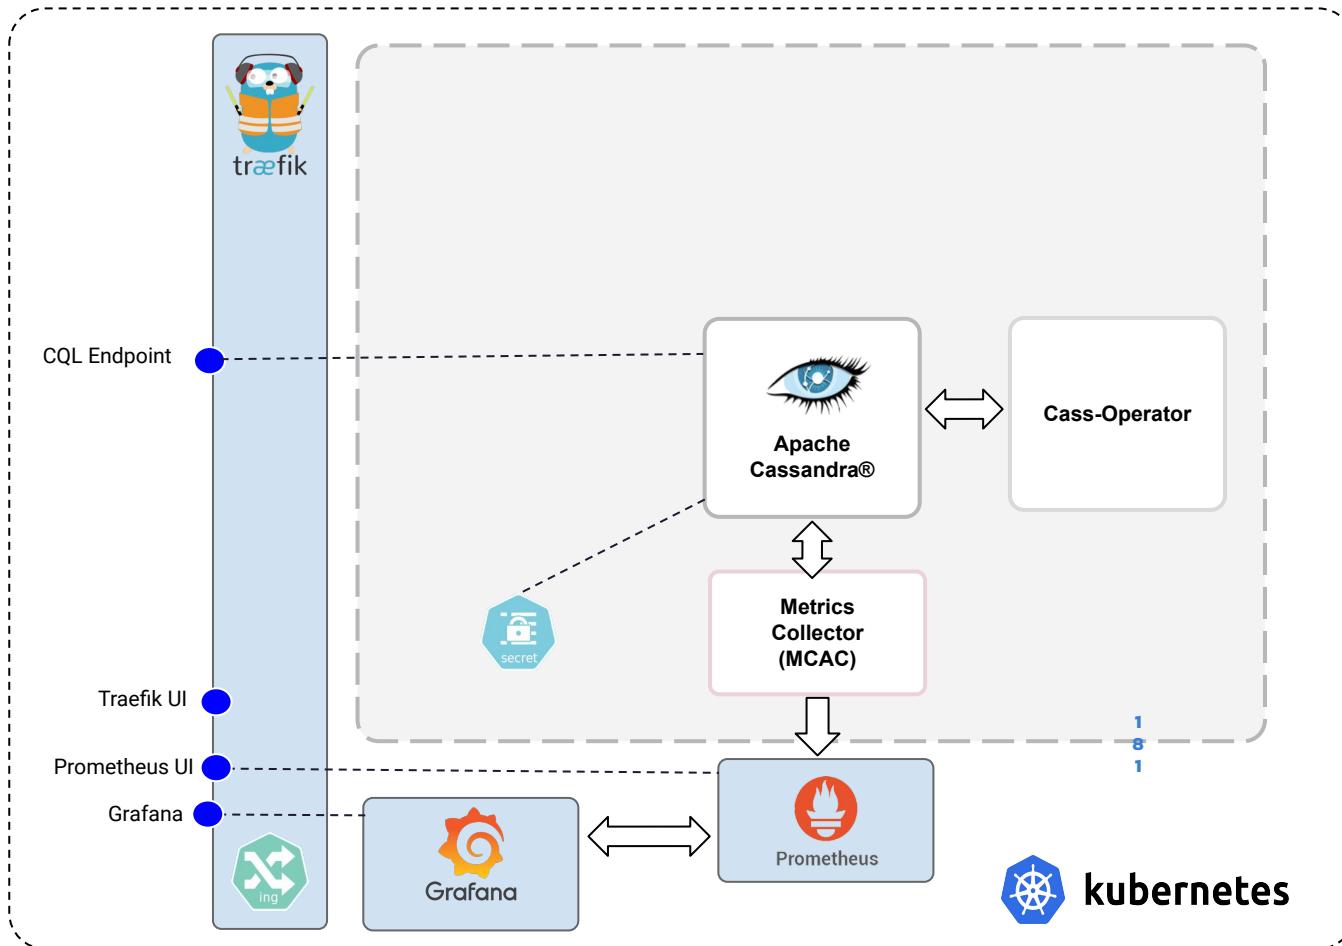
Running a Cloud-Native Data Tier with K8ssandra hosted by DataStax
Jeff Carpenter — 04.23.2021 — ARTICLES

Highly Scalable Cloud Application Architecture



K8ssandra Components in Context

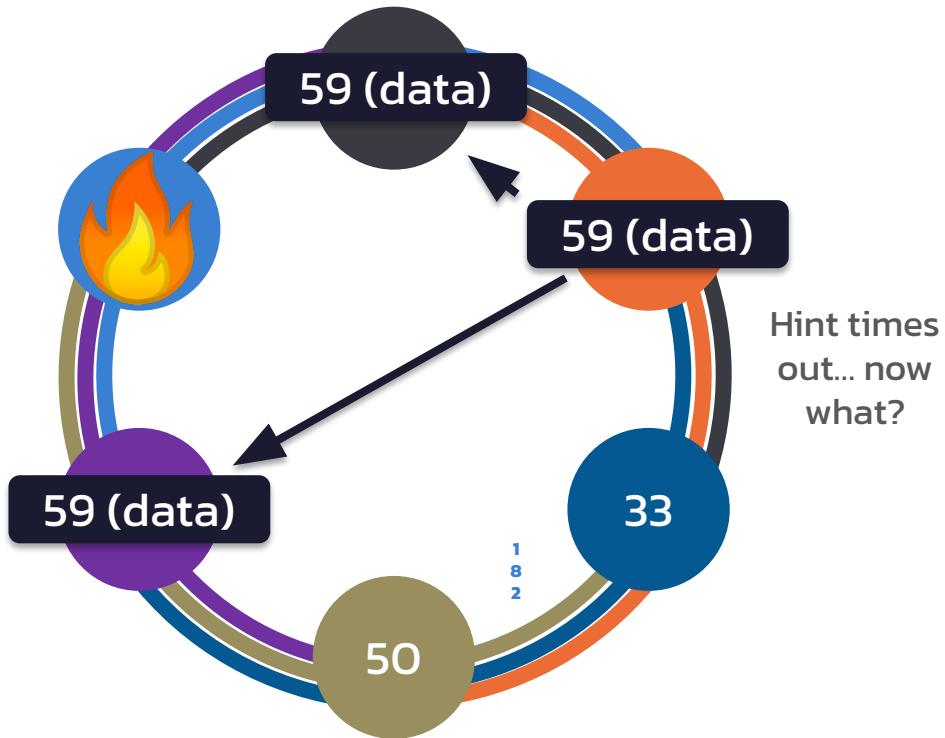




Repairs in Cassandra

- Repairs are the “last resort” to keep data across a Cassandra cluster in sync
- If repairs don’t happen the health of the data will slowly degrade over time due to Entropy

“Without repairs, your database is *hopefully consistent*”
– Vinay Kumar Chella, Principal Cloud Database Architect, Netflix



Cassandra Reaper

The screenshot shows the 'Cluster' section of the Cassandra Reaper interface. A red arrow points to the 'kIsandra' node entry, which displays the following information:

- Nodes: 1
- Total load: 568 kB/s
- Running repairs: 0
- Forget cluster: [Info](#)

The screenshot shows the 'Schedules' section of the Cassandra Reaper interface. A red arrow points to the 'Add schedule' button at the top of the page.

Repair

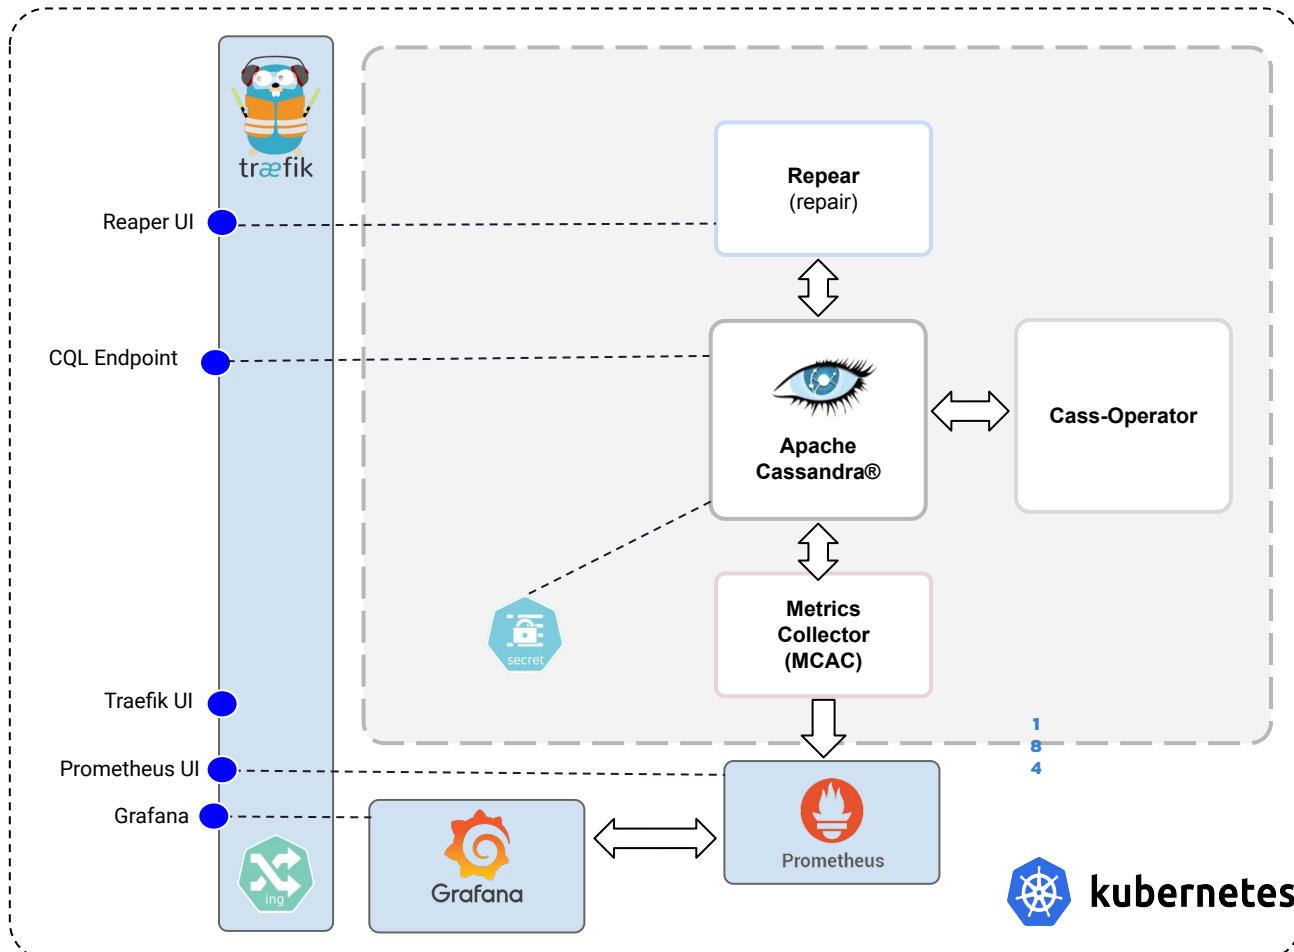
The screenshot shows the 'Repair' section of the Cassandra Reaper interface. It displays a table of repair tasks:

| ID | State | Cluster | Keyspace | CFs | Incremental | Repaired | Actions |
|--------------------------------------|-------------|------------|----------|-----|-------------|----------|---|
| 90bc2910-5022-11e7-b7cc-039e197701f7 | PAUSED | newcluster | test | | false | 162/201 | Activate Abort |
| a8854a70-4ffc-11e7-8025-872a3b38b99e | NOT_STARTED | newcluster | test | | false | 0/201 | Activate Delete |

The screenshot shows the 'Repair' section of the Cassandra Reaper interface. It displays a table of completed repair tasks:

| ID | State | Cluster | Keyspace | CFs | Incremental | Repaired | Actions |
|--------------------------------------|-------|------------|----------|-----|-------------|----------|------------------------|
| 832607a0-46fc-11e7-b3c0-2dc25dfb993e | DONE | twcs308 | booya | | true | 3/3 | Delete |
| c7d63470-4ff7-11e7-a8a1-b957dacf6ec2 | DONE | newcluster | test | | false | 12/12 | Delete |
| 3bd206c0-54c5-11e7-a8f0-7336a93cc91f | DONE | twcs308 | booya | | true | 3/3 | Delete |
| f44c1700-54c4-11e7-a8f0-7336a93cc91f | DONE | twcs308 | booya | | true | 3/3 | Delete |

Repair



Medusa

Orchestrated backup and restore for Apache Cassandra and DSE.

Designed by  Spotify®

Implemented by  THE LAST PICKLE

Maintained and supported by  DATASTAX®

Storage backends

- NFS shares
- AWS S3
- Google Cloud Storage
- IBM Cloud Object Storage
- Azure Blob Storage



amazon
S3



Google Cloud Storage

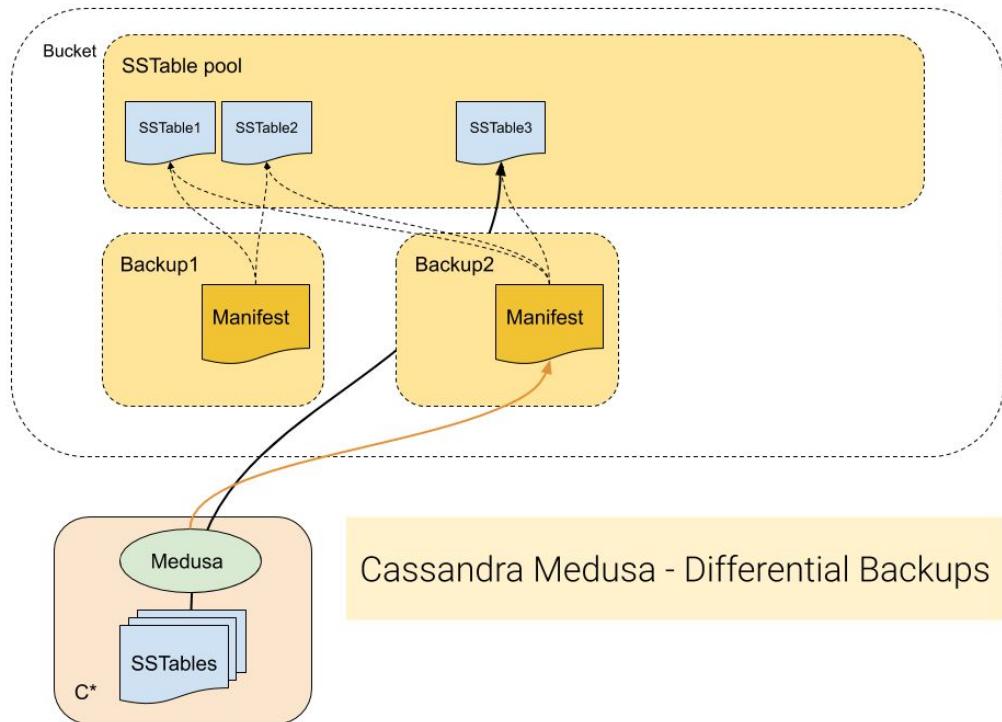
Microsoft Azure
Blob Storage

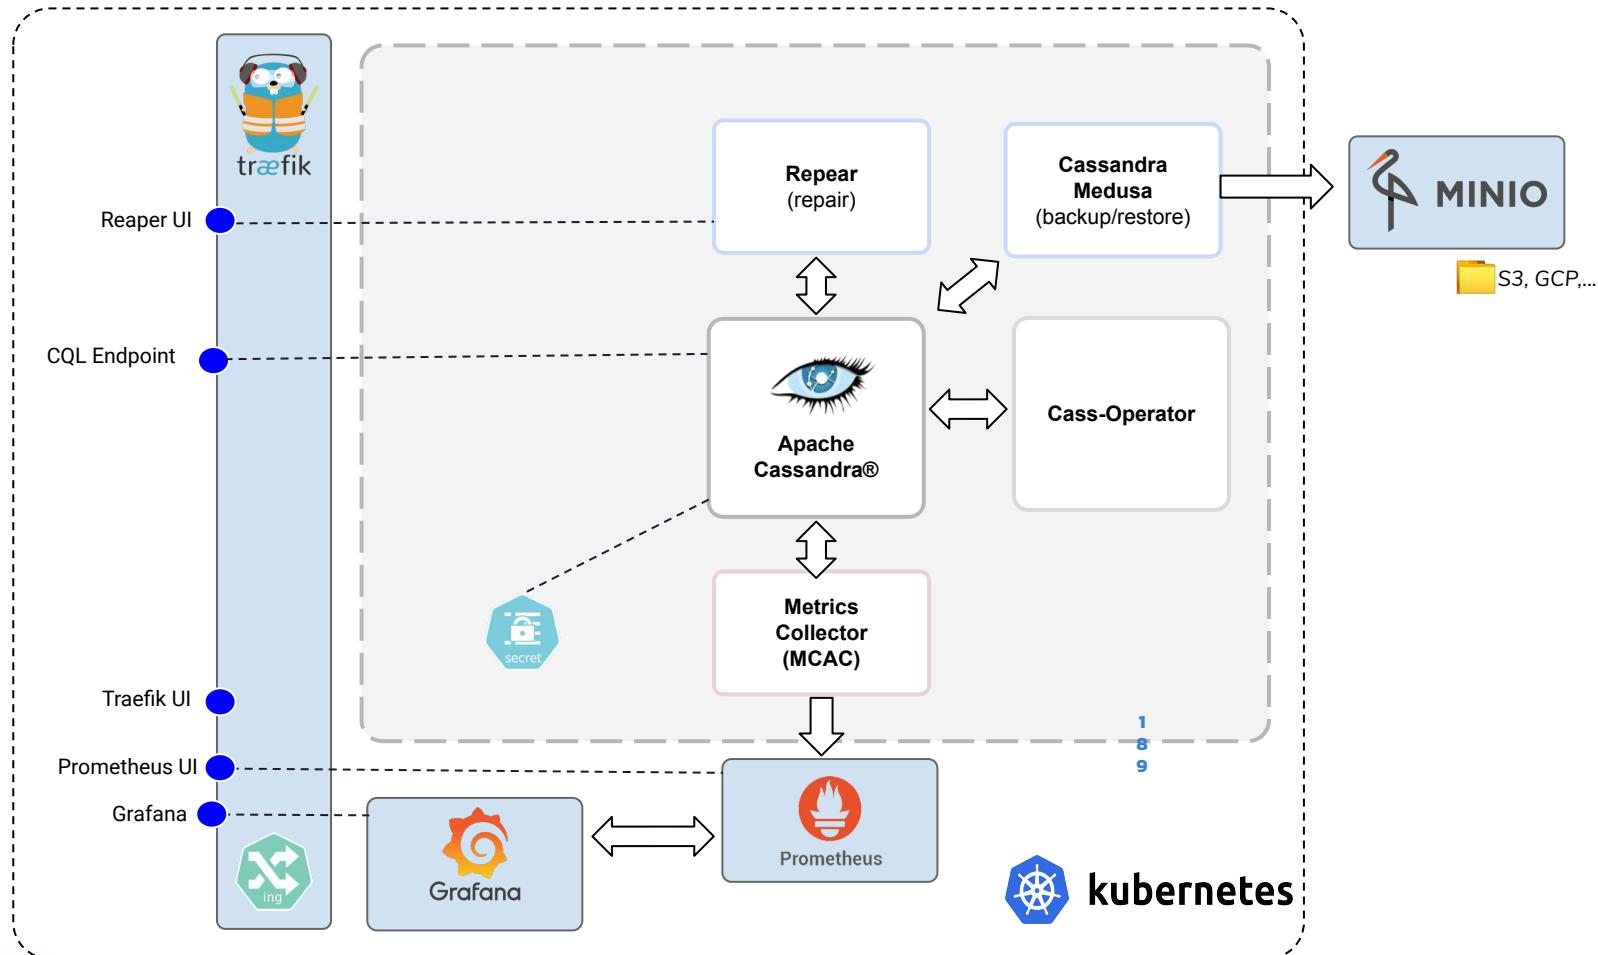


IBM Cloud
Object Storage

Differential backups (default)

SSTables are uploaded only **once** in their lifetime to save space and bandwidth





Stargate Overview

An open source API framework for data

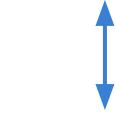
Stargate makes it easy to use a database for any application workload by adding plugin support for new APIs, data types, and access methods



MICROSERVICES
DEVELOPERS



STARGATE



cassandra

Stargate.io

stargate / stargate Public

Code Issues 246 Pull requests 13 Discussions Actions Projects Wiki Security ...



STARGATE



OPEN SOURCE DATA GATEWAY

Stargate is an open source data gateway that sits between your app and your databases. Stargate brings together an API platform and data request coordination code into one OSS project.

Get Started  Github

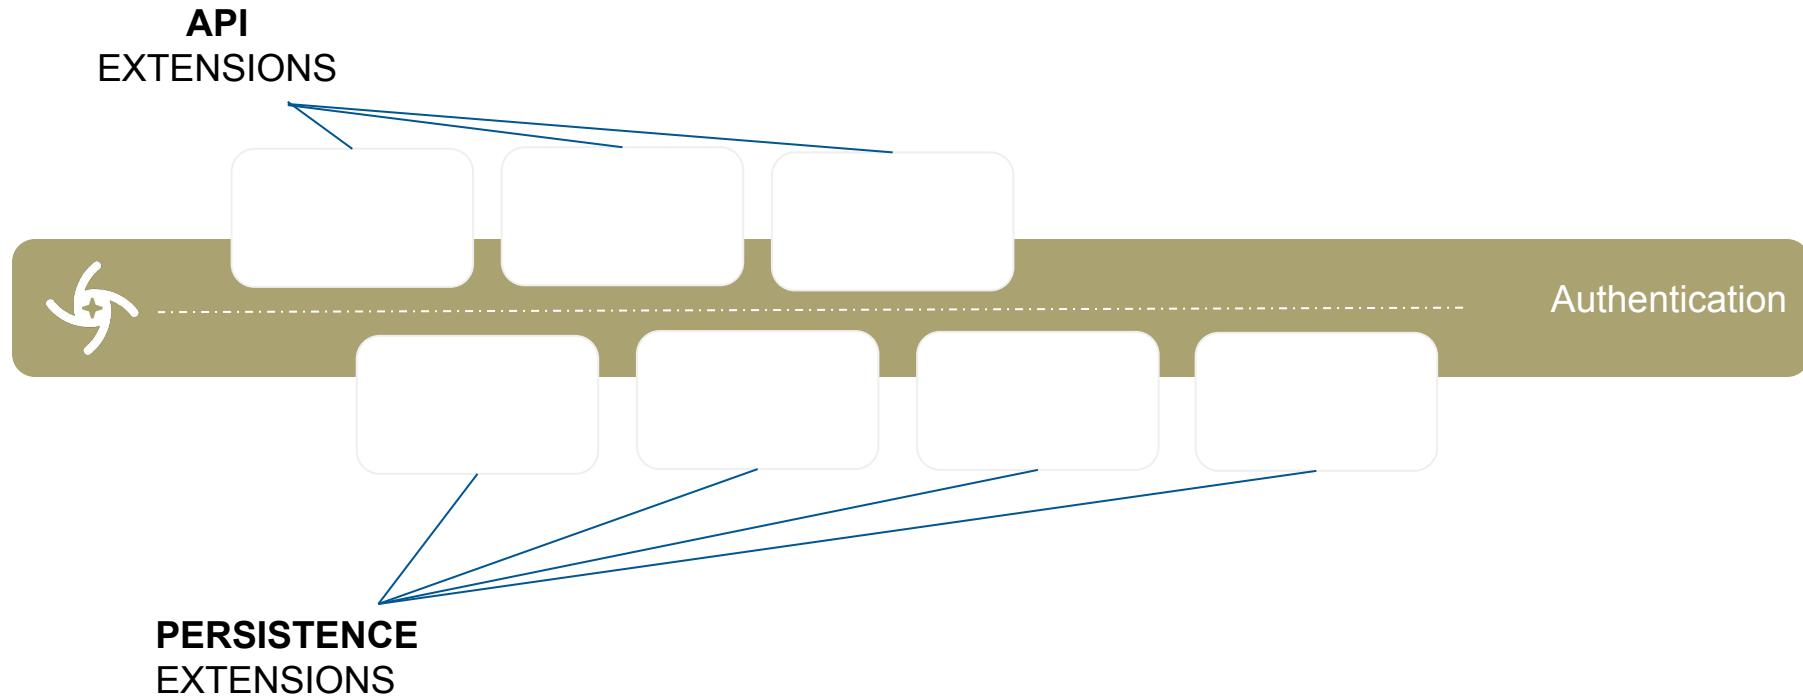
Stargate is the official data API for DataStax Astra and DataStax Enterprise! Try It Now

grpc Reference example for gRPC java client (#1236) yesterday

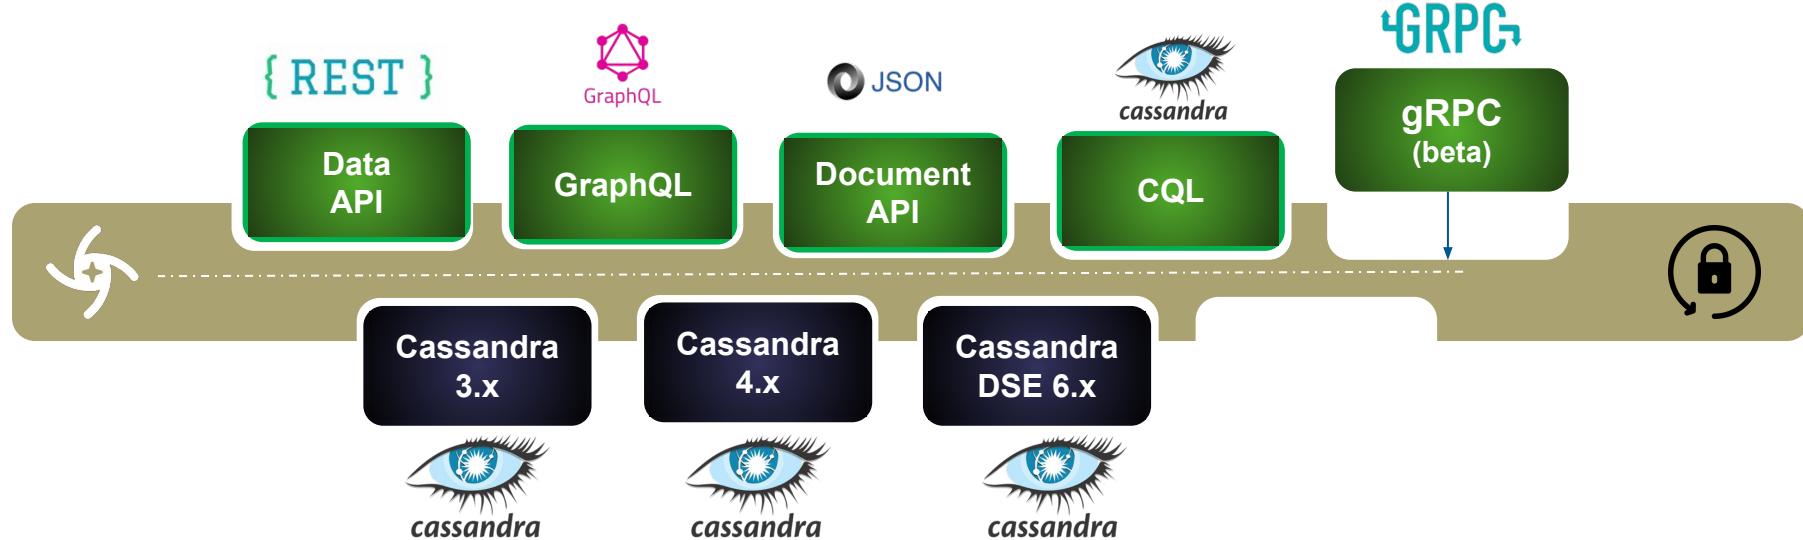
Contributors 30 

DataStax Developers

API and Persistence Extensions



API Extensions and Persistence Extensions



9042

8090

8082

8080

8084

CQL

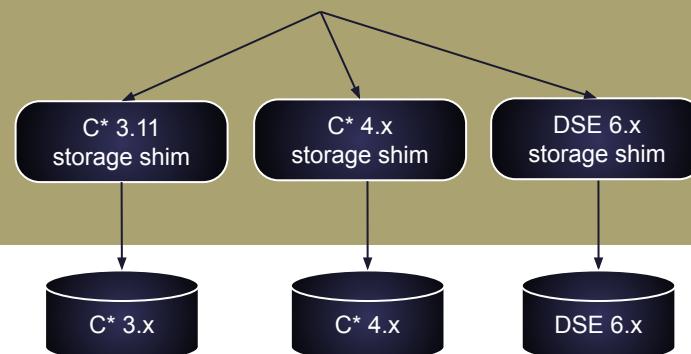
GRPC

Data + Document APIs

GraphQL APIs



STARGATE



DataStax Developers

9042

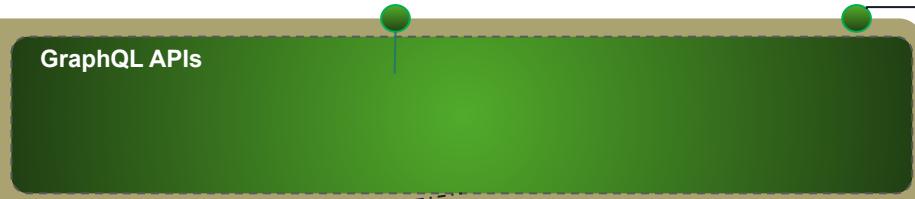
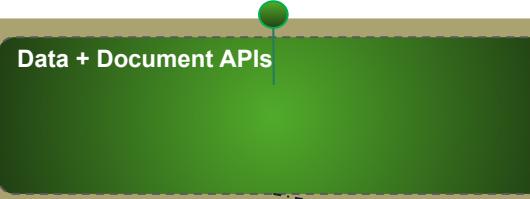
8090

8081

8082

8080

8084



Authentication Service

Rate Limiting Service

Metrics Services



Persistence Service

C* 3.11
storage shimC* 4.x
storage shimDSE 6.x
storage shim

STARGATE

OSGI
Service Registry

OSGI™ Java

DataStax Developers

9042

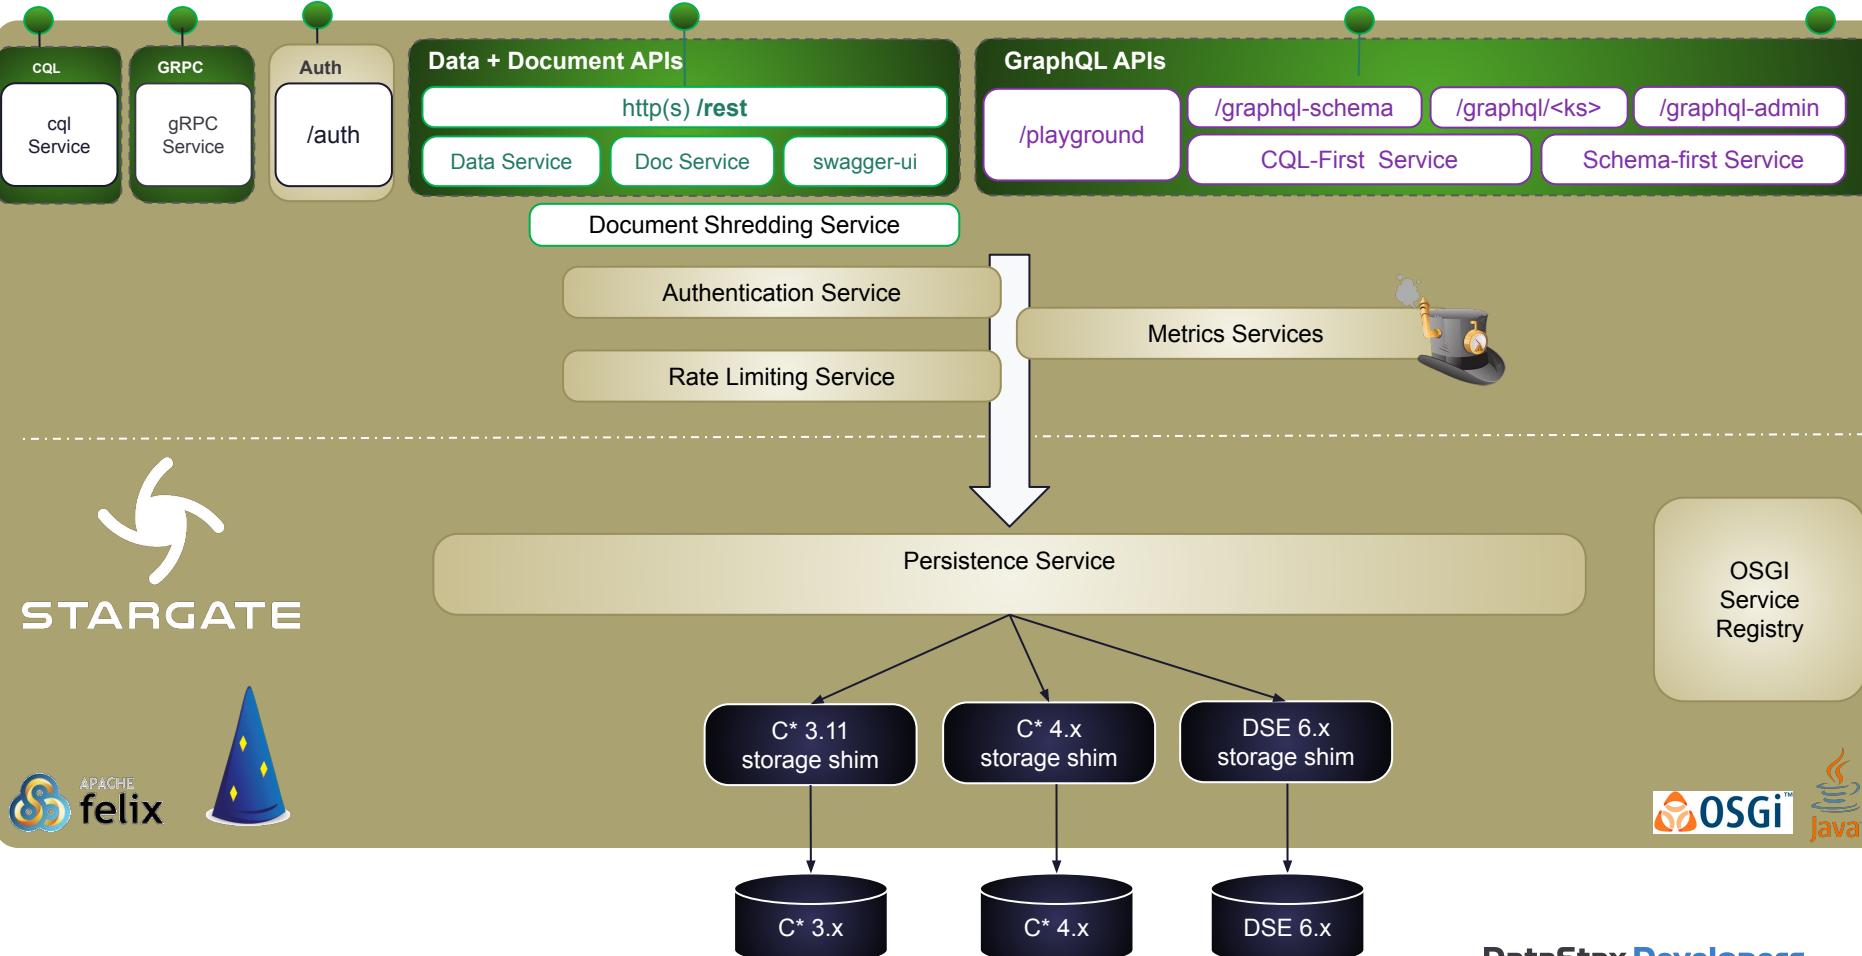
8090

8081

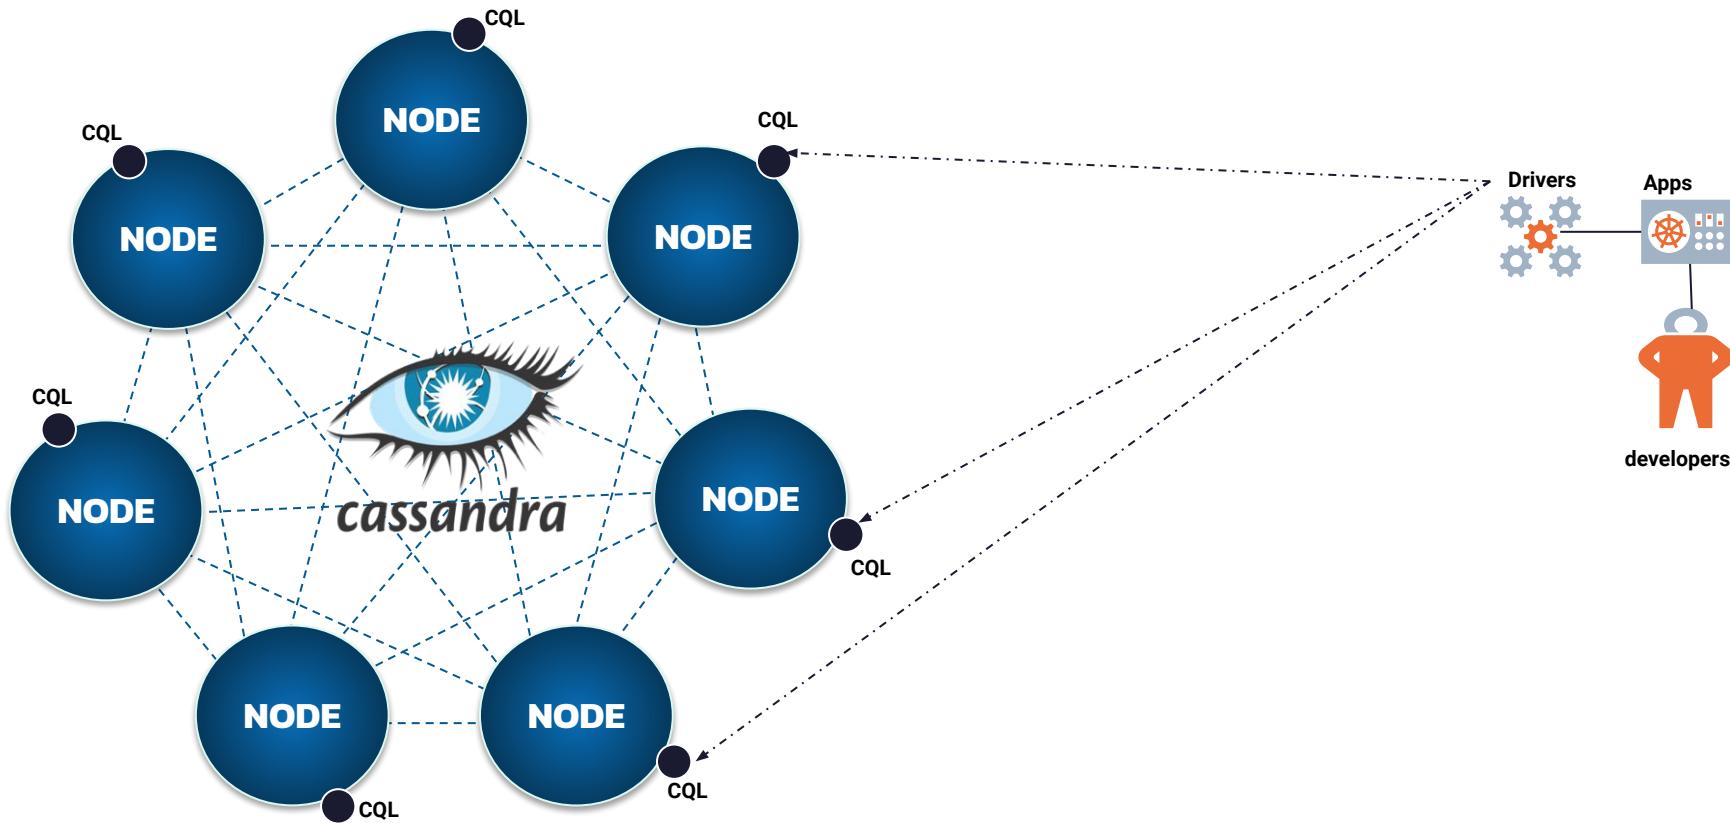
8082

8080

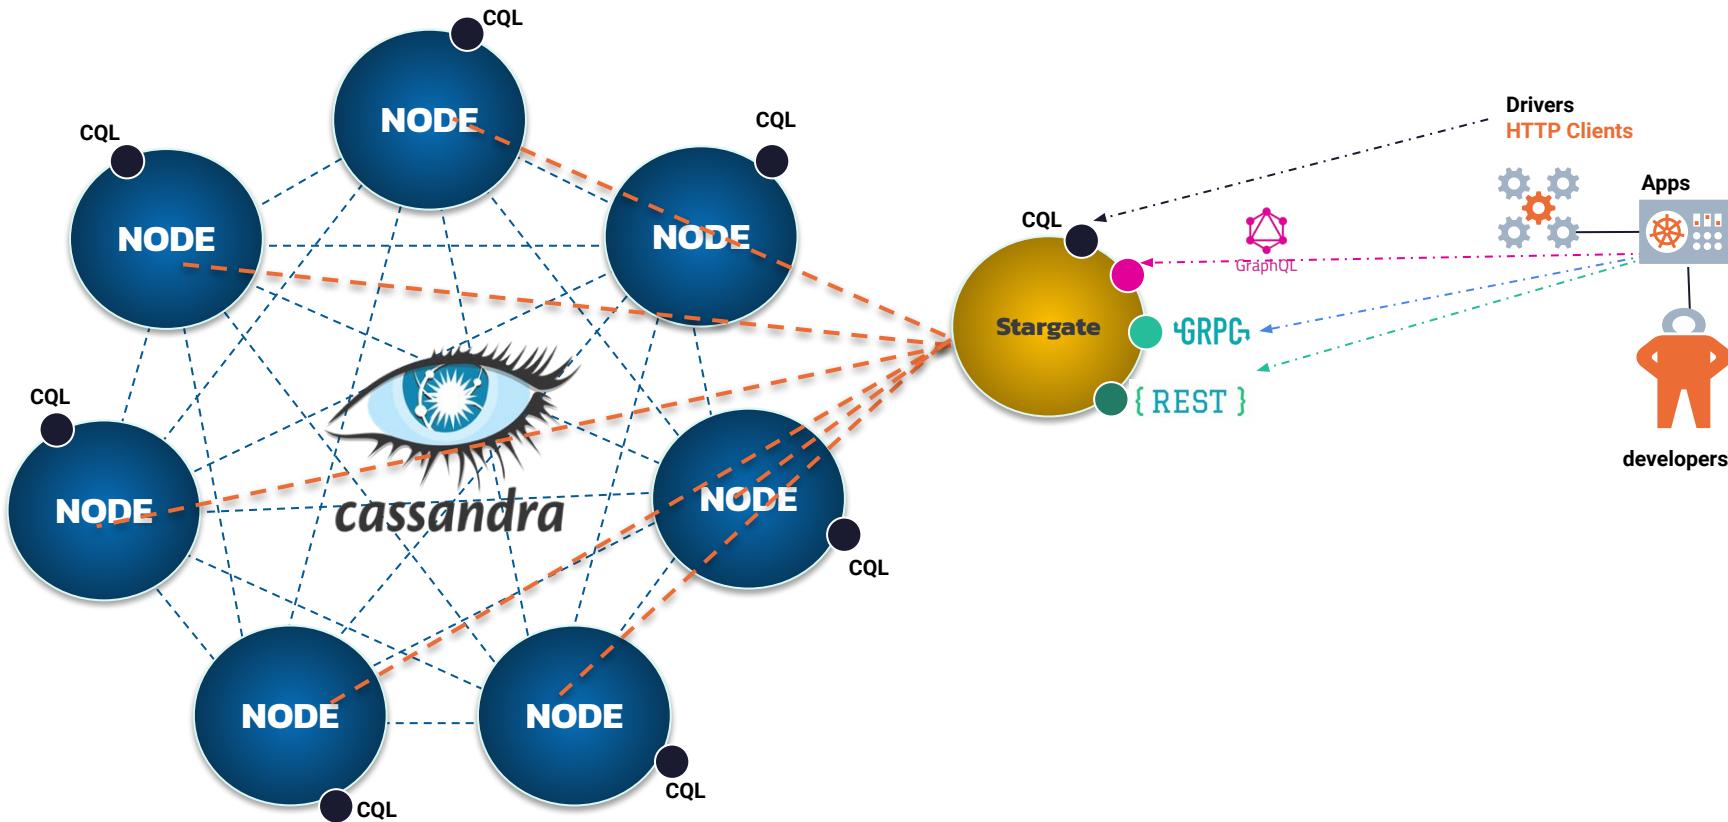
8084



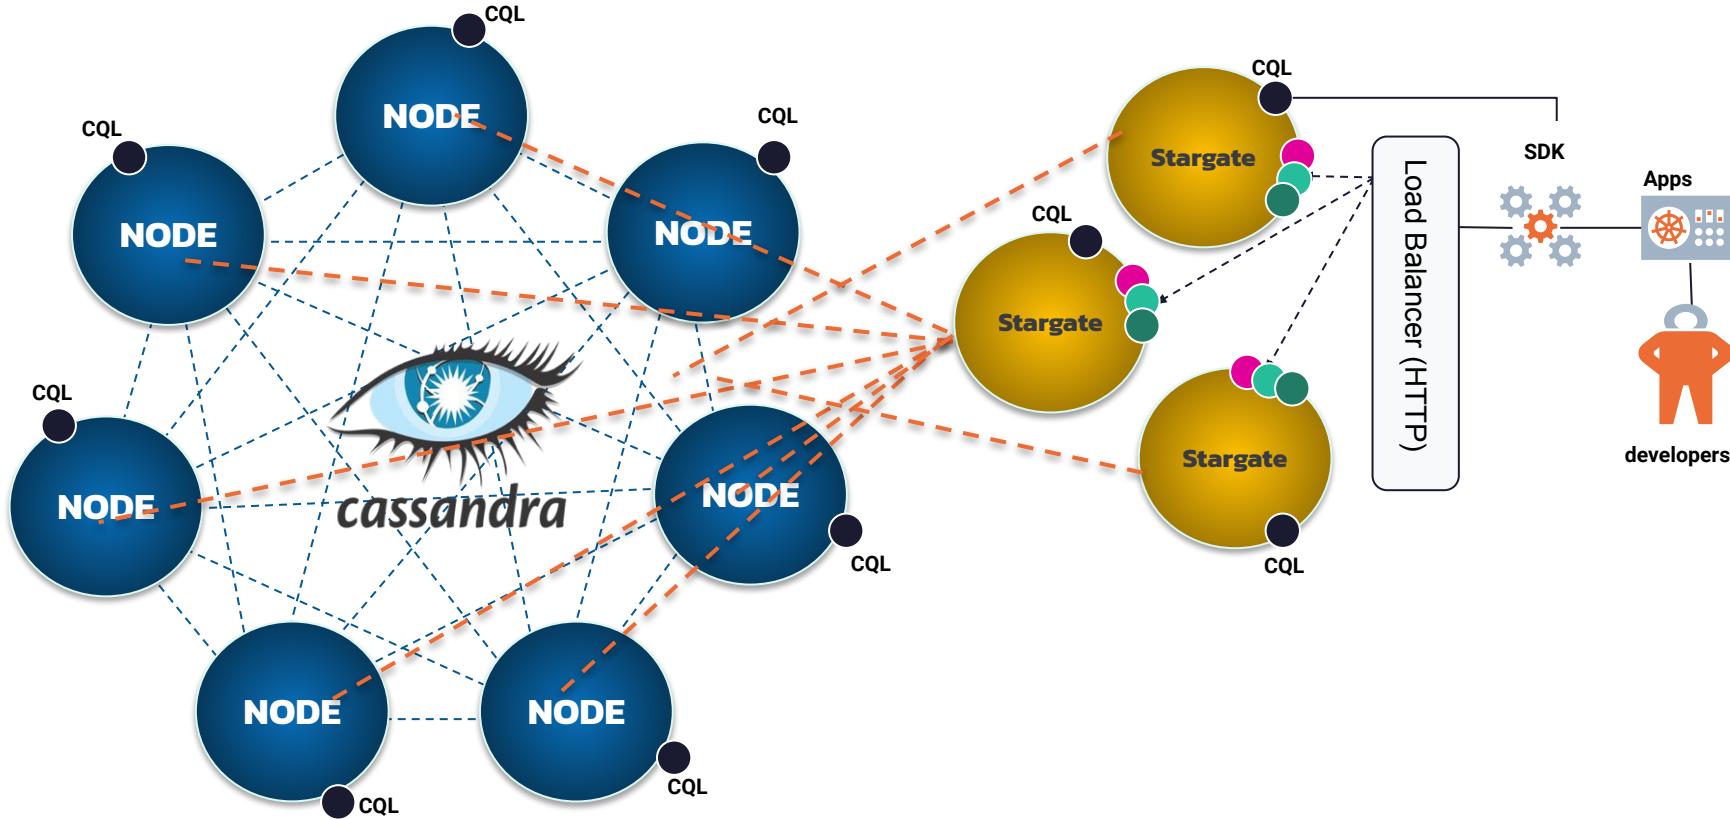
Connecting to your cluster (Before)



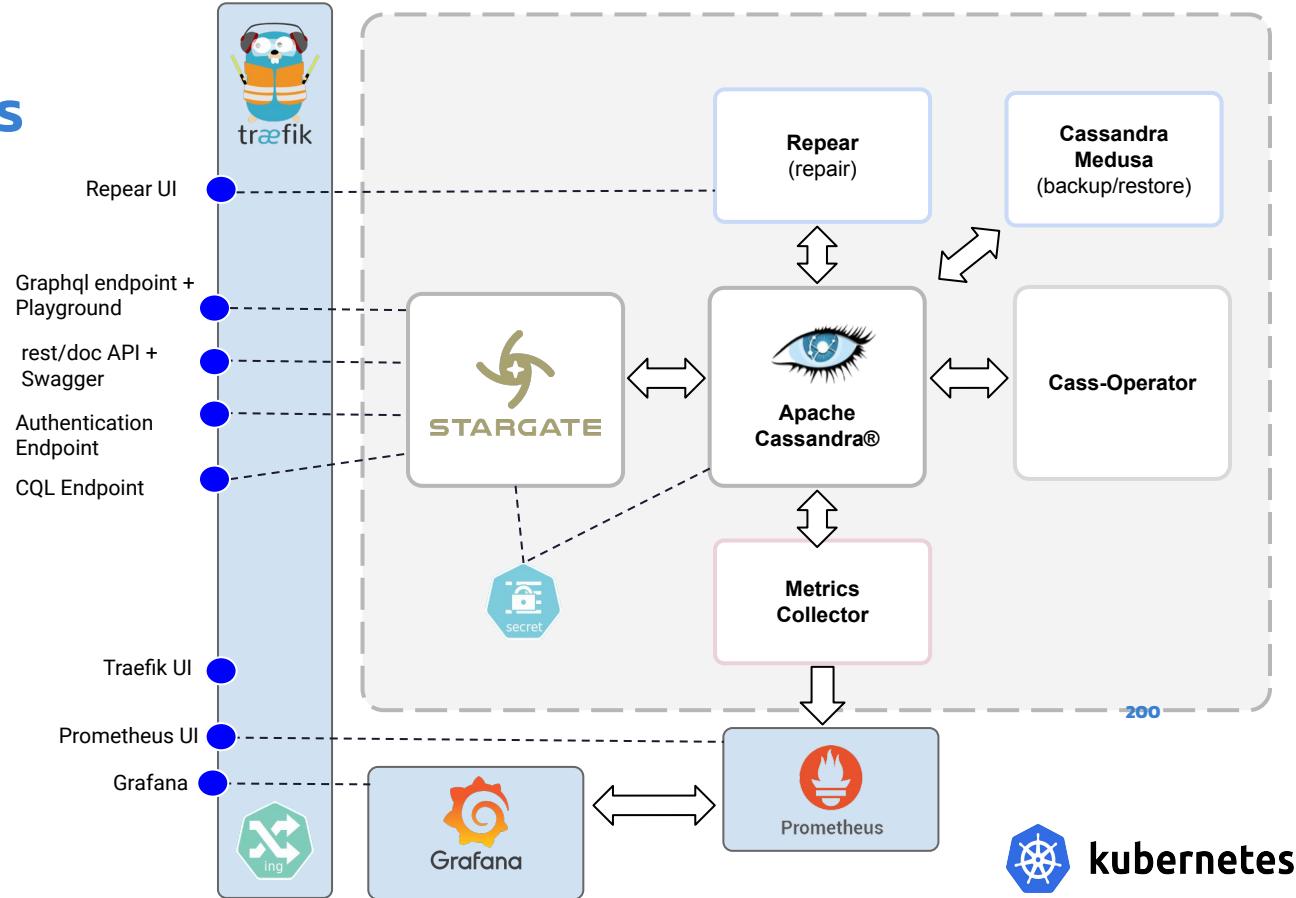
Connecting to your cluster (with Stargate)



Connecting to your cluster (with Stargates)



Data Access



Join the K8ssandra Community!



Sign up for newsletter ->



forum.k8ssandra.io



Discord live chat



@k8ssandra



github.com/k8ssandra

The screenshot shows the homepage of k8ssandra.io. The header features the K8SSANDRA logo with a stylized anchor icon above the word "K8SSANDRA". The navigation bar includes links for About, Get Started, Docs, Community, Forum, Blog, and Events. A search bar is located in the top right corner. The main hero section has a blue background with white stars and text. It reads "K8ssandra" in large letters, followed by a description: "Built on the rock-solid Apache Cassandra® NoSQL database, K8ssandra brings together a complete operational data platform for Kubernetes including APIs, monitoring, and backups." Below this are two buttons: "Get Started" (red) and "Read the Docs" (white). A yellow sidebar on the left contains the text "K8ssandra and the Community" and a snippet from an article by Séan Bogard. The main content area features a yellow banner with the same text. To the right, there's a "Latest Articles" section with three entries: "K8ssandra and the Community" by Séan Bogard (04.28.2021), "KubeCon / CloudNativeCon – North America 2021" by Kate Sandra (04.29.2021), and "Running a Cloud-Native Data Tier with K8ssandra hosted by DataStax" by Jeff Carpenter.

Hands-ON

<https://github.com/k8ssandra/workshop>

Apache Cassandra Operations in Kubernetes



Introduction

Deploy a Sample Application with Cassandra in Kubernetes

START LESSON



Running Cassandra in Docker

Use Docker containers to deploy an Apache Cassandra™ cluster



Managing Cassandra Clusters in Kubernetes Using the Cass-Operator

Use a Kubernetes operator to deploy and manage a Cassandra Cluster

START LESSON



Monitoring Cassandra Clusters in Kubernetes with Prometheus and Grafana



Access Cassandra Via Stargate APIs



Running K8ssandra

Learn how K8ssandra provides a

Data on Kubernetes Community – DoK Community



Community

Schedule

Blog

Dok Day 2021

An open community for
data on Kubernetes



Can you run databases on K8s?

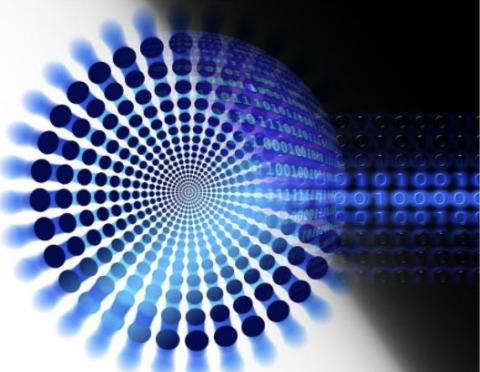
THE NEW STACK Podcasts ▾ Events Ebooks ▾ Newsletter Sponsorship

Architecture ▾ Development ▾ Operations ▾ 

DATA / KUBERNETES / SPONSORED / CONTRIBUTED

A Case for Databases on Kubernetes from a Former Skeptic

5 Apr 2021 9:00am, by Christopher Bradford



Kubernetes is everywhere. Transactional apps, video streaming services and machine learning workloads are finding a home on this ever-growing platform. But what about databases? If you had asked me this question five years ago, the answer would have been a resounding “**No!**” — based on my experience in development and operations. In the following years, as more resources emerged for stateful applications, my answer would have changed to “*Maybe*,” but always with a qualifier: “It’s fine for development or test environments...” or “If the rest of your tooling is Kubernetes-based, and you have extensive experience...”

But how about today? Should you run a database on Kubernetes? With complex operations and the requirements of persistent, consistent data, let’s retrace the stages in the journey to my current answer: “In a cloud native environment? **Yes!**”



Christopher Bradford

Christopher has a passion for enabling efficiency through automation. From promoting effortless scaling via Cassandra to DevOps pipelines with infrastructure automation and containers, he is here to get work done and enable operators to rest easy.

<https://thenewstack.io/a-case-for-databases-on-kubernetes-from-a-former-skeptic/>

Pushing Helm to its limit

DATASTAX MODERATED / OPEN SOURCE / UNSORED / CONTRIBUTED

We Pushed Helm to the Limit, then Built a Kubernetes Operator

8 Sep 2021 6:04am, by [Jeff Carpenter](#) and John Sanda



Jeff Carpenter

Jeff works in developer relations at DataStax. He has worked on large-scale systems in the defense and hospitality industries and is co-author of 'Cassandra: The Definitive Guide.'



John Sanda

John is a DataStax engineer working on the K8ssandra project. He is passionate about Cassandra and Kubernetes and loves being involved in open source. Prior to joining DataStax, John worked as an Apache Cassandra consultant.

<https://thenewstack.io/we-pushed-helm-to-the-limit-then-built-a-kubernetes-operator/>

<https://bit.ly/cassandra-workshop>

Join our 10k Discord Community

The Fellowship of the RINGS

A screenshot of the Fellowship of the Rings Discord server's main chat room. The left sidebar shows various channels: server-conduct, upcoming-events, useful-resources, moderator-only, WORKSHOPS (with a pinned #main-chat-room), Breakout room 1, Breakout room 2 (with member Shruti), Breakout room 3, Hidden Channel, and TOPICS (with pinned #grafana-cassandra-data... and #docker-training). The main channel has a pinned message from Jack Fryer (@14:30) announcing a new workshop: "thank you everyone! we hope to see you again in a future workshop! 🎉". Below it is a message from Cedrick Lunven (@16:18) about a Spring + Cassandra talk. A large embed from Eventbrite for the "Cloud-Native Workshop: Connecting Cassandra and Kubernetes!" is displayed, showing the event date (March 3 or March 4), time (Cloud-native | Beginner), and a Level Up Developers logo. The right sidebar lists PRESENTER—3 (Aleks Volochnev, David Jones-Gilardi, jscarp), HELPER—1 (John Sanda), and EN LIGNE—222 (Abhiprada, Absurdism, hiya, Adalberto, aditya_dhunna, adnaneCord, Adrigunz, Aemilius Gaurav, Aemilius gaurav, Aguvas, ajscilingo, akashTaxvisor).

thank you everyone! we hope to see you again in a future workshop! 🎉

Jack Fryer Aujourd'hui à 14:30
NEW WORKSHOP ALERT

Cassandra meets Kubernetes!

Come and meet K8ssandra! A cloud-native distribution of Apache Cassandra™ built for running on Kubernetes

<https://www.eventbrite.co.uk/e/cloud-native-workshop-connecting-cassandra-and-kubernetes-tickets-142078180663>

Eventbrite

Cloud-Native Workshop: Connecting Cassandra and Kubernetes!

Come and meet K8ssandra! A cloud-native distribution of Apache Cassandra™ built for running on Kubernetes

LIVE hands-on workshop >/>

Apache Cassandra™ meets Kubernetes!

March 3 or March 4 | Cloud-native | Beginner

LEVEL UP Developers

Cedrick Lunven Aujourd'hui à 16:18
Hey Community I will be live tonight for 50 min talk on Spring + Cassandra to the Virtual Java User Group
<https://www.youtube.com/watch?v=nuyPKDQn1gl> come and say hi ?

Envoyer un message à #main-chat-room

<https://www.datastax.com/workshops>

Workshops Hebdomadaires

The image shows a composite view of the DataStax Developers YouTube channel and its website.

YouTube Channel Header: Features a white astronaut icon, the text "LEVEL UP with the DataStax Developers", and a "SUBSCRIBE" button.

YouTube Channel Navigation: Includes links for "ACCUEIL", "VIDÉOS", "PLAYLISTS", "COMMUNAUTÉ", "CHAÎNES", and "À PROPOS".

YouTube Channel Content: Shows a grid of video thumbnails. Examples include "Microservices with Cassandra + Spring" (2:23:56), "Advanced Data Modeling" (2:41:51), "An Introduction to Apache Cassandra™" (1:28:14), "Streaming to C* with Apache Pulsar!" (2:11:40), "Building Microservices with Cassandra + Spring" (1,1 k vues), "Advanced Data Modeling in Apache Cassandra™" (1,3 k vues), "Polka DataOps Meetup: Introduction to Apache..." (151 vues), "Bring Streaming to Cassandra with Apache..." (703 vues), "Introduction Cassandra™" (1,6 k vues), "Taking your K8s app to the Cloud!" (1 k vues), "RESTful Data Access from a Spring Boot App" (215 vues), "Ask Me Anything Hackathon Special" (667 vues), "Connecting Apache Cassandra™ and Kubernetes" (1,1 k vues), and "An Easy Back-end API with Ap..." (2,3 k vues).

Website Header: Features the "LEVEL UP with the DataStax Developers" logo and the text "for weekly content on building".

Upcoming Live Events: A section listing five workshops:

- Apache Cassandra™ Certification Preparation** (Multiple Dates | NoSQL | Beginner)
- Build Microservices with Apache Cassandra™!** (Feb 17 or Feb 18 | NoSQL | Beginner)
- Certification Exam Preparation Workshop** (Multiple Dates)
- Cloud-Native Workshop: Build Spring Microservices with Apache Cassandra™** (Multiple Dates)
- Learn how to build a Serverless Game!** (Feb 24 or Feb 25 | Game Development | Beginner)

Bottom Right Event: "Cloud-Native Workshop: Build Microservices w/ Apache Cassandra™ + Quarkus!" (THU MAR 11 2021 | Register Now)

DataStax Developers

Thank you!



GitHub

@clun



@clunven



DataStax Developers